

# CPSC 213

---

## January 8th - Unit 0: Introduction

- Staff
  - Instructor: Geoffrey Tien
    - [gctien@cs.ubc.ca](mailto:gctien@cs.ubc.ca)
  - Course Coordinator: Gale Chen
    - [cpsc213-admin@cs.ubc.ca](mailto:cpsc213-admin@cs.ubc.ca)
    - Administrative queries
- Course Logistics
  - Canvas
  - PrarieLearn
  - iClicker Cloud
  - Piazza
- Grading
  - Participation
    - 1% - lab participation (80% of 12 labs)
    - 1% - lecture iClicker questions for 80% lectures
  - Assignments
    - 25% for 10 assignments
    - best of 8 of A1-A9 + A10
  - Quizzes
    - 33% for 3 quizzes
    - in-person in CBTF (self-schedule 1-hour slot during quiz weeks)
  - Final exams
    - 40% final exam
- MUST
  - Achieve at least 50% on quiz/final average
  - Submit at least 80% of homework
- What is a computer?
  - Computer
    - Stores information and can perform a simple set of operations

- Uses existing information plus **work** to produce a new information
- Performs a sequence of operations determined by stored instructions
  - stored instructions are part of the machine's stored information
- Can be modeled by a Turing machine
- Computer System
  - Abstraction of a computing machine
  - Used in various combinations by programmers
  - Abstraction examples
    - Hardware instruction-set architectures
    - Operating systems
    - Compilers and high-level languages
    - Middleware and libraries
  - Programming
    - Building new abstractions in software
    - Use existing abstractions
- Why care?
  - Mental model of computation is rooted in the machine
  - Understanding abstraction will
    - Understand and respect limits
    - Be more flexible and powerful
    - Holistic view of computational stack

```

public class A {
    static int s = 0;
    int i;
    public A() {
        s = s + 1;
        i = s;
    }
    public static void main(String[] args) {
        A a0 = new A();
        A a1 = new A();
        a0 = a1;
        a1.i = a1.i + 1;
        System.out.printf("%i, %i", a0.i, a1.i);
    }
}
  
```

In the code on the left, what prints when **main()** executes?

A. 0, 2  
 B. 1, 3  
 C. 2, 3  
 D. 3, 3  
 E. Something else

**Draw a picture!**

- Correct Answer: D
- Explanation: Notice that s is a singleton (static variable), and a0, a1 both refer to the same object (a0 = a1).
- Inside the machine

- All states are stored in a single, **main memory**
- Each *byte* has an unique numeric *address*
- A *variable* is a name for a memory location that has an **address, size, value**
  - Statistically allocated variables: address is *specified at compilation time*
  - Dynamically allocated variables: address is *determined at runtime*
- A *pointer* is a variable whose value is a memory address

```
public class B {
    int val;
    B(int val) {
        this.val = val;
    }
    int add(B x, B y) {
        this.val = this.val + x.val;
        this.val = this.val + y.val;
        return this.val;
    }
    public static void main(String[] args) {
        B a = new B(4);
        B b = a;
        B c = a;
        System.out.println( a.add(b,c) );
    }
}
```

In the code on the left, what is printed when `main()` executes?

- A. 4
- B. 8
- C. 12
- D. 16
- E. 20

[Draw a picture!](#)

- Correct Answer: D
- Explanation: All a, b, and c are referring to the same object. Thus when `add()` method was called, `this.val` will be updated for all three references. `this.val = 4 → 8 → 16`; a, b, and c are pointers
- Inside the machine (Continued)
  - A new object is created only when the "new" keyword is called
  - Object is identified by its *memory address*, a number
  - **Pointer variables** are stored in memory too
    - Their value is the **memory address** of some object

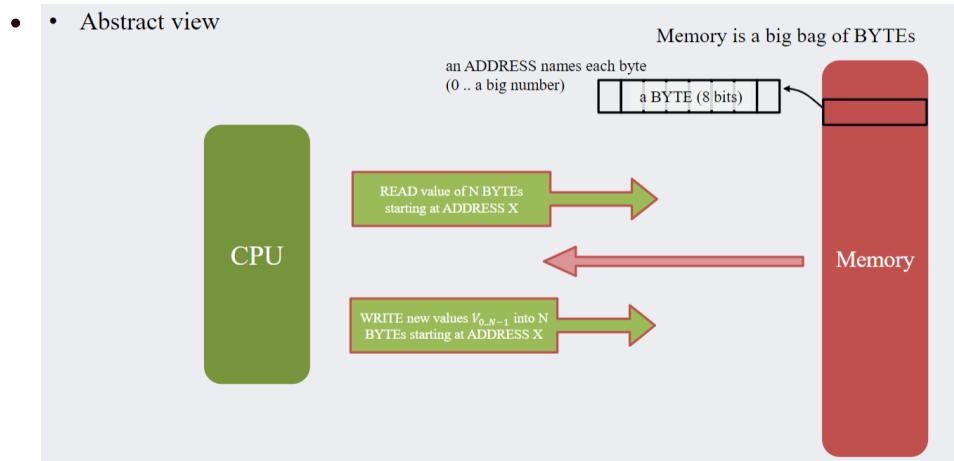
## January 10th - Unit 1a: Numbers and Memory

- Simple Computing Machine
  - Memory
    - Stores data encoded as bits
    - Program instructions and state
  - CPU
    - Reads instruction and data from memory

- Performs specified computation and writes result back to memory
- Example
  - $C = A + B$
  - Memory: stores add instruction, variables A, B, C
  - CPU: reads instruction and values of A, B; adds values; writes result to C
- Memory (Endianness will be talked about later)

• ADDRESS	BIT 0	1	2	3	4	5	6	7
0000	*	*	*	*	*	*	*	*
0001	*	*	*	*	*	*	*	*
0010	*	*	*	*	*	*	*	*
0011	*	*	*	*	*	*	*	*

- CPU: READ value of  $N$  BYTES starting at ADDRESS X;  
Memory: Gives the requested data to CPU  
CPU: WRITE new values of  $N$  BYTES starting at ADDRESS X



- Naming
  - Unit of addressing is **byte** (8 bits)
  - Every byte of memory has an **unique address**
  - Some machines (this class: SM213) have 32-bit memory addresses, some have 64
    - $2^{32}$  bytes in memory is about 4GB.
- Access
  - Many things are too big to fit in a single byte
  - CPU accesses memory in **contiguous, power-of-two-sized** chunks of bytes
  - Address of a chunk is address of its **first byte**

• # BYTES	# BITS	C	JAVA	ASM
• 1	8	char	byte	b - byte

# BYTES	# BITS	C	JAVA	ASM
2	16	short	short	w - word
4	32	int	int	l - long
8	64	long	long	q - quad

- We will only use 32-bit integers (4 bytes)
- Numbers and Representation
  - Integer value of a chunk of bytes
    - Binary to Decimal
  - The bits themselves
    - Use power-of-two sized way to identify addresses
    - Binary: **0b**
    - Octal: **0**
    - Hexadecimal: **0x**
      - 4 bits = hex "digit"
      - 32-bit address = 8-hexit address
      - 1 byte = 2 hexits

Which of the following statements is true?

- A. The Java constants **16** and **0x10** are exactly the same integer
- B. The Java constants **16** and **0x10** are different integers
- C. Neither of the statements above is always true
- D. I don't know
- E. I'm just choosing E for participation credit

- • Correct Answer: A
- Hexadecimal operations
  - For memory addresses
  - Addition / Subtraction
    - Convert to decimal
    - Calculate in hex directly
    - Subtraction: 2's complement
    - Carry for addition, Borrow for subtraction

- Object A is at address **0x10d4**, and object B is at **0x1110**. They are stored contiguously in memory (i.e. they are adjacent to each other without any gap). How big is A?

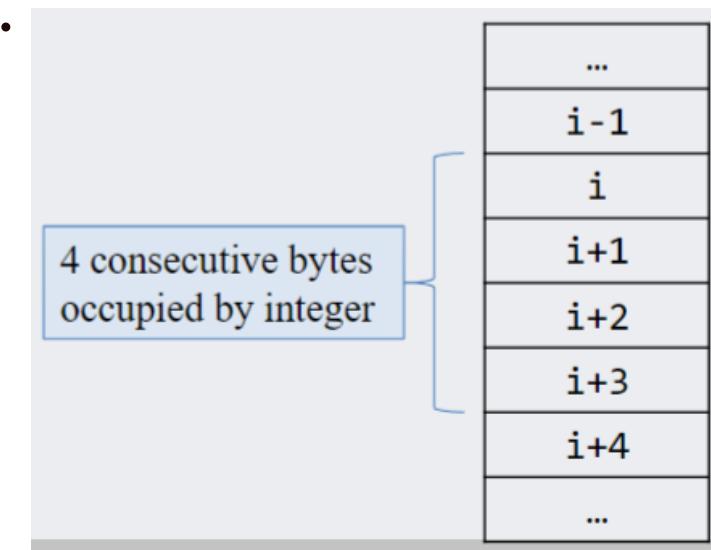
- A. 16 bytes
- B. 48 bytes
- C. 60 bytes
- D. 80 bytes
- E. Not enough information to tell



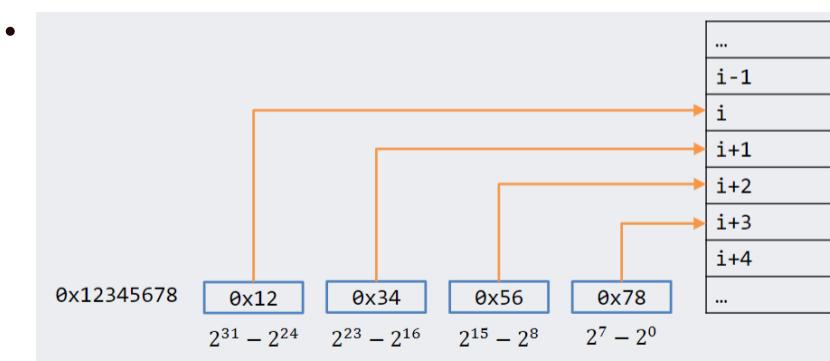
- • Correct Answer: C
- Explanation: **0x1110 - 0x10d4 = 0x003c = 60<sub>10</sub>**

## January 12th - Unit 1a: Cont.d

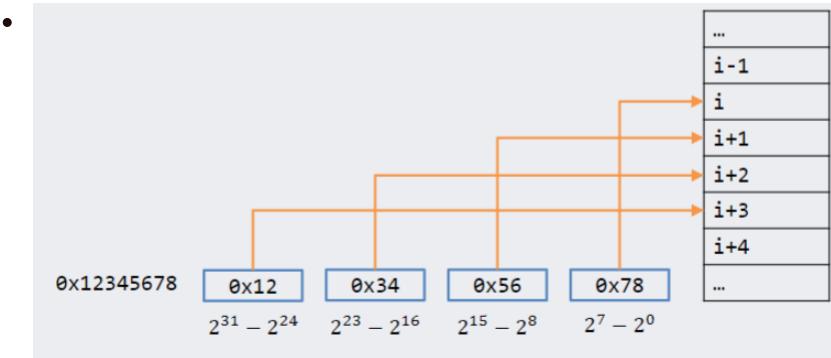
- Making Integers from Bytes
  - Each memory address hold 1 byte
  - Architectural decision
    - Integers = multiple bytes of memory
    - 32-bit integer (**int** in Java / C)
      - 4 bytes in memory
      - If memory address start at  $i$ , we also need addresses  $i+1, i+2, i+3$



- Endianness
  - Big-endian
    - First byte stored is the most significant byte (bit 31 - 24)
    - Used in old IBM servers, network connections



- If integer is signed, first byte can tell us +/-.
- Little-endian
  - First byte is the least significant byte (bit 8-0)
  - Used in most Intel-based system



- First byte can tell us even/odd; In C, we can use 0/1 represent F/T, then the least significant byte is useful, as we only need the least significant byte.

- What is the Little-endian 4-byte integer value at address 0x4?

- A. 0xc1406b37  
 B. 0x1c04b673  
 C. 0x73b6041c  
 D. 0x376b40c1  
 E. 0x739a8732

Address	Value
0x0:	0xfe
0x1:	0x32
0x2:	0x87
0x3:	0x9a
0x4:	0x73
0x5:	0xb6
0x6:	0x04
0x7:	0x1c

- Correct Answer: B
- Address alignment
  - 4-byte integer can be placed at almost any address
  - Forcing address alignment is better for hardware
    - Consider an 8-byte memory, if we place the integer on addresses 1, 2, 3, then we cannot place another 4-byte integer, this leads to memory waste (decrease in efficiency)
    - Aligned address: addresses whose numeric value is a **multiple** of the object size (arrays and structs can be more complicated)
    - Smaller data types can fit inside larger things
      - $2 \times$  shorts =  $1 \times$  int
      - Significant for arrays, as we can move in multiples of the object size.
    - Some CPUs do not support misaligned address
- Power-of-2 alignment
  - Address computation in arrays is achieved by shifting **bits**.
    - If we want to access element  $i$ , where each element occupies  $2^j$  bits, then:

```
&a[i] == &a[0] + i * (size == 2^j) == &a[0] +
i << j
```

- Memory implementation detail
  - Memory is organized internally into larger chunks called **blocks**
  - Suppose a block is 16 bytes
  - Every memory access requires accessing one of these blocks
  - This relates to memory caches (caching - provide these blocks to CPU when requested)
- Alignment and Memory access
  - CPU
    - Read/write  $N$  bytes starting at address  $A$
  - Memory converts above request into
    - Read/write  $N$  bytes starting at the  $O^{th}$  byte of block  $B$ 
      - $O$  - block offset;  $B$  - block number
      - Blocks are numbered so that Block 0 contains address 0...15
      - If generally, Block  $B$  contains address  $X \dots X + blocksize - 1$
      - Then,  $X \leq A \leq X + blocksize - 1$
    - So to access this address given  $A$ , we need to find the Block number and Block offset, then consider a tuple that satisfy some function of  $A$ :  $(B, O) = f(A)$
  - The process is simplified if
    - Block size = power of 2
    - Object size ( $N$ ) = power of 2
    - Absolute address ( $A$ ) = aligned to  $N$
    - e.g Assume Block size = 8;  $N = 4$ . We want 4 bytes starting at address  $A = 84_{10}$ 
      - $(B, O) = f(84) \rightarrow (B, O) = (10, 4)$ , want 4 bytes from 10<sup>th</sup> block offset by 4 bytes
        - $B = A/8, O = A \% 8$
        - Since block size is power of 2, calculating  $B$  is just a right shift
        - Calculating  $O$  is obtained by **masking** the least significant bits (in this case, the last 3 bits)

## Jan. 15th - Unit 1a: Cont.d

- Address alignment

- Which of the following statement(s) is/are true?
  - The address 6 ( $110_2$ ) is aligned for addressing a **short**
  - The address 6 ( $110_2$ ) is aligned for addressing an **int** (i.e. 4 bytes)
  - The address 20 ( $10100_2$ ) is aligned for addressing a **long** (i.e. 8 bytes)
  - Two or more of the above
  - None of the above

- Correct answer: A
- Explanation: **short** are block sizes of 2 bytes, and  $6 \% 2 = 0$ ; the other data types have remainders.

- Which of the following statement(s) is/are true?
  - The address  $0x14$  is aligned for addressing a **short**, but not an **int**
  - The address  $0x14$  is aligned for addressing a **short** and **int**, but not a **long**
  - The address  $0x14$  is aligned for addressing **short**, **int**, and **long**
  - None of the above
  - P-A-R-T-I-C-I-P-A-T-I-O-N!

- Correct answer: B
- Explanation:  $0x14 = 20$ ; only 2 bytes and 4 bytes divide 20.

- Extending integers (Small type to large type)

- Extending

- Increasing the number of bits (bytes) used to store an integer

- ```
byte b = -6;
int i = b;
out.printf("b: 0x%x, %d; i: 0x%x, %d", b, b,
i, i);
```

- Output: "b: 0xfa, -6; i: 0xfffffffffa, -6"

- Signed extension

- Used with signed numbers (e.g. in Java)
- Copies signed bit into upper empty bits of the extended number (use 2's complement to ensure the same number)

- Zero extension

- Used with unsigned number (e.g. in C)
- Sets upper empty bits to 0
- Can be forced by **masking** using bitwise AND operator

- ```
int u = b & 0xff;
printf("u: 0x%x, %d\n", u, u);
```

- Output: "u: 0xfa, 250"
- Truncating integers (Large type to small type)
  - We can use fewer bits to represent a value
  - ```
int i = -6;
byte b = i;
out.printf("b: 0x%02x, %d; i: 0x%02x, %d");
```
  - This piece of code segment will run with error "possible loss of precision"
  - Cast explicitly to hide such a warning:
  - ```
byte b = (byte) i
```
  - This would run without big issues
  - If we use  $i = 256 = 0x100$ ,  $b$  would only be  $0x00 = 0$ ; If we use  $i = 128$ ,  $b$  would be treated as  $0x80$ , with leading 0s, which is  $b = -128$ .
- Bit operations
  - $a \ll b$ : shift all bits in  $a$  to the left  $b$  times, fill remaining right bits with zero
  - $a \gg b$ : shift all bits in  $a$  to the right  $b$  times
    - C: if  $a$  is unsigned, zero extension, otherwise sign extension
    - Java:  $\gg$  sign extension;  $\ggg$  zero extension
  - $a \& b$ : bitwise AND
  - $a | b$ : bitwise OR
  - $a ^ b$ : bitwise XOR
  - $\neg a$ : bitwise NOT
  - Bit shifts
    - Shifting to the left: multiplication by  $2^b$
    - Shifting to the right: division by  $2^b$
    - Negative numbers? (Signed shift ✓)

- What is the value of *i* after this Java statement executes?

```
int i = 0xff8b0000 & 0x00ff0000;
```

- A. 0xffff0000
- B. 0x0000008b
- C. 0x008b0000
- D. 0xff8b0000
- E. None of the above

- Correct answer: C

- What is the value of *i* after this Java statement executes?

```
int i = 0x8b << 16;
```

- A. 0x8b
- B. 0x0000008b
- C. 0x008b0000
- D. 0xff8b0000
- E. None of these

- Correct answer: C (Assumed leading 0s)

- What is the value of *i* after this Java statement executes?

```
int i = 0x0000008b << 16;
```

- A. 0x008b
- B. 0x0000008b
- C. 0x008b0000
- D. 0xff8b0000
- E. None of these

- Correct answer: C

- What is the value of *i* after this Java statement executes?

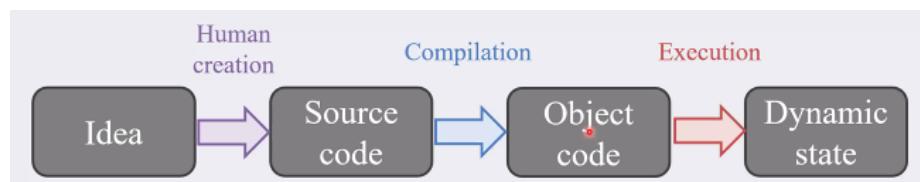
```
int i = ((byte) 0x8b) << 16;
```

- A. 0x8b
- B. 0x0000008b
- C. 0x008b0000
- D. 0xff8b0000
- E. None of these
- F. I don't know

- Correct answer: D
- Explanation: When 0x8b is casted into byte, it is interpreted as a negative value as there is a leading 1. Thus, when left shift 16 bits, we then would yield an integer 0xff8b0000 to preserve sign.

## Jan. 17th - Unit 1b: Static Scalars and Arrays

- Develop a model of computation, Analyze the processor, Explore the language
- The CPU
  - CPU executes instructions
  - Execution
    - Fetch: Load the next instruction from memory
    - Decode: Figure out what the instruction is
    - Execute: Do what the instruction specifies
  - These executing steps are looped over and over again
  - Instructions
    - Read/load a value from memory
    - Write/store a value to memory
    - Add/subtract/AND/OR two numbers
    - Shift a number
    - Control flow
    - Some operations are carried out in Arithmetic & Logic Unit (ALU), used in execution stage
- Phases of computation
  - Human creation: design program in high-level language
  - Compilation: Convert high-level language to machine-executable text
  - Execution: a physical machine executes the code
    - Parameterized by input values **unknown** at compilation
    - Producing output values that are **unknowable** at compilation



- Object code is binary machine code
  - Anything that **the compiler can compute** is called ***static***
- Anything that can **only be discovered during execution** is called ***dynamic***

- Consider the code below:

```
int a;
a = 5;
```

The value assigned to variable **a** is:

- A. static
- B. dynamic
- C. neither static nor dynamic
- D. it depends

- Correct answer: A
- Explanation: Compiler gives space to *a*, then assigned value to *a*, so it is static, as it is in the source code.

- Assume that a variable of type **int** is assigned a value that is read from the user through an input dialog.

This value is:

- A. static
- B. dynamic
- C. neither static nor dynamic
- D. it depends

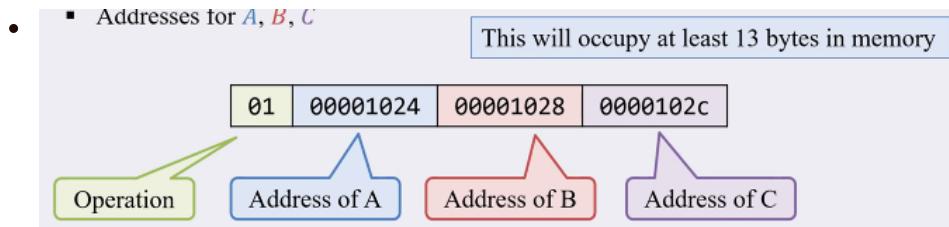
- Correct answer: B
- Explanation: Compiler still makes space for the variable, but cannot assign a value to the variable until runtime

- The processor (CPU)

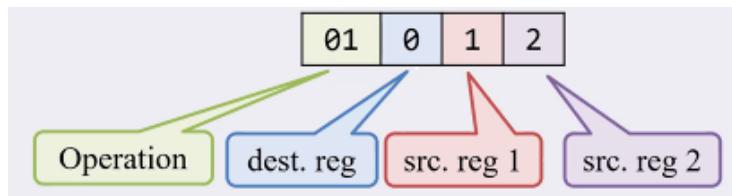
- Implements a set of instructions
- Each instruction is implemented using **logic gates**
  - Built from **transistors**
- Instruction design philosophies
  - **RISC**: fewer and simpler instructions makes processor design simpler
  - **CISC**: having more types of instructions and more complex instructions allows for shorter/simpler program code and simpler compilers

- First proposed instruction: +

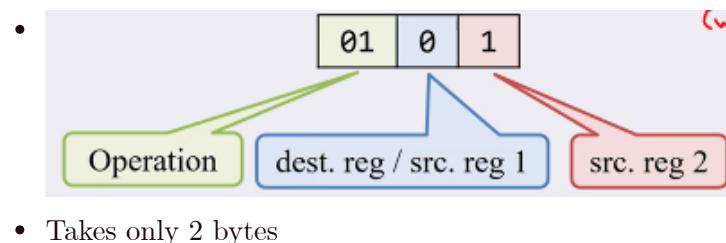
- We want to do:  $A \leftarrow B + C$
- Instruction parameters
  - Addresses of *A, B, C*: 32-bits
- Instruction is encoded as a sequence of bits
  - Operation name
  - Addresses for variables



- Accessing memory is SLOW
  - ~100 CPU cycles for every memory access
  - goal: fast programs avoid accessing memory when possible
- Big instructions are costly
  - Memory addresses are big (instruction is then big)
  - Big instructions → big programs
  - Reading memory is slow
  - Large instructions use more CPU resources (transfer, storage)
- General purpose registers
  - Register file
    - Small, fast memory in CPU itself
    - roughly **single-cycle** access
  - Register
    - Named by a number (0-15)
    - Size: architecture's common integer
- Improve using registers
  - Memory instructions handle only memory
    - Load data from memory to register
    - Store data from register to memory
  - Accessing data in registers
    - Small and fast
  - Then it becomes:



- Takes only about 3 bytes
- To further improve instruction size; we can share register for one source and destination



- $B \leftarrow B + C$
- Minimize interaction with memory will increase processing speed
  - Suppose we are adding more than 2 variables
  - Direct memory access will require the above non-register-involved process to run at least twice
  - However, in the register-involved method, only memory accesses are loading each variable's address and store the final result to a designated memory
  - This will reduce the times to access memory, speeding up the program.
- Special purpose registers
  - e.g. Program Counter (PC): contains address of next instruction to execute
  - Instruction Register (IR): contains the machine instruction that has just been fetched from memory
    - IR needs to be large enough to contain the longest instruction

## Jan. 19th - Unit 1b: Cont.d

- Instruction Set Architecture (ISA)
  - Formal interface to a processor implementation
    - Defines the instructions the processor implements
    - Defines the format of each instruction (in binary/hexidecimal)
  - Types of instructions
    - Math and logic
    - Memory access
    - Control transfer: "goto" and conditional "goto"
  - Design
    - Instruction format
      - Sequence of bits: **opcode** and **operand** values
      - All represented as numbers → defines machine code
    - Assembly language
      - Symbolic representation of machine code
- Representing Instruction Semantics
  - RTL: pseudo language to describe semantics
    - Right-to-left format, describes machine steps
  - Syntax
    - Line form:  $LHS \leftarrow RHS$

- $LHS$  is a memory or register that receives a value
- $RHS$  is a constant, memory, register, or expression on 2 registers
- $m[a]$  is a memory in address  $a$
- $r[i]$  is a register with number  $i$  (0 to 7, for general purpose registers)
- Example 1

•  $r[0] \leftarrow 0x2000$

Register 0 receives 0x2000 (as an address)

$r[1] \leftarrow m[r[0]]$

Register 1 receives memory whose address is  $r[0]$

$r[2] \leftarrow r[2] + r[1]$

Register 2 increases by the value of Register 1

- ```
// Assume at 0x2000 address, we have a value
of 7
int n = 7;
int r_2 = 10;

// A pointer that represents Register 0, that
stores the address of n
int *r_0 = &n;
// r_0 now stores 0x2000

// Register 1 stores the value at the address
stored in Register 0, which in this case, the
value of variable n
int r_1 = *r_0;
// r_1 now stores 7;

r_2 += r_1;
// r_2 would now store the value of 17.
```

- Example 2

- Assume at address 0x1234, value 22 is stored, consider the following instructions

•  $r[0] \leftarrow 10$

$r[1] \leftarrow 0x1234$

$r[2] \leftarrow m[r[1]]$

$r[2] \leftarrow r[0] + r[2]$

$m[r[1]] \leftarrow r[2]$

- Step by step:

- Register 0 stores the value of 10
- Register 1 stores the address of 0x1234
- Register 2 stores the value at the address stored in Register 1 (0x1234), which is 22
- Register 2 increases the value by the value stored in Register 0, which is 10, resulting in 32
- Value at memory stored in Register 1 (0x1234) now stores the value in Register 2, which is 32
- Variables
  - Named storage locations for values
  - Features
    - Name - static
    - Type/Size - static
    - Scope (where in the code we can access the variables) - static
    - Lifetime (availability of the variable at a certain part of the program) - dynamic
    - Memory location (address) - dynamic/static
    - Value - dynamic/static
  - Static variables (built-in types)
    - In Java
      - Static data members are allocated to a class, not an object
      - Store built-in scalar types or references to arrays or objects
    - In C
      - Global variables and any other variable declared static
      - Can be static scalars, arrays or structs or pointers
      - ```
int a;
int b[10];

void foo() {
    a = 0;
    b[a] = a;
}
```
- Allocation
  - Assigning a memory location to a variable
- Static vs. Dynamic computation
  - Global/static variables can exist before program starts and live until after it finishes

- Compiler allocates variables, giving them a constant address
- No dynamic computation is required to allocate, they just exist
- Compiler tracks free space during compilation - it is in control of program's memory

• Assume `a` is a global variable in C. When is space for `a` allocated? In other words, when is its address determined?

- A. The compiler assigns the address when it compiles the program
- B. The compiler calls the memory to allocate `a` when it compiles the program
- C. The compiler generates code to allocate `a` before the program starts running
- D. The program locates available space for `a` before it starts running
- E. The program locates available space as soon as the variable is used for the first time

- Correct answer: A
- Explanation: B. Memory itself is not involved; C. Compilation happens even before program starts running; D. The program is already running; E. Dynamic allocation

## Jan. 22nd - Unit 1b: Cont.d

- What is a proper RTL instruction for  
`a = 0;`

- A. `r[0x1000] ← 0`
- B. `m[0x1000] ← 0`
- C. `0x1000 ← r[0]`
- D. `m[r[0x1000]] ← 0`
- E. `0x1000 ← m[0]`

Static memory layout

0x1000: value of a
...
0x2000: value of b[0]
0x2004: value of b[1]
...
0x2024: value of b[9]

- Correct Answer: B
- Variables - Cont.d
  - Static variable access
    - Observations
      - Value of `a` is generally unknown at compilation time
      - Even though addresses of `b[0...]` are known **statically**, which one being used is **dynamic**
      - In general, given `b[a] = a`, the compiler does not know the address of `b[a]`
    - Array access is computed from a **base address** and **index**
      - Address = base + offset
      - Offset = index × element size

- The base address and element size are static, but the index is dynamic

• What is a proper RTL instruction for  
 $b[a] = a;$   
Assume that the compiler does not know the current value of  $a$ .

- A.  $m[0x2000+m[0x1000]] \leftarrow m[0x1000]$   
B.  $m[0x2000+4*m[0x1000]] \leftarrow m[0x1000]$   
C.  $m[0x2000+m[0x1000]] \leftarrow 4*m[0x1000]$   
D.  $m[0x2000]+4*m[0x1000] \leftarrow m[0x1000]$   
E.  $m[0x2000]+m[0x1000] \leftarrow m[0x1000]$

Static memory layout  
0x1000: value of a  
...  
0x2000: value of b[0]  
0x2004: value of b[1]  
...  
0x2024: value of b[9]

- Correct answer: B
- Alternative for instructions sets
  - For  $a = 0$ 
    - Option 1:
      - Static address and value:  $m[0x1000] \leftarrow 0x0$
      - 9 bytes (instruction code + two integers)
      - If the instruction is  $a = b$ ,  $b$  can come from user input, so we may require dynamic value
    - Option 2:
      - Static address, dynamic value: $r[0] \leftarrow 0x0$   
 $m[0x1000] \leftarrow r[0]$ 
      - 5 bytes for immediate (constant) value instruction (instruction + register + one integer)
      - 5 bytes for memory instruction (instruction + register + one integer)
      - op codes are always half a byte, register is always half a byte
    - Option 3:
      - Dynamic address, dynamic value: $r[0] \leftarrow 0x0$   
 $r[1] \leftarrow 0x1000$   
 $m[r[1]] \leftarrow r[0]$ 
      - 2 bytes for memory instruction: instruction code + two registers
  - For  $b[a] = a$ 
    - Option 1:
      - Static base address, index, and value: $m[0x2000 + 4*m[0x1000]] \leftarrow m[0x1000]$ 
      - 13 bytes (instruction + 3 integers)

- Option 2
    - Calculate address explicitly

```

r[0] ← 0x1000
r[1] ← m[r[0]]
r[3] ← r[1] << 2
r[2] ← 0x2000
r[2] ← r[2] + r[3]
m[r[2]] ← r[1]

```

  - Array access are very common: calculating address every time is costly
- Option 3
    - Dynamic base address, index, and value:

```

r[0] ← 0x1000
r[1] ← m[r[0]]
r[2] ← 0x2000
m[r[2] + r[1] << 2] ← r[1]

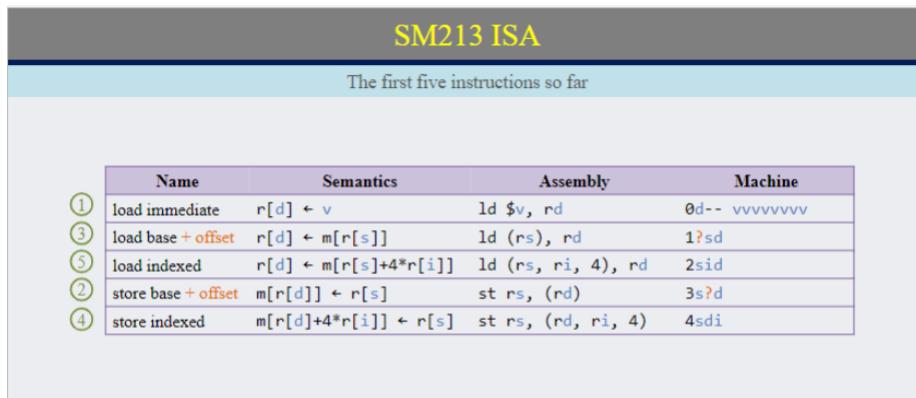
```

  - The indexed memory instruction require 2 bytes: instruction + 3 registers

## Jan. 24th - Unit 1b: Cont.d

- ISA design goals (RISC paradigm)
  - $a = 0; m[0x1000] \leftarrow 0x0$ 
    - $r[0] \leftarrow 0x0$
    - $r[1] \leftarrow 0x1000$
    - $m[r[1]] \leftarrow r[0]$  (store)
  - $b[a] = a; m[0x2000 + 4 * m[0x1000]] \leftarrow m[0x1000]$ 
    - $r[0] \leftarrow 0x2000$
    - $r[1] \leftarrow m[r[0]]$  (load)
    - $r[2] \leftarrow 0x2000$
    - $m[r[2] + 4 * r[1]] \leftarrow r[1]$
  - Minimizes the number of memory instructions in ISA (at most 1 memory access per instruction)
  - Minimize the total number of instructions in ISA
  - Minimize the size of each instruction
- Instructions so far
  - Load a constant value into a register:  $r[0] \leftarrow v$
  - Store a value in a register into memory at some address:  $m[r[y]] \leftarrow r[x]$

- Load a value in memory into a register:  $r[x] \leftarrow m[r[y]]$
- Store value in a register into memory at address in register \* 4 + base address
  - $m[r[z] + 4 * r[x]] \leftarrow r[y]$
- Load value in memory at address in register \* 4 + base address into register
  - $r[y] \leftarrow m[r[z] + 4 * r[x]]$

- 

• s: source; i: index; d: destination; ?: 4 bits to encode offset

- Code translation

- High-level code  $\rightarrow$  RTL  $\rightarrow$  Assembly (1 : 1 correspondence)  $\rightarrow$  Machine code (1 : 1 correspondence)
- .pos: position directive, next item places into memory at the specified address.

<code>int i; int a[10]; a[2] = a[i];</code>	<b>Name</b>	<b>Semantics</b>	<b>Assembly</b>
	load immediate	$r[d] \leftarrow v$	<code>ld \$v, rd</code>
	load base + offset	$r[d] \leftarrow m[r[s]]$	<code>ld (rs), rd</code>
	load indexed	$r[d] \leftarrow m[r[s]+4*r[i]]$	<code>ld (rs, ri, 4), rd</code>
	store base + offset	$m[r[d]] \leftarrow r[s]$	<code>st rs, (rd)</code>
	store indexed	$m[r[d]+4*r[i]] \leftarrow r[s]$	<code>st rs, (rd, ri, 4)</code>

• Which correctly implements  $a[2] = a[i];$ ?

A. `st ($a, $i, 4), ($a, 2, 4)`

B. `ld ($a, $i, 4), r0  
st r0, ($a, 2, 4)`

C. `ld $a, r0  
ld $i, r1  
ld (r1), r1  
ld (r0, r1, 4), r2  
st r2, 2(r0)`

D. `ld $a, r0  
ld $i, r1  
ld (r1), r1  
ld (r0, r1, 4), r2  
ld $2, r1  
st r2, (r0, r1, 4)`

E. None of these / I don't know

- Correct Answer: D
- Explanation: C. offset must be multiple of 4; D.

```

•      r0 <- &a;           # r0: address
      of a
      r1 <- &i;           # r1: address
      of i
      r1 <- *r1;          # r1: value of
      i
      r2 <- m[r0 + 4 * r1]; # r2: value of
      a[i]
      r1 <- 2;            # r1: 2
      m[r0 + 4 * r1] <- r2 # a[2]: value
      of r2

```

- SM213 ISA
  - Architecture
    - Register file: eight 32-bit general purpose registers (0-7)
    - CPU: one cycle per instruction (fetch + execute)
    - Main Memory: byte addressed, big endian integers
  - Instruction format
    - 2 or 6 byte instructions (each character is a hexit)
      - x-01, xxsd, x0vv, or x-sd vvvvvvvv
      - x is op-code
      - -: means unused
      - s and d are operand register numbers
      - vv, vvvvvvvv are immediate/constant values

## Jan. 26th - Unit 1b: Cont.d

- Machine code
  - [addr:] x-01 [vvvvvvvv]
    - addr: sets starting address for subsequent instructions
    - x-01 hex value of an instruction with opcode x and operands 0 and 1
    - vvvvvvvv hex value of optional extended value
- Assembly code
  - [label:] [instruction | directive] [# comment]
    - directive :: (.pos number) | (.long number)
    - instruction :: opcode operand+
    - operand :: \$literal | reg | offset(reg) | (reg, reg, 4)
    - reg :: r0..7
    - literal :: number
    - offset :: number
    - number :: decimal | 0xhex

Name	Semantics	Assembly	Machine
load immediate	$r[d] \leftarrow v$	ld \$v, rd	0d-- vvvvvvvv
load base + offset	$r[d] \leftarrow m[r[s]+(o=4*p)]$	ld o(rs), rd	1psd
load indexed	$r[d] \leftarrow m[r[s]+4*r[i]]$	ld (rs, ri, 4), rd	2sid
store base + offset	$m[r[d]+(o=4*p)] \leftarrow r[s]$	st rs, o(rd)	3spd
store indexed	$m[r[d]+4*r[i]] \leftarrow r[s]$	st rs, (rd, ri, 4)	4sdi

- 4 addressing modes
  - **immediate**: constant value stored as part of instruction
  - **register**: operand is register number; register stores value
  - **base+offset**: operand is register number; register stores memory address of value (+ offset = p \* 4), p encodes 0..15, so offset ranges from 0 to 60.
  - **indexed**: two register-number operands; store base memory address and value of index

• Arithmetic			
Name	Semantics	Assembly	Machine
register move	$r[d] \leftarrow r[s]$	mov rs, rd	60sd
add	$r[d] \leftarrow r[d] + r[s]$	add rs, rd	61sd
and	$r[d] \leftarrow r[d] \& r[s]$	and rs, rd	62sd
inc	$r[d] \leftarrow r[d] + 1$	inc rd	63-d
inc address	$r[d] \leftarrow r[d] + 4$	inca rd	64-d
dec	$r[d] \leftarrow r[d] - 1$	dec rd	65-d
dec address	$r[d] \leftarrow r[d] - 4$	deca rd	66-d
not	$r[d] \leftarrow \sim r[s]$	not rd	67-d
• Logical (shifting), NOP, and halt			
Name	Semantics	Assembly	Machine
shift left	$r[d] \leftarrow r[d] \ll v=s$	shl \$v, rd	7dss (ss>0)
shift right	$r[d] \leftarrow r[d] \gg v=s$	shr \$v, rd	7dss (ss<0)
halt	halt machine	halt	F0--
nop	do nothing	nop	FF--

- 4 incrementing/decrementing is useful when calling functions
  - Shifting: distinguished by the sign of the last byte (2's complement)

• What instruction is represented by the bytes: 0x46 0x24?

- ld (r4, r4, 4), r6
- st r4, (r6, r2, 4)
- st r6, (r2, r4, 4)
- st r6, 2(r4)
- st r6, 8(r4)

Name	Semantics	Assembly	Machine
load immediate	$r[d] \leftarrow v$	ld \$v, rd	0d-- vvvvvvvv
load base + offset	$r[d] \leftarrow m[r[s]+(o=4*p)]$	ld o(rs), rd	1psd
load indexed	$r[d] \leftarrow m[r[s]+4*r[i]]$	ld (rs, ri, 4), rd	2sid
store base + offset	$m[r[d]+(o=4*p)] \leftarrow r[s]$	st rs, o(rd)	3spd
store indexed	$m[r[d]+4*r[i]] \leftarrow r[s]$	st rs, (rd, ri, 4)	4sdi

- Correct answer: C
  - Explanation: 0x4624 →  $m[r[d]+4*r[i]] \leftarrow r[s] \rightarrow st r6, (r2, r4, 4)$

- What instruction is represented by the bytes: `0x75 0xF8`?

- `shl $5, r8`
- `shl $8, r5`
- `shr $5, r8`
- `shr $8, r5`
- `shl $0xf8, r5`

Name	Semantics	Assembly	Machine
shift left	$r[d] \leftarrow r[d] \ll v=s$	<code>shl \$v, rd</code>	<code>7dss</code> ( $ss>0$ )
shift right	$r[d] \leftarrow r[d] \gg v=s$	<code>shr \$v, rd</code>	<code>7dss</code> ( $ss<0$ )

- Correct answer: D
- Explanation: `0x75f8`, f8 is negative  $\rightarrow r[d] \leftarrow r[d] \gg / \ll v=s \rightarrow r[5] \leftarrow r[5] \gg f8 \rightarrow shr \$8, r5$
- CPU implementation
  - Internal state
    - PC (program counter): address of **next** instruction
    - Instruction: the value of current instruction
  - Cycle stage
    - Fetch
    - Execute
  - Internal Registers
    - `insOp0, insOp1, insOp2, insOpImm, insOpExt, pc`
    - Read using `get()`
    - Change using `set(value)`
  - General purpose registers
    - `reg.set`
    - `reg.get`
  - Main Memory
    - `mem.readInteger`
    - `mem.writeInteger`
- Global dynamic arrays
  - Java
    - array variable stores reference to array allocated dynamically with **new** statement
  - C
    - array variables can store static arrays, or
    - pointers to arrays allocated dynamically with call to **malloc** library procedure
    - `// dynamic array`

```

int a;
int * b; // global variables still statically
          // allocated
          // pointer to integer (array)

void foo() {
    b = malloc(10 * sizeof(int)); // malloc
happens when program runs
    b[a] = a;
}

// static array
int a;
int b_data[10];
int * b = &b_data[0];

```

- C pointers
  - Use the term *pointer* instead of *reference* (in Java)
  - Declaration
    - the type is a pointer to the type of its elements, indicated with a \*
    - store internally as an address, regardless of element type
  - Type safety
    - any pointer can be typecast to any other pointer type
  - Bounds checking
    - C performs no array bounds checking
    - out-of-bounds access manipulates memory that is not part of the array
    - Performance is better, but a major source of vulnerability
- Dynamic allocation in C
  - done by **malloc**
  - Returns the **address of a contiguous block of bytes**, with **no associated type or initialization**
    - Element type and dereferencing determines how the binary data are interpreted
  - Dynamic allocation will be discussed in more details in future units
- Static vs. Dynamic allocation
  - For static arrays, the compiler allocates the array
  - For dynamic arrays, the compiler allocates a pointer

- The dynamic array is allocated by `malloc`,  
e.g. at address 0x3000
- Variable `b` is set to the memory address of  
this array

## Jan. 29th - Unit 1b: Cont.d and Unit 1c: Instance Variables and Dynamic Allocation

- Summary: Arrays in Java and C
  - Similarities
    - Array: list of items of the same type
    - Array elements are named by non-negative integers starting at 0
    - syntax for accessing element `i` of array `b` is `b[i]`
  - In Java
    - variable stores a reference to the array
  - In C
    - variable can store a pointer to the array, or the array itself
      - distinguished by variable type (e.g. pointer variable vs. array variable)
    - array element access via pointer can be done by
      - `b[i]`
      - `*(b + i)`
  - Declaring C pointers
    - Declaration
      - `TYPE * ptr;`
    - `int * x, * y;`
  - Pointers can be assigned the address of a variable of type `TYPE`
    - `int a = 5;`  
`int * ptr_a = &a;`
  - Access the value being pointed by `dereferencing` with `*`
    - `int a = 5;`  
`int * ptr_a = &a;`  
`*ptr_a = 37;`

- Pointer Arithmetic

- Purpose: an alternative way to access array elements compared to  $a[i]$
- Address are numbers
  - Can add / subtract an integer index to a pointer (address)
  - Results in the address of something else
  - Value of the pointer is offset by **index**  $\times$  **size** of the pointer's referent (+ 3, is equivalent to a pointer value 12 bytes larger than the original)
- Subtracting two pointers of the same type
  - Gives us the number of elements between two addresses
  - $\&a[7] - \&a[5] == 2 == (a + 7) - (a + 5)$

- ```
int * a;
a[7] = 5;

int * a;
*(a + 7) = 5;
```

- Which element of the original 10-element array is modified with the highlighted instructions?

- A. Element with index 3  
 B. Element with index 6  
 C. Element with index 0  
 D. No element is modified  
 E. More than one element is modified

```
int* c;
void foo() {
    c = malloc(10 * sizeof(int));
    c = &c[3];
    *c = *c[3];
}
```

|        |
|--------|
| 0x1000 |
| ...    |
| 0x2000 |
| 0x2004 |
| 0x2008 |
| 0x200c |
| 0x2010 |
| 0x2014 |
| 0x2018 |
| 0x201c |
| 0x2020 |
| 0x2024 |

- Correct answer: A (c's base address has been changed, 6th element of the original c is now  $c[3]$ )
- Explanation:

```
int * c;

void foo() {
    c = malloc(10 * sizeof(int)); // Dynamically allocate 10-element array
                                // c == 0x2000

    c = &c[3]; // c = 0x200c (0x200c becomes the new base address)
    *c = *&c[3]; // *c == c[3]
                  // *&c[3] == c[6]
}
```

- Summary

- C yes, Java no
  - Static arrays
  - Arrays can be accessed using pointer dereference operator
  - pointer arithmetic

- C no, Java yes
  - Type-safe dynamic allocation
  - Automatic array bounds checking

- Unit 1c - Instance Variables

- Variables that are an instance of a class or struct
  - Created dynamically (usually)
  - Many instances of the same class/struct can coexist
- Java vs C
  - Java: objects are instances of non-static variables of a class
  - C: structs are named variable groups, instance also called a struct
- Accessing an instance variable
  - requires a reference(pointer to a particular object/struct)
  - then variable name chooses a variable in that object/struct

- Structs in C

- A struct is a collection of variables of arbitrary type
  - allocated and accessed together
  - only variables, no functions
- Declaration
  - similar to declaring a Java class without methods
  - name is "struct" plus name provided by programmer
- The definition of a struct only describes layout of member variables in memory when an instance is created

- ```
struct D {
    int e;
    int f;
}

struct D d0; // Static declaration
struct D * d1; // Dynamic declaration

d0.e = d0.f; // Static access
d1->e = d1->f; // Dynamic access
```

- Static struct are allocated by the compiler

```

•      .pos 0x1000
d0.e:          .long 0xffffffff
•      .pos 0x1004
d0.f:          .long 0xffffffff

```

- Address of the first member is the same as address of the struct itself
- Dynamic struct are allocated runtime

- the variable that stores the struct pointer may be static or dynamic

- the struct itself is allocated when the program calls **malloc**

```

•      .pos 0x1000
d1:          .long 0x00001000

•      .pos 0x2000 (malloc)
d1->e:          .long 0xffffffff

```

- `struct D * d1;`

```

void foo() {
    d1 = malloc(sizeof(struct D));
}

```

- Struct Access

- Struct members can be accessed using **offset** from base address
  - Offset to each member from base of struct is also **static**
- As before, static and dynamic differ by an extra memory access

## Jan. 31st - Unit 1c: Cont.d

- Memory addressing modes

- Scalar

```
i = a;
```

- Address in register
- Access memory at address in register: ld (r1), r0
- Arrays

```
i = a[j]
```

- **base address** in register
- dynamic **index** in register
- access memory at base + index × element size: ld (r1, r2, 4), r0
- Struct members

```
i = a.j;
i = b -> k;
```

- Base address in register
- Static **offset**
- Access memory at base plus offset
- Equivalent to array access with static index (index can be used but requires extra register, index must be manually computed)

Given:

<code>struct X {     int i;     int j; }</code>	<code>struct X s; // s is a global variable</code>
	<code>ls \$s, r0 # r0: address of s</code>

Which of the following loads `s.j` into `r1`?

- |                             |  |
|-----------------------------|--|
| A. ld \$4, r1               | D. ld (r0), r0<br>add 4, r0<br>ld (r0), r1 |
| B. ld 4(r0), r1             | E. None of these                           |
| C. ld (r0), r1<br>add 4, r1 |  |

- Correct answer: B
- Offset approaches

Base+Offset vs Index			
array pointer with dynamic index		struct pointer	
<code>int* b;</code>	<code>load b[i] into r0</code>	<code>struct S {     int i;     int j; } *s;</code>	<code>load s-&gt;j into r0</code>
<ul style="list-style-type: none"> <li>• using <code>ld (rx, ry, 4), r0</code></li> </ul>	<ul style="list-style-type: none"> <li>• using <code>ld x(ry), r0</code></li> </ul>	<ul style="list-style-type: none"> <li>• using <code>ld (rx, ry, 4), r0</code></li> </ul>	<ul style="list-style-type: none"> <li>• using <code>ld x(ry), r0</code></li> </ul>
<code>ld \$b, r0      # r0: &amp;b ld (r0), r0    # r0: b (base address) ld \$i, r1      # r1: &amp;i ld (r1), r1    # r1: i <b>ld (r0, r1, 4), r0 # r0: b[i]</b></code>		<code>ld \$s, r0      # r0: &amp;s ld (r0), r0    # r0: s (struct address) <b>ld 4(r0), r0 # r0: s-&gt;j</b></code>	<code>ld (rx, ry, 4), r0</code>
<ul style="list-style-type: none"> <li>• using <code>ld x(ry), r0</code></li> </ul>			

- Array pointer with dynamic index: **cannot** use "ld x(ry), r0", because "x" here is static, but we need it to be dynamic
- Struct pointer: **can** use "ld (rx, ry, 4), r0", because the offset required is static, we can use a dynamic register to offset it. "ry" would contain the "index" of "j", which is 1.

From previous slide:

array pointer with dynamic index

```
int* b;  
int i;
```

load b[i] into r0

- using ld (rx, ry, 4), r0

```
ld $b, r0          # r0: &b  
ld (r0), r0        # r0: b (base address)  
ld $i, r1          # r1: &i  
ld (r1), r1        # r1: i  
ld (r0, r1, 4), r0 # r0: b[i]
```

Suppose we have run the first 4 assembly statements. How else can we load b[i] into r0?

- |                                 |   |
|---------------------------------|---|
| A. ld i(r0), r0                 | D. shl \$2, r1<br>add r0, r1<br>ld (r1), r0 |
| B. ld r1(r0), r0                |   |
| C. shl \$2, r1<br>ld r1(r0), r0 | E. add r0, r1<br>ld (r1), r0                |

- Correct answer: D
- Explanation: r1 becomes  $4 * i$ ; r1 becomes  $r0 + 4 * i$ ; r1 is now the address of b[i]
- Variations in struct declaration
  - Struct variables can be declared inside other structs

```
• struct D {  
    int e;  
    int f;  
}  
  
struct X {  
    int i;  
    struct D d;  
    int j;  
}  
  
struct Y {  
    int i;  
    struct D * d;  
    int j;  
}  
  
struct Z {  
    int i;  
    struct Z * z;  
    int j;  
}
```

- struct Z needs to be terminated somehow
- struct members can be arrays or pointers

- ```

struct W {
    int i;
    int a[10];
    int j;
}

struct U {
    int i;
    int * a;
    int j;
}

```

- What is the offset (in bytes) of member **j** in **struct X** below?

|                                                                  |                                                                                         |
|------------------------------------------------------------------|-----------------------------------------------------------------------------------------|
| <code> <b>struct D {</b>     int e;     int f; <b>}</b>; </code> | <code> <b>struct X {</b>     int i;     <b>struct D d;</b>     int j; <b>}</b>; </code> |
|------------------------------------------------------------------|-----------------------------------------------------------------------------------------|

- A. 0
- B. 4
- C. 8
- D. 12
- E. something else

- Correct answer: D
- Struct size and alignment

- Alignment rules apply inside of a struct
- Each instance variables will be aligned within the struct according to its type size
- structs are usually aligned according to their largest instance variable type
- structs can mix data types (**padding** may be needed to align the members)

- ```

// needs 8 bytes
struct X {
    char a; // e.g. 0x1000
    char b; // 0x1001
    int i; // 0x1004
}

// needs 12 bytes
struct Y {
    char a; // e.g. 0x1000
    int i; // 0x1004
    char b; // 0x1008
}

```

- Multiple of the largest element
- What is the offset (in bytes) of member **j** in **struct X** below?  
*Assume 8-byte addresses.*

```
struct D {
    int e;
    int f;
};
```

```
struct X {
    int i;
    struct D* d;
    int j;
};
```

- A. 0  
 B. 8  
 C. 12  
 D. 16  
 E. something else

- Correct answer: D

## Feb. 2nd - Unit 1c: Cont.d

- Arrays and structs
  - Array in struct

```
struct S {
    int i;
    int j[10];
}
```

- Array of structs

```
struct S a[10];
```

- Array of pointers to structs

```
struct S * b[10];
```

- What is the offset (in bytes) of member **j** in **struct X** below?

```
struct D {
    int e;
    int f;
};
```

```
struct X {
    int i;
    struct D d[10];
    int j;
};
```

- Correct Answer: 84
- Explanation: Assume 4-byte integer, and 4-byte address;  $d[10]$  can be shifted to eliminate padding, because each data in D is still 4-byte. Thus, it is  $4 \times 1 + 8 \times 10 = 84$  bytes.

- What is the offset (in bytes) of member `j` in `struct X` below?

```
struct Y {
    char c;
    short s;
};
```

```
struct X {
    int i;
    struct Y y[10];
    int j;
};
```

- Correct answer: 44
- Explanation: `char` (1 byte), `short` (2 bytes), `int` (4 bytes); each `Y` will occupy 4 bytes with padding after `short`. Thus, the offset is  $4 \times 1 + 10 \times (1 + 1 + 2) = 44$

- Given

```
struct A {
    char a;
    int* b;
    int* p;
};
```

```
struct X {
    char c;
    struct A a;
    int j;
};
```

```
struct X s; // a global variable
```

- How many memory reads to load `s.a.b[2]` into `r1`?
  - A. 1
  - B. 2
  - C. 3
  - D. 4
  - E. Not enough information

- Correct answer: B
- Explanation:

```
ld $s, r0          # r0 <- s
ld 8(r0), r2       # r2 <- m[r0 + 8]
ld 8(r2), r1       # r1 <- m[r2 + 8]
```

- Dynamic allocation in C and Java

- Programs can allocate memory dynamically
  - Reserves a range of memory
- In Java, instances of classes are allocated by `new` keyword
- In C, byte ranges are allocated by call to `malloc` procedure
  - Bytes can be used for *any type* that can fit in them

- C

- Memory allocation

```
void * malloc(unsigned long n)
```

- `n` is the number of bytes to allocate
- return type is `void *`: just an address
  - A pointer to anything (no specific type)

- Can be cast to/from any other pointer type
- Cannot be dereferenced directly
- **sizeof** to determine the number of bytes to allocate
  - **sizeof(x)** statically computes the number of bytes in a type or variable
  - **sizeof(ptr)** gives the size of a *pointer*, not what it *points* to
  - ```
struct Foo * f = malloc(sizeof(struct Foo));
struct Foo * f = malloc(sizeof(*f))
```

- The system marks the chunk of memory as allocated: the allocation belongs to the program for the lifetime; or until explicitly deallocated
- Deallocation
  - In Java
    - Automatic garbage collection: requires keeping track of every reference to an object
  - In C
    - Dynamic memory must be deallocated explicitly by calling **free** (call on an address which was returned by **malloc** previously)
    - **free** deallocates memory immediately; does not check to see if it is still in use
  - **Memory Leak**
    - Dynamically allocated data is not deallocated
  - Memory heap
    - The **heap** is a large section of memory from which **malloc** allocates objects
    - In Java, all objects are stored on the heap
  - Issues with explicit deallocation
    - **free(x)** does not do
      - change the value of **x**, the pointer still points to the freed object
      - other variables may still point there too
      - the binary data stored at address **x** is not erased

```

• struct MBuf * create() {
    struct MBuf * mBuf =
malloc(sizeof(*mBuf));
    //
    return mBuf;
}

void destroy(struct MBuf * buf) {
    //
    free(buf);
}

```

- What if access buf after destroying?  
Dangling pointer
- What if call destroy twice on **buf**?  
Double-free
- What if forget to call destroy?  
Memory leak

## Feb. 5th - Unit 1c: Cont.d

- **malloc** and **free**
  - Organization of memory
    - statically allocated: code, static data
    - the rest is unused or dynamically allocated
  - **malloc/free**
    - implemented in the library **<stdlib.h>**
    - manage a large chunk of memory called the heap
    - keep track of what's allocated or free
  - **malloc**
    - find a free chunk of memory of requested size
    - mark it as allocated
    - return a pointer to it (keeping track of the size of allocation) - address of the first byte
  - **free**
    - use pointer to determine allocated size
    - mark referenced memory as free
    - can only be called on an address associated with an allocation
- Dangling pointers

- A pointer to an object that has been freed / could point to unallocated memory / to another object (a later `malloc` re-used the address)

- Often caused by use-after-free
- Can happen with incorrect pointer manipulation

- ```

struct D {
    int e;
    int f;
};

struct D * d; // global

void foo() {
    d = malloc(sizeof(struct D));
}

// execution
void someFunction() {
    foo();
    d -> e = 5;
    d -> f = 12;
    // some other stuff
    free(d);
    // some other stuff
    short * s = malloc(8 * sizeof(short));
    // s happen to be

    // address in d
    s[0] = -13;
    s[1] = 2;
    s[2] = 5;

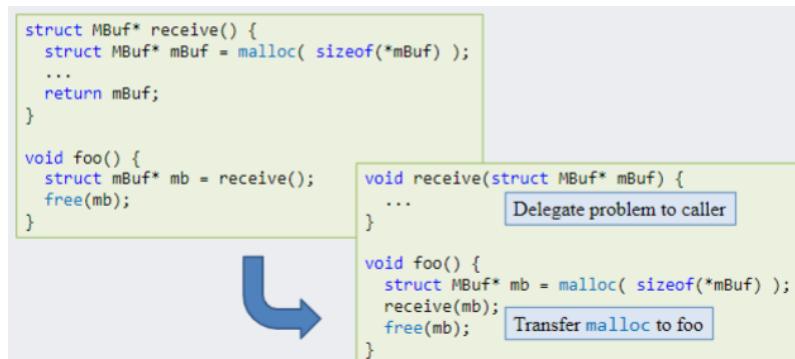
    printf("%d\n", d -> e); // grabs s[0] and
    s[1] together in binary
    // as e
}

```

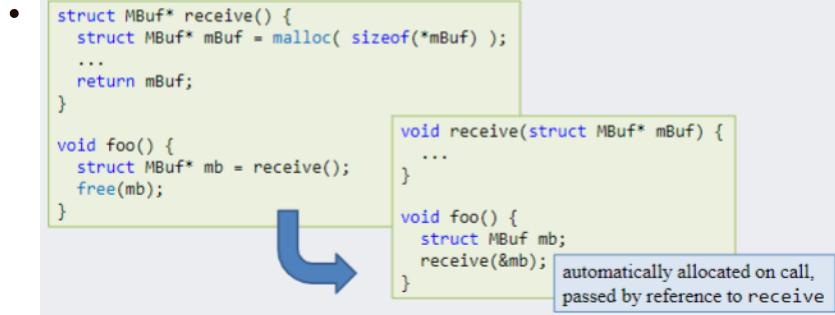
- Why a problem?
  - program thinks it is writing to object of type X, but it isn't actually, it may be writing to an object of type Y
- Memory leak
  - A dynamic memory object with no pointers pointing to it
    - Usually happens if a program doesn't free object properly
    - May also happen if a pointer to a valid object is changed to another value

- ```
//  
d = malloc(sizeof(struct D));  
// ...  
d = malloc(sizeof(struct D));
```

- Why problem?
  - If contains useful information, it can no longer be accessed
  - If object is large, or many memory leaks happen at the same time, the program uses up too much memory
- Avoiding memory problems in C
  - Avoid the problem cases, **if possible**
    - restrict dynamic allocation / free to single procedure
    - don't write procedures that return pointers
    - use local variables instead - automatically allocated on call and freed on return
  - Engineer for memory management
    - Define rules for which procedure is responsible for deallocation
    - implement explicit **reference counting** if multiple potential deallocator
    - define rules for which pointers can be stored in data structures
    - use coding conventions and documentation to ensure rules are followed
  - Co-locate allocation and deallocation
    - allocate the object in its caller, and pass pointer to it - to the procedure (callee)
    - good if caller itself does both **malloc/free** itself



- as long as **receive()** does not pass along the reference
- Use local variables instead
  - If a procedure does both **malloc/free**
    - Use a "static" local variable instead of **malloc**, and don't give another procedure a pointer to the local variable
  - Local variables are allocated on call and deallocated on return



- C strings

- A string "like this" in C
  - It is an array of `chars`
  - Has its end indicated by the first null '\0' in the array
  - Every string has a maximum length (capacity of array -1), and a printable length (determined by the position of the first null character)
  - Each character resides at some address (`char *`)
- Standard C library <string.h> has operations on strings
  - `strlen(s)` returns the printable length of a string, starting at address `s`
- How can we create a new string that copies the existing string?

|                                                                                                                                                                                                                                        |                                                                                                                                                                                |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"> <li>•           <pre> char* copy(char* s) {     int len = strlen(s);     char* d = malloc(len + 1);      for (int i = 0; i &lt;= len; i++)         d[i] = s[i];      return d; } </pre> </li> </ul> | <pre> // in your application program void foo(char* s) {     char* d = copy(s);     printf("%s\n", d);      free(d); }  void bar() {     foo("Hello, World!"); } </pre>        |
| <b>C library equivalent: strdup</b>                                                                                                                                                                                                    |                                                                                                                                                                                |
| <pre> void copy(char* d, char* s) {     int len = strlen(s);      for (int i = 0; i &lt;= len; i++)         d[i] = s[i]; } </pre>                                                                                                      | <pre> // in your application program void foo(char* s) {     int len = strlen(s);     char* d = malloc(len + 1);     copy(d, s);     printf("%s\n", d);      free(d); } </pre> |
| <b>C library equivalent: strcpy</b>                                                                                                                                                                                                    |                                                                                                                                                                                |

- - Pro: `malloc` and `free` are in the same procedure.
  - Con: Need to make sure we `malloc` sufficient space, which can potentially lead to out-of-bound if `d` is too short

```

void copy(char* d, int dsize, char* s) {
    int len = strlen(s);

    if (len > dsize-1) len = dsize-1;

    for (int i = 0; i < len; i++)
        d[i] = s[i];

    d[len] = '\0';
}

// in your application program
void foo(char* s) {
    int len = strlen(s);
    char* d = malloc(len + 1);
    copy(d, len+1, s);
    printf("%s\n", d);

    free(d);
}

```

C library equivalent: `strncpy`

- `len` will be set to minimum between two strings
- `strncpy` copies minimum of  $\{n, \text{len}(\text{source\_string})\}$
- Detecting memory problems
  - Valgrind
    - performs dynamic analysis of the runtime of a program
    - Monitors dynamic allocation and deallocation
    - Memory leaks, use after free (dangling pointers) etc
  - "... if possible..."
  - A reference may needs to be passed among multiple modules
  - Correctly `free` an object
    - `free` should only happen if the object is still in use
    - need to keep track of all its users
      - start and stop

```

struct MBuf* receive() {
    struct MBuf* mBuf = malloc(sizeof(*mBuf));
    ...
    return mBuf;
}

void foo() {
    struct MBuf* mb = receive();
    bar(mb);
    free(mb);
}

struct MBuf* aMB;

void bar(struct MBuf* mb) {
    aMB = mb;
}

void bat() {
    aMB->x = 0;
}

```

• what might happen in `bar()`, and  
• why a subsequent call to `bat()` would expose a serious bug?

## Feb. 7th - Unit 1c: Cont.d

- Detecting memory problems

- ```

void foo() {
    struct MBuf * mb = receive();
    bar(mb);
    free(mb);           // How do we know we can do
this?                  // Need information about
                        whether the pointer is used
}

void bar(struct MBuf * mb) {
    if (smth_complicated_private) {
        aMB = mb;           // may or may not occur
    }
}

```

- Reference counting

- Initialize the reference count to 1 (the caller has a reference)
- any procedure that stores a reference (starts using the object) increments the count
- any procedure that discards a reference (stops using the object) decrements the count
- never call **free** directly, only **freed** when the reference count goes to 0

<pre> struct MBuf* receive() {     struct MBuf* mBuf = rc_malloc(sizeof(*mBuf));     ...     return mBuf; } </pre>	<pre> struct MBuf* aMB = 0;  void foo() {     struct MBuf* mb = receive();     bar(mb);     rc_free_ref(mb); } </pre>	<pre> void bar(struct MBuf* mb) {     // free existing reference in aMB (why?)     if (aMB != 0)         rc_free_ref(aMB);     rc_keep_ref(mb);     aMB = mb; } </pre>
--	---	--

- replace **malloc** with a reference-count **rc\_malloc**; replace **free** with a reference-count **rc\_free\_ref**
- Reference counter can be part of a **struct**
  - updates to reference counter are done through special methods
  - no longer explicitly call **malloc/free** directly when creating the struct instances

<pre> struct buffer* rc_malloc() {     struct buffer* buf = malloc(sizeof *buf);     buf-&gt;ref_count = 1;     return buf; }  void rc_keep_ref(struct buffer* buf) {     buf-&gt;ref_count++; }  void rc_free_ref(struct buffer* buf) {     buf-&gt;ref_count--;     if (buf-&gt;ref_count == 0)         free(buf); } </pre>	<pre> struct buffer {     ...     // the actual data attributes     ...     int ref_count; } </pre>
---	---

- Problem: every data type requires their own struct definition, and their reference counting procedures.
- Implementing reference counting

- `refcount.h`

```
void * rc_malloc(int nbytes); //  
alternative malloc with room for  
// ref_count  
void rc_keep_ref(void * p); // call when  
reference added  
void rc_free_ref(void * p); // call when  
reference removed
```

- `refcount.c`

```
void * rc_malloc(int nbytes) {  
    // Assume the largest atomic data type is  
    8 bytes  
    int * ref_count = malloc(nbytes + 8);  
    *ref_count = 1;  
    return ((void *) ref_count) + 8;  
}  
  
void rc_keep_ref(void * p) {  
    int * ref_count = p - 8;  
    (*ref_count) += 1;  
}  
  
void rc_free_ref(void * p) {  
    int * ref_count = p - 8;  
    (*ref_count) -= 1;  
    if (*ref_count == 0) {  
        free(ref_count);  
    }  
}
```

- Reference counting and procedure calls

- When reference is a parameter
  - the caller has a reference and maintains it for the duration of the call
  - the callee need not perform a `keep_ref` unless
    - callee saves the pointer someplace that will outlast the call (global variable)
    - callee returns a pointer to the caller
- When reference is a return value
  - the callee must have a reference to the value

- it passes that reference to the caller
  - the callee implicitly gives up its reference as part of the return
  - the reference is transferred to the caller
- the caller must call **free\_ref** when it no longer stores the reference
- When reference is stored in a local variable
  - that variable goes away implicitly when the procedure returns
  - the procedure must call **free\_ref** before it returns

## Feb. 9th - Unit 1c: Cont.d + Unit 1d: Static Control Flow

- What if there is more to the quiz structure?

```
struct quiz* quiz_create(char* name) {
    struct quiz* quiz = rc_malloc(sizeof(*quiz));
    int name_size = strlen(name) + 1;
    quiz->name = rc_malloc(name_size);
    strcpy(quiz->name, name, name_size);
    quiz->grades = NULL;
    return quiz;
}

void addGrades(struct quiz* quiz, int* grades, int numGrades) {
    quiz->grades = rc_malloc(sizeof(*quiz->grades) * numGrades);
    for (int i = 0; i < numGrades; i++)
        quiz->grades[i] = grades[i];
}

void quiz_delete(struct quiz* quiz) {
    rc_free_ref(quiz->name);
    rc_free_ref(quiz);
}
```

```
struct quiz {
    char* name;
    int* grades;
    int num_grades;
};
```

How should we fix  
quiz\_delete?

- - ```
void quiz_delete(struct quiz * quiz) {
            rc_free_ref(quiz -> name);
            if (quiz -> grades != NULL) {
                rc_free_ref(quiz -> grades);
            }
            rc_free_ref(quiz);
        }
```

- Garbage Collection (Java)
  - In Java, objects are deallocated implicitly
    - Program never calls **free**
    - the runtime system tracks every object reference
    - garbage collector runs periodically to deallocate unreachable objects
  - Advantage compared to explicit **free** in C
    - No dangling pointers

- (almost) No memory leaks
- Reference cycles are not a problem
  - Object a has pointer to b; Object b has pointer to a
  - Cause memory leak when using reference counting alone
- Discussion
  - Advantages of explicit `free`
    - fast, efficient, forces us to be careful about memory management
  - Advantages of reference counting
    - Reduce coupling between modules
  - Advantages of garbage collection
    - convenient, automatic, works well for most cases
  - Should we ignore deallocation in Java
    - No, `HashMap`
- Memory management in Java
  - Memory "leaks" will occur
    - Garbage collector fails to reclaim **unneeded** objects
    - Still a significant problem for long-running programs where garbage accumulates
  - Why would an object not be reclaimed
    - Garbage collector only reclaims **unreachable** objects
    - Unneeded / unused objects that keep references are not reclaimed
    - Collections and maps may maintain old unneeded objects.

## Static Control Flow

- Control flow
  - Definition: sequence of instruction executions performed by a program
  - Every program execution can be described by such a linear sequence
- Loops

- ```

public class Foo {
    static int s = 0;
    static int i;
    static int a[] = new int[] {2, 4, 6, 8, 10, 12,
    14, 16, 18, 20};

    static void foo() {
        for (int i = 0; i < 10; i += 1) {
            s += a[i];
        }
    }
}

```

- ```

int s = 0;
int i;
int a[] = {2, 4, 6, 8, 10, 12, 14, 16, 18, 20};

void foo() {
    for (i = 0; i < sizeof(a); i++) {
        s += a[i];
    }
}

```

- ```

int s = 0;
int * p;
int a[] = {2, 4, 6, 8, 10, 12, 14, 16, 18, 20}

void foo() {
    p = a;
    while (p < a + 10) {
        s += *p++
    }
}

```

- `s += *a++` is not valid because `a` is a static array (base address), cannot change
  - even if we could, the loop condition needs to be updated
- Implementing loops in SM213

- ```

int s = 0;
int i;
int a[] = {2, 4, 6, 8, 10, 12, 14, 16, 18, 20};

void foo() {
    for (i = 0; i < sizeof(a); i++) {
        s += a[i];
    }
}

```

- This loop can be unrolled.

- Dissecting loops (using `goto`)

- ```

for (i = 0; i < n; i += 1) {
    s += a[i];
}

```

- ```

i = 0
loop:      goto end_loop if not (i < n)
            s += a[i]
            i++
            goto loop
end_loop: // whatever code comes after

```

- Control flow in the machine

- Program counter: contains address of **next instruction** to execute
- For sequential instructions, PC is updated in the fetch stage
  - **incremented** by 2 or 6 bytes depending on instruction size
  - To break sequential execution, we need to change PC from what it would contain (`goto`)

- Control flow: ISA extensions

- New instructions: jumps and conditional jumps
  - `pc <- <address>`
  - `pc <- <address> if <condition>`
- Options for evaluating
  - **Unconditional**
  - Conditional: common in RISC
    - `goto <address> if <register> <condition> 0`

- Conditional: based on result of last executed ALU instruction
- **j address:** jump to **address**
- **br address:** branch to **address**
- **beq r0, address:** branch to **address** if register equal to 0
- **bgt r0, address:** branch to **address** if register strictly greater than 0
- Jump instruction size
  - Memory addresses are big (32 bits in SM 213)
  - And control flow instructions are common
  - Observations
    - jumps *usually* move a short distance (forward or backward)
- PC-relative addressing
  - instead of specifying the destination address completely, specify the **offset** from the current location in code
  - use the current value of PC as **base address**
  - Offset must be a signed number
  - Assembly still specifies the actual address/label
  - jumps that use PC-relative addressing are called **branches**
  - Example
    - Suppose we want to do: **0x1000: goto 1008**
    - Option 1: absolute addressing (**X--- 00001008**), 6 bytes
    - Option 2: PC-relative addressing (branch), PC updated to **0x1002** already
      - **Y-06**, the offset is 6 from 1002 to 1008
      - The address usually offsets by 2 bytes or 6 bytes (even), we can expand our range by encoding offsets /2. The **Y-06** is in fact **Y-03**

## Feb. 12th - Unit 1d: Cont.d

- ISA for static control flow
  - At least one absolute jump
  - At least one PC-relative jump
    - specify relative distance using real distance / 2
    - PC-relative offset is signed
  - Some conditional jumps

| Name              | Semantics                                          | Assembly  | Machine        |
|-------------------|----------------------------------------------------|-----------|----------------|
| branch            | $pc \leftarrow a$ (or $pc + p^*2$ )                | br a      | 8-pp           |
| branch if equal   | $pc \leftarrow a$ (or $pc + p^*2$ ) if $r[c] == 0$ | beq rc, a | 9cpp           |
| branch if greater | $pc \leftarrow a$ (or $pc + p^*2$ ) if $r[c] > 0$  | bgt rc, a | a_cpp          |
| jump              | $pc \leftarrow a$                                  | j a       | b--- aaaaaaaaa |

- Convert the bgt instruction to machine code:

```
.pos 0x10
bgt r0, L1
.pos 0x20
L1: halt
```

- A. 0xa0 0x10  
 B. 0xa0 0x08  
 C. 0xa0 0x0e  
 D. 0xa0 0x07  
 E. 0xa0 0x20

| Name              | Semantics                                         | Assembly  | Machine |
|-------------------|---------------------------------------------------|-----------|---------|
| branch if greater | $pc \leftarrow a$ (or $pc + p^*2$ ) if $r[c] > 0$ | bgt rc, a | a_cpp   |

- Correct answer: D
- Explanation: **bgt r0, L1** is a 2-byte instruction, so after fetching, PC has value 0x12. 0x20 is 0x0e away from 0x12, so half the real distance is 0x07. This means **pp** in machine code corresponds to **0x07**

- Convert the br instruction to machine code.

- note: each instruction here is 2 bytes

```
loop: mov r0, r5
      add r4, r5
      beq r5, end_loop
      inc r0
      br loop
```

- A. 0x80 0xf5  
 B. 0x80 0xf8  
 C. 0x80 0xfc  
 D. 0x80 0xfb  
 E. 0x80 0xf6

| Name   | Semantics                           | Assembly | Machine |
|--------|-------------------------------------|----------|---------|
| branch | $pc \leftarrow a$ (or $pc + p^*2$ ) | br a     | 8-pp    |

- Correct answer: D
- Explanation: After we fetch **br loop**, PC holds the address right after the instruction, which is 10 bytes away from **loop**. Half the distance is 5, but since we are moving backwards, it should be **-5**, which corresponds to **0xfb**.

- What does the following instruction do?

0xa0 0x00

- A. infinite loop  
 B. sets PC to zero  
 C. sets PC to beginning of program  
 D. nothing  
 E. something else

| Name              | Semantics                                          | Assembly  | Machine        |
|-------------------|----------------------------------------------------|-----------|----------------|
| branch            | $pc \leftarrow a$ (or $pc + p^*2$ )                | br a      | 8-pp           |
| branch if equal   | $pc \leftarrow a$ (or $pc + p^*2$ ) if $r[c] == 0$ | beq rc, a | 9cpp           |
| branch if greater | $pc \leftarrow a$ (or $pc + p^*2$ ) if $r[c] > 0$  | bgt rc, a | a_cpp          |
| jump              | $pc \leftarrow a$                                  | j a       | b--- aaaaaaaaa |

- Correct answer: D

- Explanation: PC value is updated after fetch, but the offset is **0x00**, so the machine just continues running to the next instruction. This is equivalent to no operation.

- What is the value of **r0** after this code executes?

```

ld $1, r0
ld $4, r1
loop: beq r1, end
      shl $1, r0
      dec r1
      br loop
end: halt
    
```

- A. It never terminates

- B. 4

- C. 8

- D. 16

- E. 32

| Name              | Semantics                                         | Assembly  | Machine        |
|-------------------|---------------------------------------------------|-----------|----------------|
| branch            | $pc \leftarrow a$ (or $pc + p^2$ )                | br a      | 8-pp           |
| branch if equal   | $pc \leftarrow a$ (or $pc + p^2$ ) if $r[c] == 0$ | beq rc, a | 9cpp           |
| branch if greater | $pc \leftarrow a$ (or $pc + p^2$ ) if $r[c] > 0$  | bgt rc, a | acpp           |
| jump              | $pc \leftarrow a$                                 | j a       | b--- aaaaaaaaa |

- Correct answer: D

- ```

int a = 1;
int i = 4;
while (i != 0) {
    a *= 2;
    i -= 1;
}
    
```

- Implementing **for** loops

- ```

for (i = 0; i < 10; i += 1) {
    s += a[i];
}
    
```

- General form:

```

for (<init>; <continue-condition>; <step>) {
    <statement-block>
}
    
```

- Each of init, continue, and step can be compound expression

- pseudo-code template

|                                                                                                                                                                                |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>&lt;init&gt;</b><br><b>loop:</b> goto <b>end_loop</b> if not <b>&lt;continue-condition&gt;</b><br><b>&lt;statement-block&gt;</b><br><b>&lt;step&gt;</b><br>goto <b>loop</b> |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

- By the template:

```

for (i=0; i<10; i++)
    s += a[i];
    
```



|                                                                                                        |
|--------------------------------------------------------------------------------------------------------|
| <b>loop:</b> goto <b>end_loop</b> if not <b>(i &lt; 10)</b> ??<br>s += a[i]<br>i++<br>goto <b>loop</b> |
|--------------------------------------------------------------------------------------------------------|

- ISA does not support comparison to 10, but we can compare to 0, and there is no need to store **i** (or **s**) in memory for each loop iteration

- `i' == 0`

| Assembly: assume that all variables are global |           |
|------------------------------------------------|-----------|
| ld \$a, r1 # r1 = a = &a[0]                    |           |
| ld \$s, r2 # r2 = &s                           |           |
| ld (r2), r2 # r2 = s = s'                      |           |
| ld \$0x0, r0 # r0 = i' = 0                     |           |
| ld \$-10, r4 # r4 = -10                        | init      |
| loop: mov r0, r5 # r5 = t' = i'                |           |
| add r4, r5 # r5 = i' - 10                      | cont/exit |
| beq r5, end_loop # if i' = 10, goto +8         |           |
| ld (r1, r0, 4), r3 # r3 = a'[i']               |           |
| add r3, r2 # s' += a'[i']                      |           |
| inc r0 # i'++                                  |           |
| br loop # goto -14                             | step      |
| end_loop: ld \$s, r1 # r1 = &s                 |           |
| st r2, (r1) # s = s'                           |           |
| st r0, 4(r1) # i = i'                          |           |

- Implementing conditionals

- General form:

```

if <condition> {
    <then-statements>
} else {
    <else-statements>
}

```

- 

|                                            |  |
|--------------------------------------------|--|
| <code>c' = not &lt;condition&gt;</code>    |  |
| <code>else: &lt;else-statements&gt;</code> |  |
| <code>        goto end_if</code>           |  |
| <code>then: &lt;then-statements&gt;</code> |  |
| <code>end_if:</code>                       |  |

|                                            |  |
|--------------------------------------------|--|
| <code>c' = &lt;condition&gt;</code>        |  |
| <code>else: &lt;else-statements&gt;</code> |  |
| <code>        goto end_if</code>           |  |
| <code>then: &lt;then-statements&gt;</code> |  |
| <code>end_if:</code>                       |  |

or

- 

| Assembly:                         |  |
|-----------------------------------|--|
| ld \$a, r0 # r0 = &a              |  |
| ld (r0), r0 # r0 = a              |  |
| ld \$b, r1 # r1 = &b              |  |
| ld (r1), r1 # r1 = b              |  |
| mov r1, r2 # r2 = b               |  |
| not r2 # c' = !b                  |  |
| inc r2 # c' = -b                  |  |
| add r0, r2 # c' = a - b           |  |
| bgt r2, then # if (a>b) goto then |  |
| else: mov r1, r3 # max' = b       |  |
| br end_if # goto end_if           |  |
| then: mov r0, r3 # max' = a       |  |
| end_if: ld \$max, r0 # r0 = &max  |  |
| st r3, (r0) # max = max'          |  |

- Which of the choices corresponds to this C code:

```
Common setup:
ld $a, r0
ld (r0), r0 # r0: a
ld $b, r1
ld (r1), r1 # r1: b
ld $c, r2 # r2: &c
ld $0, r3 # r3: 0
ld $1, r4 # r4: 1
```

```
if (a == 0 && b > 0)
    c = 1;
else
    c = 0;
```

A. `beq r0, L0  
bgt r1, L0  
st r3, (r2)  
L0: st r4, (r2)`

B. `beq r0, L0  
bgt r1, L0  
st r4, (r2)  
br L1  
L0: st r3, (r2)  
L1:`

C. `beq r0, L0  
br L1  
L0: bgt r1, L2  
L1: st r3, (r2)  
br L3  
L2: st r4, (r2)  
L3:`

D. `beq r0, L0  
br L1  
L0: bgt r1, L2  
L1: st r4, (r2)  
br L3  
L2: st r3, (r2)  
L3:`

E. `beq r0, L0  
bgt r1, L0  
st r3, (r2)  
br L1  
L0: st r4, (r2)  
L1:`

- Correct answer: C

## Feb. 14th - Unit 1d: Cont.d, Unit 1e: Procedure class and the Stack

- What does this code do?

- 

```
.pos 0x1000
ld $0, r0
ld $0, r1
ld $1, r2
ld $j, r3
ld (r3), r3
ld $a, r4
L0: beq r3, L9
    ld (r4, r0, 4), r5
    and r2, r5
    beq r5, L1
    inc r1
L1: inc r0
    dec r3
    br L0
L9: ld $o, r0
    st r1, (r0)
    halt
```

- 

```
.pos 0x2000
j: .long 2
a: .long 1
        .long 2
o: .long 0
```

- Step 1: comment each line
- Step 2: Refine the comments to look more like C

- ```
int i = 0;
int j = 2;
int a[2] = {1, 2};
int o;
```

- Step 3: Look for basic blocks by examining branches
- Step 4: Associate control structure with C

- ```
for (j' = j, i' = 0; j' != 0; j' -= 1, i' += 1) {
    if (a'[i']) & 1) {
        o' += 1;
    }
}
```

- Step 5: Deal with what's left

- ```
for (i = 0; i != j; i += 1) {
    if (a[i] & 1) {
        o += 1;
    }
}
```

- Static procedure calls

- Java: method; C: procedure: subroutine with a name, arguments, and local scope
- Control flow during procedure calls

- ```
void foo() {
    ping();
}

void ping() {
    // smth
}
```

- Caller ↔ Callee
- Caller
  - goto ping: `j ping`
  - continue executing from callee's return
- Callee
  - from caller: then do whatever `ping()` does
  - goto `foo()` just after call to `ping()`, but return to where?
- The jump from `ping()` is dynamic
- Implementing procedure return
  - Return address**
    - the address to where the procedure jumps when it completes
    - the address of **the instruction following the call** that caused it to run
    - a dynamic property of the program (a procedure may be called from different locations within a program)
  - Questions

- How does the procedure know the return address?
- How does it jump to a dynamic address?
- How are recursive calls handled?
- Only the caller can provide the address
  - the caller must save it **before** it makes the call
  - By SM213 convention, caller will save the return address in **r[6]**
  - We need a new instruction to read the PC contents: **gpc**
- Jumping back to the return address
  - Callee assumes caller saved address in **r[6]**
  - We need a new instruction to jump to dynamic addresses stored in a register
- ISA instructions for static control flow
  - New requirements
    - Read value of the PC
    - jump to dynamically determined target address
  - New control flow instructions

| Name              | Semantics                                                       | Assembly    | Machine        |
|-------------------|-----------------------------------------------------------------|-------------|----------------|
| branch            | $pc \leftarrow a \text{ (or } pc + p^*2)$                       | br a        | 8-pp           |
| branch if equal   | $pc \leftarrow a \text{ (or } pc + p^*2) \text{ if } r[c] == 0$ | beq rc, a   | 9cpp           |
| branch if greater | $pc \leftarrow a \text{ (or } pc + p^*2) \text{ if } r[c] > 0$  | bgt rc, a   | a cpp          |
| jump              | $pc \leftarrow a$                                               | j a         | b--- aaaaaaaaa |
| get pc            | $r[d] \leftarrow pc + o \text{ (or } pc + p^*2)$                | gpc \$o, rd | 6fpd           |
| indirect jump     | $pc \leftarrow r[t] + o \text{ (or } r[t] + p^*2)$              | j o(rt)     | c tpp          |

- do **gpc** before we do the jump

```

void foo() {
    ping();
}

foo: gpc $6, r6  # r6 = address of next instruction after jump
      j ping    # goto ping()
      ...

void ping() {}

ping: j (r6)  # return to wherever r6 tells us to go (saved previously)

```

- **j ping** is 6 bytes, so the instruction after **ping()** must be offset by 6 for the PC

- What is wrong with this code?

```

main: gpc $6, r6      # r6 = pc + 6
      j ping        # ping()
      ld $5, r0      # r0 = 5
      ld $x, r1      # r1 = &x
      st r0, (r1)    # x = 5
      halt

ping: ld $10, r2      # r2 = 10
      gpc $6, r6      # r6 = pc + 6
      j pong        # pong()
      j (r6)         # return

pong: ld $20, r3      # r3 = 20
      j (r6)         # return

.pos 0x1000
x:   .long 0          # x
i:   .long 0          # i

```

- Since **r6** is updated in **ping**, after running **pong**, the program will just continue to execute **ping: j (r6)**

## Feb. 16th - Unit 1e: Cont.d

```

void b( int a0, int a1 ) {
    int 10 = 0;
    int 11 = 1;
    ...
    ...
}

```

Can **10**, **11**, **a0**, **a1** be allocated statically?

- Yes, always
- Yes, but only if **b** doesn't call itself directly
- Yes, but only if **b** doesn't call any functions
- No, one of these can be allocated statically at all

- Correct answer: C
- Examples (Local variables and recursion)

```

void b(int a0) {
    int 10 = a0;
    if (a0 > 0) {
        b(a0 - 1);
    }
    printf("%d\n", 10);
}

```

- The expected output is: 0, 1, 2
- There are 3 different 10s
- The actual output with static allocation is going to be: 0, 0, 0
- ```
void b(int a0) {
    int 10 = a0;
    c(a0);
}
```

- What if there is no apparent recursion, in this case, what if **c()** calls **b()**
- Conclusion: we should allocate arguments and local variables **dynamically**.
- Life of a local (argument/variable)
  - Scope
    - accessible **ONLY** within declaring procedure
    - each execution of a procedure has its own private copy
  - Lifetime
    - allocated when procedure starts
    - de-allocated when procedure returns
  - Activation
    - execution of a procedure
    - starts when procedure is called and ends when it returns
    - there can be many activations of a single procedure "alive" at once
  - Activation Frame
    - memory that stores an activation's state
    - including its locals and arguments
- Should we allocate activation frames from the heap?
  - call **malloc()** to create frame on procedure call and call **free()** on procedure return?
  - we do not allocate the frames in the heap
    - Order of frame allocation and deallocation is special (frames are de-allocated in the reverse order in which they are allocated)
    - We can thus build a very simple allocator for frames
      - start by reserving a **big** chunk of memory for all the frames
      - assuming one knows the address of this chunk

- how to allocate and free frames?  
(add/subtract the pointer by 1)
  - explicit allocation in heap requires more complicated and thus more costly allocation and deallocation to avoid fragmentation
  - Data-structure alike: Last In First Out → Stack
  - Restriction we place on lifetime of local variables and arguments: within the activation of their own procedure call.
- The runtime stack
  - Stack of activation frames
    - stored in memory
    - grows **UPWARDS** from bottom
  - Stack pointer (SP)
    - general purpose register
    - Use **r5**, stores the base address of current frame
  - Top and bottom
    - current frame is top of the stack
    - first activation is the bottom or base
  - Static and Dynamic
    - size of frame is static-ish
    - offset to locals and arguments is static
    - value of stack pointer is dynamic
  - Allocation and de-allocation
    - Allocation is done by decrementing SP
    - De-allocation is done by incrementing SP
- Allocating activation frames

- ```
void foo(int a) {
    int l;
    r5 = malloc(foo_frame_size); // but not really a malloc

    r5->saved_return_address = r6; // somehow
    ...
    l = a; // r5->l = r5->a;
    ...

    free(r5); // but not really a free

}
```

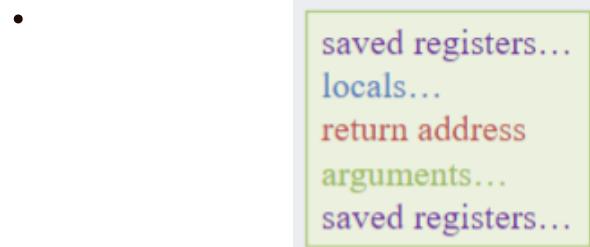
```

•
struct foo_frame {
    int l;
    void* saved_return_address;
    int a;
};
int foo_frame_size = sizeof (struct foo_frame));

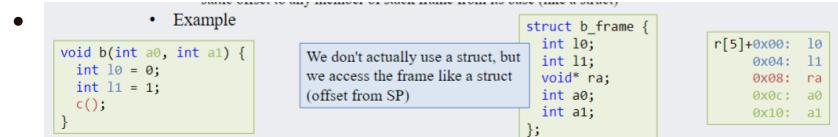
```

- Activation frame details

- Local variables and arguments
- Return address (ra)
  - previously we put this in **r6**, but does not work with consecutive/recursive function calls
  - We then save **r6** on the stack, when necessary; callee will decide
- Other saved registers
  - either or both caller and callee can save register values to the stack
  - do this so that callee has registers it can use
  - anything that is being overwritten needs to be restored
- Stack frame layout



- compiler decides
- based on order convenient for stack creation
- static offset to any member of stack frame from its base (like a **struct**)



- Accessing a local variable or argument

- access like a **struct**
- base address is in **r5** (stack pointer)
- offset is known statically

- ```

int l0 = 0;
int l1 = 1;

struct b_frame {
    int l0;
    int l1;
    void * ra;
    int a0;
    int a1;
};

```

- this becomes

```

ld $0, r0      # r0 = 0
ld r0, (r5)    # l0 = 0
ld $1, r0      # r0 = 1
ld r0, 4(r5)   # l1 = 1

```

- What would be the SM213 assembly code for the following?

```

int l1 = a1;

r[5]+0x00: l0
0x04: l1
0x08: ra
0x0c: a0
0x10: a1

```

- A. ld 16(r5), r0  
st r0, 4(r5)
- B. ld 10(r5), r0  
st r0, 0(r5)
- C. ld 16(r5), r0  
st r0, 0(r5)
- D. ld 10(r5), r0  
st r0, 4(r5)
- E. None of the above

- Correct answer: A
- Explanation: **a1** is offset 16, **l1** is offset 4

What is the value of **l** (in **foo** when it is active)? Assume no default initialization.

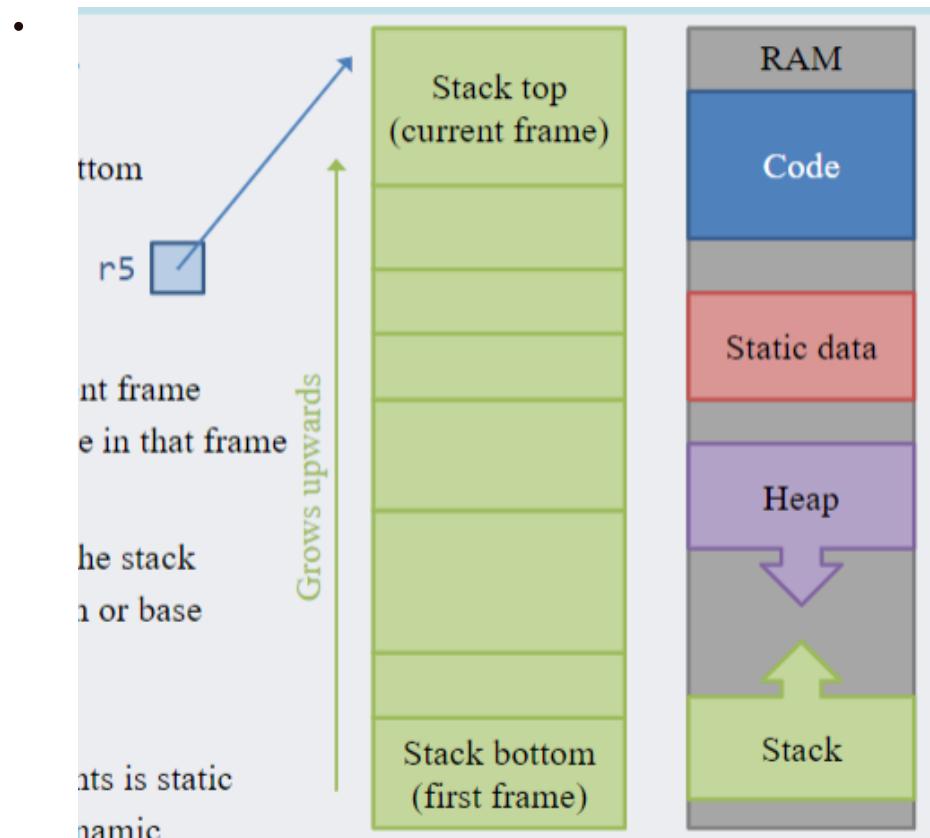
<code>int g;</code>	<code>void goo() {</code>	<code>goo();</code>
<code>void foo() {</code>	<code>    int l = 3;</code>	<code>foo();</code>
<code>}</code>	<code>}</code>	

- A. 0  
B. undefined  
C. it has no value  
D. 3  
E. I don't know

- Correct answer: D
- The stack pointer increments when **goo** returns, so the binary data was left there with no allocation; then when **foo** is just called, and **l** is declared, then the stack pointer decrements, so **l** will take the leftover value from **goo()**

## Feb. 26th - Unit 1e: Cont.d

- The runtime stack
  - Stack of activation frames
    - stored in memory
    - grows upwards from bottom
  - Stack pointer (SP)
    - general purpose register
    - we use **r5**
    - stores base address of current frame
  - Top and Bottom
    - current frame is the top of stack
    - first activation is bottom or base
  - Static and Dynamic
    - size of frame is static(ish)
    - offset to locals and arguments are static
    - value of stack pointer is dynamic



- What is wrong with these functions?

- ```
int * foo() {
    int l;
    return &l;
}
```

```

void bar() {
    int * l = malloc(100);
}

void ping() {
    //
    int * p = foo();
    // when foo returns, p will refer to a de-
    allocated space
    // this leads to a dangling pointer
}

void pong() {
    //
    bar();
    // after bar returns, we have an allocation on
    heap where no
    // variables in the stack can access, this leads
    to a
    // memory leak
}

```

- "Dynamic" arrays in the stack

- Latest version of C allows local arrays to have dynamic size

```

void foo(int n) {
    int a[n];
    int b;
    // ...
    b = 0;
}

```

- Activation frame layouts/offsets are statically determined, so the question is what is the offset for **b**, when this dynamically-sized array can occur?
- What code does the compiler generate for the last statement?
  - Compiler chooses a size for the array, large enough to handle most "reasonable" values at **n**.

This can waste a lot of space if actual **n** is small, but can also possibly overflow when the actual **n** is even larger

  - Compiler convert local array code into array pointer with **malloc** dynamic allocation, inject **free** at end of procedure
- Allocating and deallocating stack frames
  - Allocation of stack is relatively simple
    - Decrement the stack pointer to allocate

- Increment the stack pointer to free
- Whose responsibility is it to allocate in the stack
  - Caller doesn't (need to) know the size of callee's stack
  - However, caller needs to prepare (copy) the arguments
- Division of labour
  - Caller allocates space for arguments
  - Callee allocates space for the rest
- Compiler
  - generates code to allocate and free when procedures are called / return
- Procedure **prologue**
  - code that executes just before procedure starts
    - part in caller before call
    - part in callee at beginning of call
  - allocates activation frame and changes stack pointer
    - subtract frame size from the stack pointer **r5**
  - possibly saves some register values
- Procedure **epilogue**
  - code generated by compiler to execute when procedure ends
    - part in callee just before return
    - part in caller just after return
  - possibly restores some saved register values
  - deallocates activation frame and restore stack pointer
    - add frame size to **r5**
- Stack management

- ```
void foo() {
    b(0, 1);
}

void b(int a0, int a1) {
    int l0 = a0;
    int l1 = a1;
    c();
}
```

- Caller prologue in **foo()** before call
  - allocate stack space for arguments
  - save actual argument values to stack

- ```

•      r[sp] -= 8
        m[0 + r[sp]] <= 0
        m[4 + r[sp]] <= 1
    
```
- Callee prologue in **b()** at start
  - allocate stack space for return address and locals
  - save return address to stack
  - ```

r[sp] -= 12;
m[8 + r[sp]] <= r[6]
    
```
- Callee epilogue in **b()** before return
  - load return address from stack
  - deallocate stack space of return address and locals
  - ```

r[6] <= m[8 + r[sp]]
r[sp] += 12
    
```
- Caller epilogue in **foo()** after return
  - deallocate stack space of arguments
  - ```

r[sp] += 8
    
```

Diagram illustrating the assembly code flow between three functions: **foo()**, **b()**, and **c()**.

The assembly code is annotated with numbered steps corresponding to the diagram:

- 1. caller prologue**: **foo:** deca r5 # allocate callee part of foo's frame  
st r6, 0x0(r5) # save ra on stack
- 2. call**: gpc \$6, r6 # set return address  
j b # b(0, 1)
- 3. callee prologue**: ld \$8, r0 # r0 = 8 = size of caller part of b's frame  
add r0, r5 # allocate callee part of b's frame  
st r6, 0x8(r5) # store return address to stack
- 4. callee body**: ld 12(r5), r0 # r0 = a0  
st r0, 0(r5) # 10 = a0  
ld 16(r5), r0 # r0 = a1  
st r0, 4(r5) # 11 = a1  
gpc \$6, r6 # set return address  
j c # c()
- 5. callee epilogue**: ld 8(r5), r6 # load return address from stack  
ld \$12, r0 # r0 = 12 = size of callee part of b's frame  
add r0, r5 # deallocate callee parts of b's frame  
j 0(r6) # return

Annotations on the right side of the diagram:

- int c()**
- void b(*e*)**
- Jordon Johnson, Geoff**

- Creating the stack
  - Every thread starts with a hidden procedure
    - its name is **start** (or smth like **crt0**)
  - The start procedure

- allocates memory for stack
- initializes the stack pointer
- calls `main()`
- For code from previous slide
  - the main procedure is `foo()`
  - we will statically allocate stack at address `0x1000` to keep simulation simple

```

• start:  ld $stackBtm, r5          # sp =
          address of last word of stack
          inca r5                      # sp =
          address of word after stack
          gpc $0x6, r6                  # r6 = pc + 6
          j foo                         # foo()
          halt
.pos 0x1000
stackTop:    .long 0x0
...
stackBtm:    .long 0x0

```

- What is the value of `r5` in `three()`?
  - (numbers in decimal to simplify math)
  - Write the letter corresponding to your choice

- A. 1964  
 B. 2032  
 C. 1994  
 D. 2004  
 E. 1974  
 F. 2024  
 G. 1968  
 H. None of the above  
 I. I'm not sure

```

void three() {
    int i;
    int j;
    int k;
}

```

```

void two() {
    int i;
    int j;
    three();
}

```

```

void one() {
    int i;
    two();
}

```

```

void foo() {
    // r5 = 2000
    one();
}

```

- Correct Answer: A or G
- Explanation:

- for iClicker 1e.4

```

void three() {
    int i;
    int j;
    int k;
}

```

1968: i
1972: j
1976: k

OR

1964: i
1968: j
1972: k
1976: ra

```

void two() {
    int i;
    int j;
    three();
}

```

1980: i
1984: j
1988: ra

```

void one() {
    int i;
    two();
}

```

1992: i
1996: ra

```

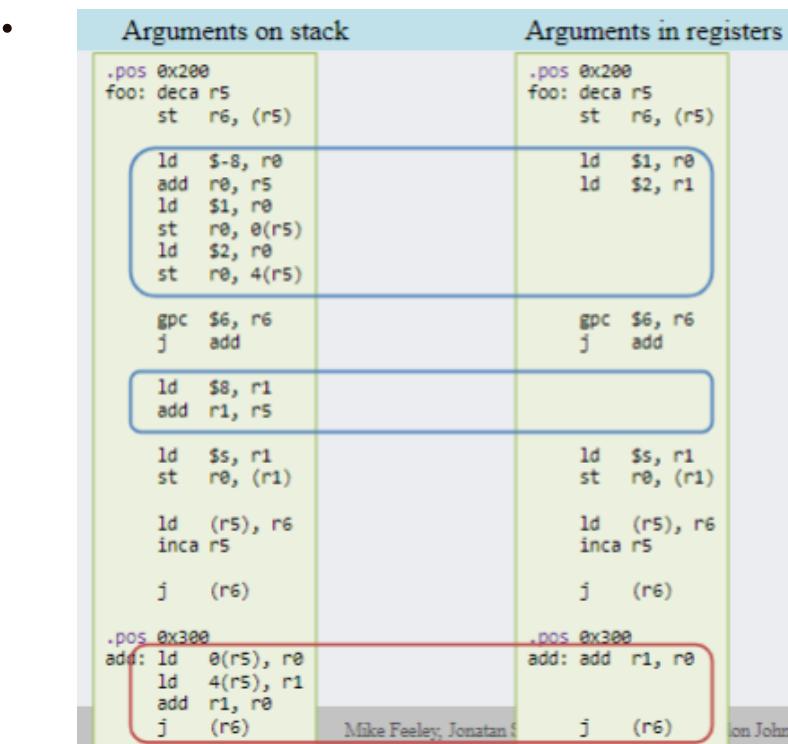
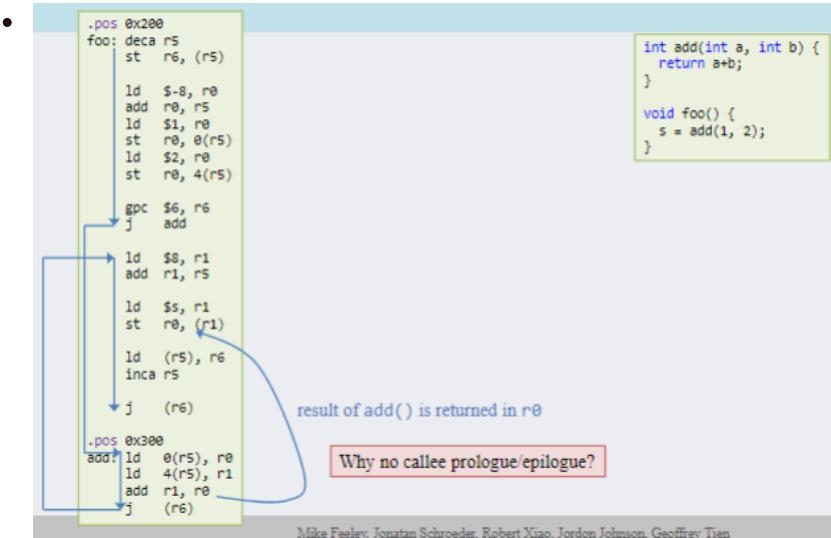
void foo() {
    // r5 = 2000
    one();
}

```

2000: ra
----------

## Feb. 28th - Unit 1e: Cont.d

- Return value, arguments, optimizations
  - **r0**: return value (or, allocate space in activation frame for return value, especially for those types that cannot fit in a register. done in caller prologue); **r5**: stack pointer; **r6**: return address
  - Return value
    - in C and Java, procedures/methods can return only a single value
    - C compilers use a designated register **r0** for this return value
  - Arguments
    - number and size of arguments are statically determined
    - value of actual arguments is dynamically determined
    - the compiler generally chooses to put arguments on the stack
      - caller prologue pushes actual argument values onto stack
      - callee reads/writes arguments from/to the stack
    - sometimes compiler chooses to avoid the stack for arguments
      - caller places argument values in registers
      - callee reads/writes arguments directly from/to these registers
      - WHY does compiler do this? Procedure does not need many local resources
      - WHEN is this a good idea?
      - This reduces memory access: fewer fetching instructions for allocation, placing arguments in activation frame, retrieving arguments in activation frame
  - Other optimizations
    - return address **r6** does not always need to be saved to the stack
      - procedure does not need to call functions
    - local variables are sometimes not needed or used
      - If computation can be entirely done by registers



- Recursion in SM 213

- Consider

```

int proc(int * a, int n) {
    if (n == 0) {
        return 0;
    } else {
        return proc(a, n - 1) + a[n - 1];
    }
}

```

- First, remove names and types for arguments (and local variables)

```

int proc(a0, a1) {
    if (a1 == 0) {
        return 0;
    } else {
        return proc(a0, a1 - 1), a0[a1 - 1];
    }
}

```

- Second, replace C-style conditional. Use comparison to 0; use `goto`; swap then and else ordering

```

int proc(a0, a1) {
    if (a1 == 0)
        goto L0;
    return proc(a0, a1 - 1) + a0[a1 - 1]
    goto L1;
L0: return 0;
L1:
}

```

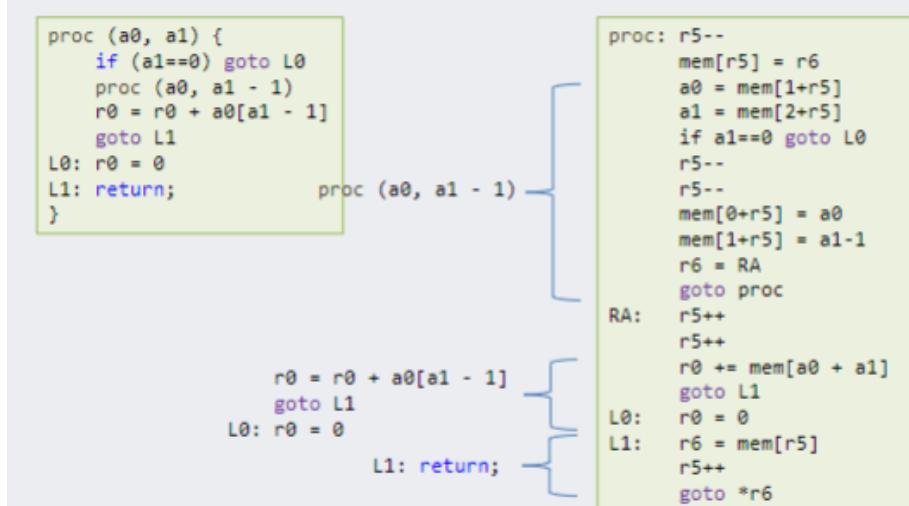
- Third, procedure return value is in `r0`

```

proc (a0, a1) {
    if (a1 == 0) goto L0
    proc(a0, a1 - 1)
    r0 = r0 + a0[a1 - 1]
    goto L1
L0: r0 = 0
L1: return
}

```

- Fourth,



No procedure calls or arrays. Save return address and use `goto` for call and return.  
 Arguments and saved value of return address are on stack, stored in memory.  
 Use global `r5` (global variable) to point to top of stack. Compute array element address.

```

proc: r5--
    mem[r5] = r6
    a0 = mem[1+r5]
    a1 = mem[2+r5]
    if a1==0 goto L0
    r5--
    r5--
    mem[0+r5] = a0
    mem[1+r5] = a1-1
    r6 = RA
    goto proc
RA:   r5++
    r5++
    r0 += mem[a0 + a1]
    goto L1
L0:   r0 = 0
L1:   r6 = mem[r5]
    r5++
    goto *r6

```

Swap the order of a few things. Use global rx variables for all temps.  
Don't trust rx variable values to remain after return from call.

```

proc: r5--
    mem[r5] = r6
    r1 = mem[2+r5]
    if a1==0 goto L0
    r1--
    r2 = mem[1+r5]
    r5--
    r5--
    mem[0+r5] = r2
    mem[1+r5] = r1
    r6 = RA
    goto proc
RA:   r5++
    r5++
    r2 = mem[1+r5]
    r1 = mem[2+r5]
    r1--
    r1 = mem[r2 + r1]
    r0 += r1
    goto L1
L0:   r0 = 0
L1:   r6 = mem[r5]
    r5++
    goto *r6

```

KW2 Mike Fealev, Jonatan Schroeder, Robert Xiao, Jordon Johnson, Geoffrev Tien

```

proc: r5--
    mem[r5] = r6
    r1 = mem[2+r5]
    if a1==0 goto L0
    r1--
    r2 = mem[1+r5]
    r5--
    r5--
    mem[0+r5] = r2
    mem[1+r5] = r1
    r6 = RA
    goto proc
RA:   r5++
    r5++
    r2 = mem[1+r5]
    r1 = mem[2+r5]
    r1--
    r1 = mem[r2 + r1]
    r0 += r1
    goto L1
L0:   r0 = 0
L1:   r6 = mem[r5]
    r5++
    goto *r6

```

Change from C syntax to  
213 assembly syntax. Global  
variables are registers

```

proc: deca r5
    st r6, (r5)
    ld 8(r5), r1
    beq r1, L0
    dec r1
    ld 4(r5), r2
    deca r5
    deca r5
    st r2, (r5)
    st r1, 4(r5)
    gpc $6, r6
    j proc
    inca r5
    inca r5
    ld 4(r5), r2
    ld 8(r5), r1
    dec r1
    ld (r2,r1,4), r1
    add r1, r0
    br L1
L0:   ld $0, r0
L1:   ld (r5), r6
    inca r5
    j (r6)

```

Mike Fealev, Jonatan Schroeder, Robert Xiao, Jordon Johnson, Geoffrev Tien

- Summary: arguments and local variables

- Stack is managed by code that the compiler generates

- grows from bottom up
- push by subtracting
- caller prologue
- callee prologue
- callee epilogue
- caller epilogue

- accessing local variables and arguments

- global variables

- address known statically

- reference variables

- variable stores address of value

- arrays

- elements, named by index

- address of element is base + index \* size of element
- instance variables
  - offset to variable from start of object/struct known statically
  - address usually dynamic
- locals and arguments
  - offset to variables from start of activation frame known statically
  - address of stack frame is dynamic
- Security vulnerabilities with activation frames

- Buffer overflow

• There is a bug in printPrefix

```

void printPrefix (char* str) {
    char buf[10];
    char *bp = buf;

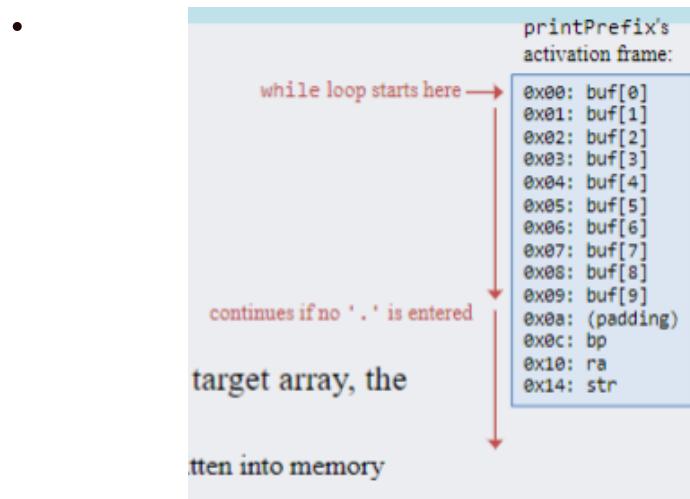
    // copy str up to . into buf
    while (*str != '.')
        *(bp++) = *(str++);
    *bp = 0;
}

// read string from standard input
void getInput (char* b) {
    char* bc = b;
    int n;
    while ((n = fread(bc,1,1000,stdin))>0)
        bc += n;
}

int main (int argc, char** argv) {
    char input[1000];
    puts ("Starting.");
    getInput (input);
    printPrefix (input);
    puts ("Done.");
}

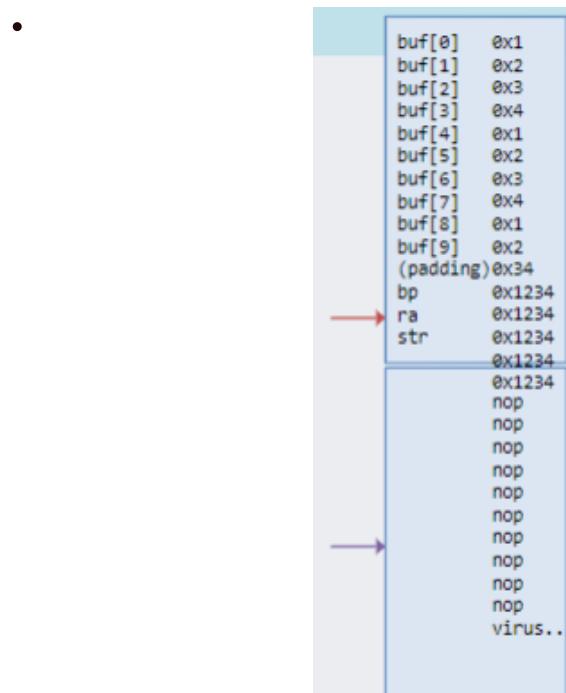
```

- What happens if the string never ends with a '!'?
- Buffer overflow principles
  - Program code is just binary values in memory
  - Program data are also just binary values in memory
  - This means we can include data that **can be interpreted as code**
    - Injected code can do **anything**
    - Most attackers want an interactive shell - inject "shellcode"
  - If the input (e.g. from the previous example) is longer than the target array, the loop will write into memory **beyond the end of the buf array**
    - If an attacker can provide the input, it can determine what values are written into memory beyond the end of **buf**



## Mar. 1st - Unit 1e: Cont.d & Unit 1f: Dynamic Control Flow

- Security Vulnerabilities
  - Buffer overflow
    - The attacker's goal: introduce code into the program from outside; trick program into running it (change the return address to change control flow)
    - Changing the return address: allows attacker to change control flow; if it points to data, then that data becomes the program.
    - Stack-smash (how to make the attack string)



- Hard parts
  - Determining the location of return address in attack string
  - Determining address to change return address to
  - Making it easier

- Approximate return address value (run the program with a big string and see where it crashes)
- Start attack string with many copies of return address
- Next in attack string is long list of **nop** instructions (called the **nop** slide/sled/ramp)
- Finally include the code for the virus
- Works if
  - return address guess is anywhere in the **nop** slide (in the graph example, **0x1234** needs to be in between **nop** and **virus**)
- System calls
  - CPU instruction that signals OS to do smth
  - typically smth that regular processes don't have permission to do
  - examples: read/write terminal, read/write file, execute programs
- Similar to function calls
  - arguments are passed in registers **r0**, **r1**, **r2**
  - values are prepared ahead of time before calling

- New ISA

- Requirements:

- Instruction encodes which system call to use
    - remaining arguments are passed in registers **r0**, **r1**, **r2**

- | Name        | Semantics      | Assembly | Machine |
|-------------|----------------|----------|---------|
| system call | system call #n | sys \$n  | f1nn    |

  - sys \$0: read(fd, buffer, size) – read data from fd (0 = stdin)
    - returns: number of bytes read, or -1 on error
  - sys \$1: write(fd, buffer, size) – write data to fd (0 = stdout)
    - returns: number of bytes written, or -1 on error
  - sys \$2: exec(buffer, size) – execute program
    - returns: 0 if successful, or -1 on error
- | n   | Description                                  | r0  | r1   | r2   | result (r0)                         |
|-----|--|-----|------|------|-------------------------------------|
| 0   | read data from file <i>fd</i><br>(0 = stdin) | fd  | buf  | size | # of bytes read<br>(-1 if error)    |
| → 1 | write data to file <i>fd</i><br>(1 = stdout) | fd  | buf  | size | # of bytes written<br>(-1 if error) |
| 2   | execute program <i>buf</i>                   | buf | size |      | 0 if success<br>-1 if error         |

file descriptor  
 address of source/destination  
 # of bytes  
 result (r0)  
 ↗ r0  
 ↗ r1  
 ↗ r2  
 ↗ result (r0)  
 ↗ buf  
 ↗ size  
 ↗ fd  
 ↗ address of string containing program path

```

.pos 0x1000
    ld  $1, r0
    ld  $str, r1
    ld  $12, r2
    sys $1
    halt

.pos 0x2000
str: .long 0x68656c6c # hell
      .long 0x6f20776f # o wo
      .long 0x726c640a # rld\n

```

- Write 12 bytes starting from address **str** to **stdoutput**
- If we call **sys \\$3**, the **buf** is the address of **str** to program path
- Protecting against buffer overflow attacks
  - What if the stack grew downwards?
  - Active frame at highest address
  - Modern protections
    - Non-executable stack
    - Canaries
    - Randomized stack addresses

## Unit 1f: Dynamic Control Flow

- Review: static procedure calls
  - Consider

```

void foo() {
    bar();
}

void bat() {
    bar();
}

void bar() {
    // does smth
}

```

- **bar()**'s return will be dynamic

- Suppose we have a Java class `A` with a procedure `foo`.
  - Suppose this class has subclasses `B` and `C`
    - `B` has its own version of `foo`, but `C` does not.
- If procedure `bar` gets called, which version of `foo` is called?
  - A. `A`'s version
  - B. `B`'s version
  - C. `C`'s version
  - D. All of them, and we choose between the result when they return
  - E. It cannot be statically determined in general

```
void bar (A a) {
    a.foo();
}
```

- Correct Answer: E
- Explanation: consider

```
A obj1 = new A();
B obj2 = new B();

void bar(A a) {
    a.foo();
}

bar(obj1);      // works, executes A's
version
bar(obj2);      // works, executes B's
version
```

- Dynamic Control Flow
- Calling different functions - a different approach via indirect jumps

- Because **functions have addresses**

- We can make **pointers to functions**

- Useful for

- Polymorphism
- Generic operations
- Jump tables

- ```
beq r0, L2
L1: ld $func1, r1 # run func1 if r0 != 0
     br L3
L2: ld $func2, r1 # run func2 if r0 == 0
L3: gpc $2, r6
     j (r1)
     ...
```

- Dynamic jumps in C
- **Function pointer**
  - a variable that stores a pointer to a procedure
  - declared as
 

```
<return-type> (*<variable-name>) (<formal-argument-list>);
```
  - used to make dynamic call
 

```
<variable-name>(<actual-argument-list>);
```

- Example

- ```
void ping() {
    // smth
}

void foo() {
    void (* ptr_func) ();
    ptr_func = ping;
    ptr_func();
}
```

- assign using procedure name without (), call using variable name with ()

`int (*x) (char *);`  
What does the variable x above represent?

- A. A function that receives a char pointer as argument and returns an int pointer
- B. A function that receives a function pointer as argument and returns an int pointer
- C. A pointer to a function that receives a char pointer as argument and returns an int
- D. A pointer to a function that receives a function pointer as argument and returns an int
- E. What you get when you fall asleep on your keyboard

- - Correct Answer: C

## Mar. 4th - Unit 1f: Cont.d

`int (*x) (char *);`  
Which of the following is a correct way to call the procedure to which x points?  
Assume we have a `char* c` that has been set to a properly allocated chunk of memory.

- A. `int i = *(*x)(c);`
- B. `int i = *x(*c);`
- C. `int i = x(c);`
- D. `int i = x(*c);`
- E. `int i = *((int*) c);`

- - Correct Answer: C
- Polymorphism
  - Invoking a method on an object in Java
    - variable that stores the object has a static type (apparent type)
    - the object reference is dynamic and so is its type

- object's actual type must be a subtype of the apparent type of the referring variable
- but object's actual type may override methods of the apparent type
- Polymorphic Dispatch

- target method address depends on the type of the reference object
- one call site can invoke different methods at different times

- ```

class A {
    void ping() {}
    void pong() {}
}

class B extends A {
    void ping() {}
    void wiff() {}
}

```

- ```

static void foo(A a) {
    a.ping();
    a.pong();
}

static void bar() {
    foo(new A());
    foo(new B());
}

```

- Polymorphic dispatch

- Method address is determined dynamically
  - compiler cannot hard code target address in procedure call
  - instead, compiler generates code to lookup procedure address at runtime
  - address is stored in memory in the object class' jump table

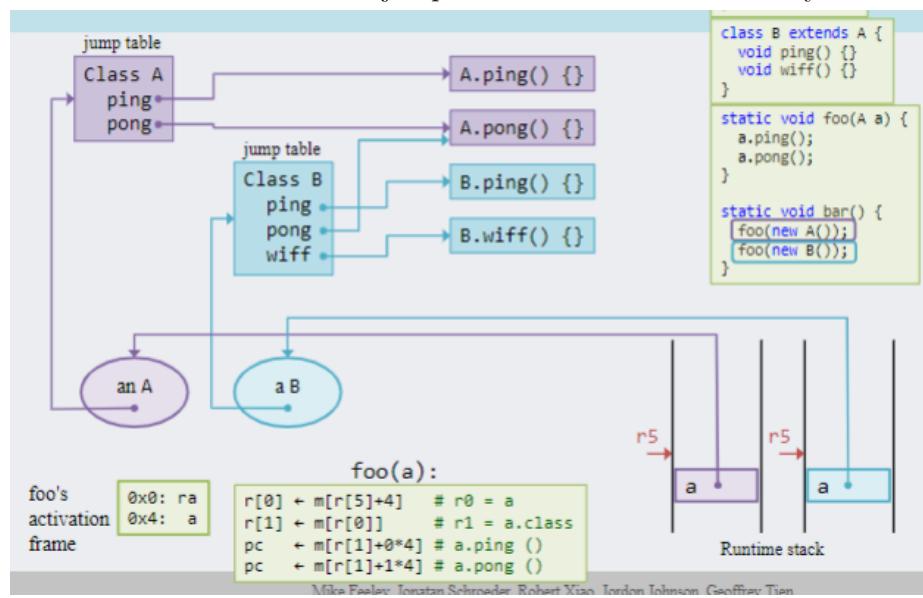
- Class **jump table**

- every class is represented by a class object
- objects store a pointer to their class object
- the class object stores the class' jump table
- the jump table stores the address of methods implemented by the class

- Static and Dynamic of method invocation

- address of jump table is determined dynamically

- objects of different actual types will have different jump tables
- method's offset into jump table is determined statically



- Same procedure should be expected to be at the same offset for subclasses as their parent class
- Polymorphism in C

- Use **struct** to store jump table
  - Declaration of class

```

struct A_class {
    void (* ping) (void *);
    void (* pong) (void *);
};

```

- Declaration of instance methods

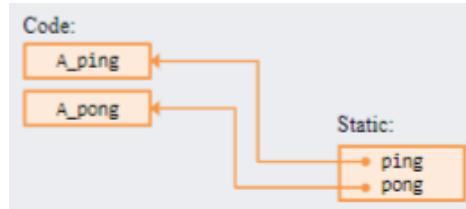
```

void A_ping(void * thisv) {
    printf("A_ping\n");
}
void A_pong(void * thisv) {
    printf("A_pong\n");
}
// ???
void A_ping(void * thisv) {
    struct A * this = thisv;
    printf("A_ping %d\n", this -> i);
}

```

- Static allocation and initialization of class object (i.e. jump table)

```
struct A_class A_class_table {
    A_ping, A_pong;
};
```



- Object (instance of class)

- Object template

```
struct A {
    struct A_class * class; // pointer to
    class object
    int i; // instance's
    attribute
};
```

- Constructor method

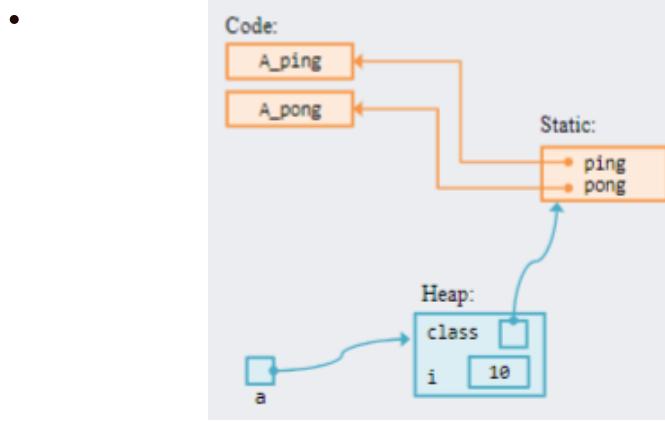
```
struct A * new_A(int i) {
    struct A * obj = malloc(sizeof(struct
        A));
    obj -> class = &A_class_table;
    obj -> i = i;
};
```

- Allocating an instance

```
struct A * a = new_A(10);
```

- Calling instance methods

```
a -> class -> ping(a);
a -> class -> pong(a);
```



- class **B extends A**
  - B's class struct (jump table) is a super-set of A's

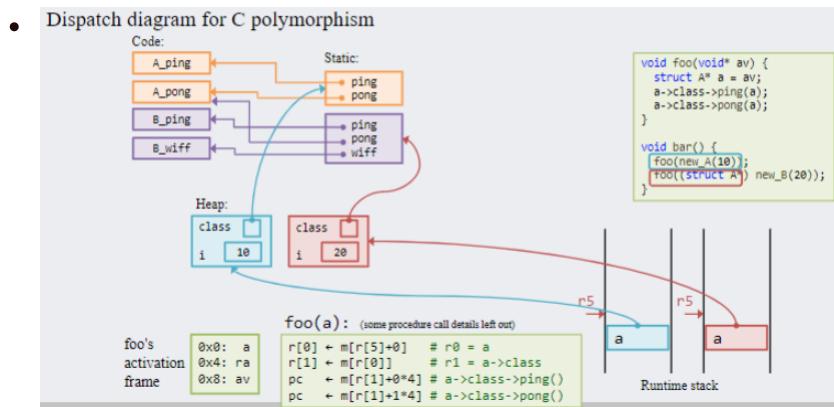
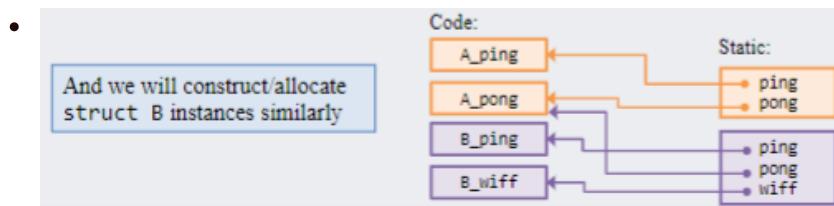
```
struct B_class {
    void (* ping) (void *);
    void (* pong) (void *);
    void (* wiff) (void *);
};
```

- B's method declarations and class object (static) allocation

```
void B_ping(void * thisv) {
    printf("B_ping\n");
}

void B_wiff(void * thisv) {
    printf("B_wiff\n");
}

struct B_class B_class_table = {
    B_ping, A_pong, B_wiff;
};
```



- ISA for polymorphic dispatch.

- How to compile `a -> class -> pong()`?
- Pseudocode: `pc <- m[r[1]] + 0 * 4`
- new ISA

Name	Semantics	Assembly	Machine
jump	<code>pc &lt;- a</code>	<code>j a</code>	<code>b--- aaaaaaaaa</code>
indirect jump	<code>pc &lt;- r[t] + o (or r[t] + pp*2)</code>	<code>j o(rt)</code>	<code>c(pp)</code>
dbl-ind jump b+o	<code>pc &lt;- m[r[t] + o] (or m[r[t] + pp*2])</code>	<code>j *(rt)</code>	<code>d(pp)</code>

- What is the difference between these two C snippets? Assume `bat` is called.

(1) <pre>void foo() { printf("foo\n"); }  void go(void(*proc)()) {     proc(); }  void bat() {     go(foo); }</pre>	(2) <pre>void foo() { printf("foo\n"); }  void go() {     foo(); }  void bat() {     go(); }</pre>
---	--

- A. (2) calls `foo`, but (1) does not
  - B. (1) is not valid C
  - C. (1) jumps to `foo` using a dynamic address and (2) a static address
  - D. They both call `foo` using dynamic addresses
  - E. They both call `foo` using static addresses
- Correct Answer: C

## Mar. 6th - Unit 1f: Cont.d

- Generic Functions
  - Consider the linked list implementation and print procedure below

```
struct node {
    int key;
    int value;
    struct node * next;
};

void print_list(struct node * list) {
    while (list) {
        printf("%d: %d\n", list -> key, list ->
value);
        list = list -> next;
    }
}
```

- Issues:
  - The `print_list` procedure is very specific
    - What if we need to print in a different format?

```

void print_list_full(struct node *
list) {
    while(list) {
        printf("key: %d, value: %d\n",
list -> key, list -> value);
        list = list -> next;
    }
}

```

- What if we want to print to a file?
- What if we want to perform another task (besides printing) with each key-value pair?
- Solution: function parameter
  - Extra parameter to `print_list`: what to do with each key-value pair
  - ```

void follow_list(struct node * list,
void (* fn)(int, int)) {
    while(list) {
        fn(list -> key, list ->
value);
        list = list -> next;
    }
}

```

- Function pointer `void (* fn)(int, int)` decides what to do with each key-value pair
- What to do with a key-value pair?

```

•
void print_element(int key, int value)
{
    printf("%10d: %10d\n", key,
value);
}

void print_element_full(int key, int
value) {
    printf("key: %d, value: %d\n",
key, value);
}

follow_list(list, print_element);
follow_list(list, print_element_full);

```

```

• follow_list:
  deca r5
  st r6, (r5) # save own r.a.
  ld 4(r5), r0 # r0 = list
loop:
  beq r8, L0 # while (list)
  ld 8(r5), r1 # r1 = fn
  ld 0(r0), r2 # r2 = list->key
  ld 4(r0), r3 # r3 = list->value
  deca r5
  deca r5 # allocate arg space for fn
  st r2, 0(r5) # arg0 = key
  st r3, 4(r5) # arg1 = value

gpc $2, r6 # prep r.a. for fn
j (r1) # call fn
inca r5
inca r5 # recover arg space for fn
ld 4(r5), r0 # r0 = list
ld 8(r0), r0 # list = list->next
br loop # repeat from loop
L0:
  ld (r5), r6 # restore r.a.
  inca r5 # recover own r.a. space
  j (r6) # return

```

- Other uses of function pointers

- e.g. Quicksort, for sorting integers

```

• int partition (int* array, int left, int right, int pivotIndex) {
    int pivotValue, t;
    int storeIndex, i;

    pivotValue      = array [pivotIndex];
    array [pivotIndex] = array [right];
    array [right]     = pivotValue;
    storeIndex = left;
    for (i=left; i<right; i++) {
        if (array [i] <= pivotValue) {
            t           = array [i];
            array [i]     = array [storeIndex];
            array [storeIndex] = t;
            storeIndex += 1;
        }
        t           = array [storeIndex];
        array [storeIndex] = array [right];
        array [right]     = t;
        return storeIndex;
    }
    void quicksort (int* array, int left, int right) {
        int pivotIndex;

        if (left < right) {
            pivotIndex = partition (array, left, right, left + (right-left)/2);
            quicksort (array, left,          pivotIndex - 1);
            quicksort (array, pivotIndex + 1, right       );
        }
    }
    void sort (int* array, int n) {
        quicksort (array, 0, n-1);
    }
}

```

- The logic/code for **quicksort/partition** is **mostly** type-independent

- change parameter to sort anything

- actually, like Java, array parameter will be a pointer to anything (or anything the same size as a pointer)

```

int partition ( void** array, int left, int right, int pivotIndex) {
    void* pivotValue, * t;
    int   storeIndex, i;           Actually, only 3 parts of the code are type-dependent
    pivotValue      = array [pivotIndex];   • array parameter – easy to deal with
    array [pivotIndex] = array [right];   • pivotValue type – easy to deal with
    array [right]     = pivotValue;      • comparison – ???
    storeIndex = left;
    for (i=left; i<right; i++) {
        if (array [i] <= pivotValue) {
            t           = array [i];
            array [i]     = array [storeIndex];
            array [storeIndex] = t;
            storeIndex += 1;
        }
        t           = array [storeIndex];
        array [storeIndex] = array [right];
        array [right]     = t;
        return storeIndex;
    }
}

```

frey Tien

- **void \*** means something
- **<=** operator may not be meaningful for all types

- In Java,

```

int partition (Comparable<T> array[], int left, int right, int pivotIndex) {
    Comparable<T> pivotValue, t;
    int storeIndex, i;

    pivotValue = array [pivotIndex];
    array [pivotIndex] = array [right];
    array [right] = pivotValue;
    storeIndex = left;
    for (i=left; i<right; i++)
        if (array [i] .compareTo (pivotValue)) <= 0) {
            ...
        }
}

class ComparableInteger<Integer> extends Integer {
    @Override
    int compareTo(Integer i) {
        return intValue() < i.intValue() ? -1: intValue() == i.intValue() ? 0: 1;
    }
}

```

- In C, use a comparator function pointer

```

int partition(void* * array, int left,
             int right, int pivotIndex, int (* cmp)
             (void*, void*)) {
    // cmp is the specific function
    // that determines how to compare the
    // elements.
}

```

- ```

int partition (void** array, ... , int (*cmp) (void*, void*)) {
    void* pivotValue, *t;
    int storeIndex, i;

    pivotValue = array [pivotIndex];
    array [pivotIndex] = array [right];
    array [right] = pivotValue;
    storeIndex = left;
    for (i=left; i<right; i++) {
        if (cmp (array [i], pivotValue) <= 0) {
            t = array [i];
            array [i] = array [storeIndex];
            array [storeIndex] = t;
            storeIndex += 1;
        }
        t = array [storeIndex];
        array [storeIndex] = array [right];
        array [right] = t;
    }
    return storeIndex;
}

```

- Side note: avoiding **void \*** confusion

- It is a pointer type that cannot be dereferenced (directly)
  - **opaque** pointer
- in C we used it to represent a pointer to **anything**
- before using, must be casted to another pointer type
  - There is no type checking on this type case
  - casting can be done implicitly; explicit cast is not needed
- code with **void\*** can be confusing, especially when encountering **verb\* \***
  - C's **typedef** statement
    - creates a new type name for a type expression
    - leads to more readable code

- common convention: end type names with a  
" \_t "

- ```
Using void*:
int partition ( void** array, ..., int (*cmp) (void*, void*)) {
    void* pivotValue, *t;
    ...
}
```

Using new type defined by `typedef`:

```
typedef void* element_t;

int partition ( element_t* array, ..., int (*cmp) (element_t, element_t)) {
    element_t pivotValue, t;
    ...
}
```

- Using the parametrized Quicksort

- ```
int cmpIntegers(void* av, void* bv) {
    int * a = av;
    int * b = bv;
    return *a < *b ? -1 : (*a == *b ? 0 : 1);
}

int a[] = {3, 8, 1};
int * pa[] = {a, a + 1, a + 2};
sort((void* *) pa, 3, cmpIntegers);
// Note: pa's contents will be
// rearranged so its dereferenced
// elements will be in increasing order;
// *pa[0] < *pa[1] < *pa[2], the array a
// is unchanged
```

- To sort strings (with `typedef`):

```
int cmpStrings (element_t av, element_t bv) {
    char* a = av;
    char* b = bv;
    return *a < *b ? -1: *a == *b ? 0: 1; } or simply:
```

```
char* array[] = {"Psyduck", "Bulbasaur", "Meowth", ..., "Eevee"};
sort ((element_t*) array, sizeof (array) / sizeof (array[0]), cmpStrings);
```

## Mar. 8th - Unit 1f: Cont.d

- Function pointers in Racket

- `map` receives a function and two lists, applies function to pairs of elements

- `(map + (list 1 4 3) (list 7 2 5)) => (list 8 6 8)`
- `(map (lambda (a b) (+ a b)) (list 1 4 3) (list 7 2 5)) => (list 8 6 8)`
- `(map max (list 1 4 3) (list 7 2 5)) => (list 7 4 5)`

- Other iterators

- `(foldl + 0 (list 1 2 3))`

- (`filter (lambda (a) (> a 3)) (list 1 2 3 4 5))`
- Implementing `map` in C

- `map` definition (using function pointer as argument)

```
void map(int (* fn) (int, int), int n, int *
s0, int * s1, int * d) {
    int i;
    for (i = 0; i < n; i += 1) {
        d[i] = fn(s0[i], s1[i]);
    }
}
```

- Example function to be applied to two elements

```
int add(int a, int b) {
    return a + b;
}
```

- Calling map with add

```
map(add, 3, a, b, c);
```

- To be even more generic, we can use arrays of `void*` instead of `int`
- Implementing `foldl`

```
int foldl(int (* fn)(int, int), int v, int * a, int
n) {
    int i;
    for (i = n - 1; i >= 0; i -= 1) {
        v = fn(v, a[i]);
    }
    return v;
}
```

```
int add(int a, int b) {
    return a + b;
}
```

```
printf("%d\n", foldl(add, 0, a, 3));
```

- Could we use this to concatenate strings? YES (as long as there is enough space to hold the strings)

- However, we need to be careful with dynamically allocated memory to avoid leaks
- TODO: look up the C library function `realloc`
- Reference counting with function pointers
  - Recall A6: `value` needs the same refcount as the element (because `value` can't be automatically be freed otherwise)

```

•
struct element {
    int num;
    char* value;
};

struct element* element_new(int num, char* value) {
    struct element* e = rc_malloc(sizeof(*e));
    e->value = rc_malloc(strlen(value) + 1);
    ...
}

void element_keep_ref(struct element* e) {
    rc_keep_ref(e->value);
    rc_keep_ref(e);
}

void element_free_ref(struct element* e) {
    rc_free_ref(e->value);
    rc_free_ref(e);
}

```

- What if our program has multiple structure definitions with nested allocations?
- Finalizers:

```

struct element {
    int num;
    char* value;
};

void element_finalizer(void* ev) {
    struct element* e = ev;
    free(e->value);
}

struct element* element_new(int num, char* value) {
    struct element* e = rc_malloc(sizeof(*e));
    e->value = malloc(strlen(value) + 1);
    ...
}

void element_keep_ref(struct element* e) {
    rc_keep_ref(e);
}

void element_free_ref(struct element* e) {
    rc_free_ref(e);
}

```

But how/when does this finalizer get called?  
How do we "automate" it?

- Fixing the refcount library:

```

struct rc_metadata {
    int ref_count;
    void (*finalizer) (void* ev); // finalizer associated with each allocation
};

void* rc_malloc(int size, void (*ef) (void* ev)) {
    struct rc_metadata* md = malloc(size + sizeof(struct rc_metadata));
    md->ref_count = 1;
    md->finalizer = ef;

    return md + 1; // address of payload
}

void rc_free_ref(void* p) { // p is a payload address
    struct rc_metadata* md = p;
    md -= 1; // set md to start of metadata
    (md->refcount) -= 1;
    if (md->refcount == 0) {
        md->finalizer(p); // call the allocation-specified finalizer
        free (ref_count); // free the entire allocation
    }
}

```

We should actually check that the provided finalizer is not NULL before calling it

- All `rc_malloc` now receives two parameters, the size of data, and function pointer to deallocate
- should be `free(md)`
- For `struct` without nested allocations, we can use `NULL` for finalizer argument

- Usage of the modified `rc_malloc`, etc.

```

struct element {
    int num;
    char* value;
};

void element_finalizer(void* ev) {
    struct element* e = ev;
    free(e->value);
}

struct element* element_new(int num, char* value) {
    struct element* e = rc_malloc(sizeof(*e), element_finalizer);
    e->value = malloc(strlen(value) + 1);
    ...
}

void element_keep_ref(struct element* e) {
    rc_keep_ref(e);
}

void element_free_ref(struct element* e) {
    rc_free_ref(e);
}

```

Just use `rc_keep_ref(e)` and `rc_free_ref(e)`, for all allocations

- `switch` statement

<pre> int i; int j;  <b>void foo() {</b>     switch (i) {         case 0: j = 10; break;         case 1: j = 11; break;         case 2: j = 12; break;         case 3: j = 13; break;         default: j = 14; break;     } } </pre>	<pre> <b>void bar() {</b>     if (i == 0)         j = 10;     else if (i == 1)         j = 11;     else if (i == 2)         j = 12;     else if (i == 3)         j = 13;     else         j = 14; } </pre>
--	--

- Semantically, the same as simplified nested/cascading if statements
  - where condition of each `if` tests the same variable
  - unless you leave out the `break` at the end of a case block
- Why bother?
  - for humans, ease of writing, and reading of code?
  - for compilers, permitting a more efficient implementation?

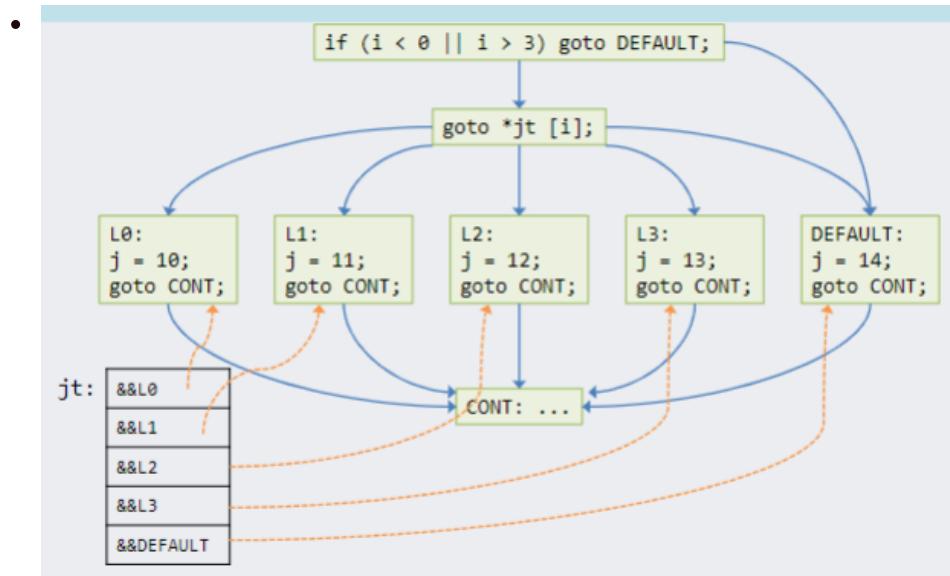
- Implementing switch statements
    - anything more to consider than just conditionals?
  - Human vs compiler
    - Benefits for human: the syntax models a common idiom, choosing one computation from a set
    - Restrictions
      - Case labels must be **static**, **cardinal** values
        - Cardinal value is a number that specifies a position relative to the beginning of an ordered set
        - integers are cardinal, strings are not
      - Case labels must be compared for equality to a single dynamic expression
        - some languages permit the expression to be an inequality
- ```
switch (pokemonType) {
    case "water": ...
    case "grass": ...
    case "psychic": ...
}
```

```
switch (i, j) {
    case i > 0: ...
    case i == 0 && j > a: ...
    case i < 0 && j == a: ...
}
```
- Why compiles like **switch** statements
    - Switch condition evaluates to a number
    - Each case arm has a distinct number
    - So, build a table with the address of every case arm, indexed by case value
    - Switch by indexing into this table and jumping to matching case arm (a **jump table**)
    - Example:
 

|                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                  |                                                                                                               |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|
| <pre><code>switch (i) {     case 0: j = 10; break;     case 1: j = 11; break;     case 2: j = 12; break;     case 3: j = 13; break;     default: j = 14; break; }</code></pre> | <pre><code>void* jt[4] = { &amp;&amp;L0, &amp;&amp;L1, &amp;&amp;L2, &amp;&amp;L3 }; if (i &lt; 0    i &gt; 3) goto DEFAULT; goto *jt [i]; L0: j = 10; goto CONT; // "break" L1: j = 11; goto CONT; L2: j = 12; goto CONT; L3: j = 13; goto CONT; DEFAULT: j = 14; goto CONT; CONT:</code></pre> | <p>Note:<br/> <b>&amp;&amp;</b> is a construct that "proper" C doesn't support, but the GCC compiler does</p> |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|

## Mar. 11th - Unit 1f: Cont.d, Unit 2a: IO & Asynchrony

- **switch** statement control flow with jump table



- Happy compilers mean happy people
  - Computation can be much more efficient (compare the running time to **if**-based alternative)
    - save time on condition testing
  - Could this all go horribly wrong?
    - construct a switch statement where this implementation technique is a really bad idea
- Basic implementation strategy
  - General form of a **switch** statement

```

switch(<cond>) {
    case <label_i>: <code_i>
    default:         <code_default>
}

```

- Naive strategy

```

goto address of code_default if cond >
max_label_value
goto address of in jumptable[label_i]

statically: jumptable[label_i] = address of code_i
forall label_i

```

- 2 additional considerations
  - Case labels are not always continuous
    - **case 0: ...; case 1: ...; case 2: ...; case 5: ...;**
  - the lowest case label is not always 0
    - **case 20: ...; case 21: ...; case 22: ...;**
- Non-contiguous case labels

```

switch(i) {
    case 0:      j = 10; break;
    case 3:      j = 13; break;
    default:    j = 14; break;
}

```

- Problem: indexing based on case number, jump table will have some gaps
- Solution: fill gaps in jump table using default branch

- **OFFSET**

|   |           |
|---|-----------|
| 0 | &&L0      |
| 1 | &&DEFAULT |
| 2 | &&DEFAULT |
| 3 | &&L3      |
| 4 | &&DEFAULT |

- Case labels not starting at 0

```

switch(i) {
    case 1000:      j = 10; break;
    case 1001:      j = 11; break;
    case 1002:      j = 12; break;
    case 1003:      j = 13; break;
    default:        j = 14; break;
}

```

- Problem: invalid indices (negative numbers), or many unused (default) entries in the table
- Solution: Normalize the case labels relative to 0.

- Implementing **switch** statements

- Choose strategy
  - use jump table unless case labels are sparse or there are very few of them
  - use nested **if**-statements otherwise
- Jump table strategy
  - Statically: build jump table for all label values between lowest and highest
  - Generate code to:
    - **goto** default if condition is less than minimum case label or greater than maximum
    - normalize condition value to lowest case label

- use jump table to go directly to code-selected case arm

- ```
goto address of code_default if cond < min_label_value
goto address of code_default if cond > max_label_value
goto address in jumptable [cond - min_label_value]

statically: jumptable [i - min_label_value] = address of code_i
            forall i: min_label_value <= i <= max_label_value
```

```
switch (i) {
    case -20: ... break;
    case 0: ... break;
    case 5: ... break;
    case 42: ... break;
    default: ... break;
}
```

Which of the following implementations is the most appropriate in this case?

- A. A jump table
- B. A sequence of "if" statements
- C. Either option is equally appropriate
- D. Neither – we're not even going to bother with implementing this

- Correct Answer: B
- Explanation: the case label range is 62, large gaps between case labels, base + offset access can only support maximum of 254 bytes of offset (maximum size should be around 63 entries)

```
switch (i) {
    case -207: ... break;
    case -208: ... break;
    case -209: ... break;
    case -210: ... break;
    case -211: ... break;
    case -212: ... break;
    case -214: ... break;
    case -215: ... break;
    default: ... break;
}
```

Which of the following implementations is the most appropriate in this case?

- A. A jump table
- B. A sequence of "if" statements
- C. Either option is equally appropriate
- D. Neither – we're not even going to bother with implementing this either

- Correct Answer: A
- Explanation: Normalize the case labels relative to 0

```
switch (i) {
    case 20: j = 10; break;
    case 21: j = 11; break;
    case 23: j = 13; break;
    default: j = 14; break;
}
```

```
static const void* jt[4] = { &&L20, &&L21, &&DEFAULT, &&L23 };
if (i < 20 || i > 23) goto DEFAULT;
goto *jt [i-20];
L20: j = 10;
      goto CONT; // "break"
L21: j = 11;
      goto CONT;
L23: j = 13;
      goto CONT;
DEFAULT:
      j = 14;
      goto CONT;
CONT:
```

- ```

foo:    ld $i, r0          # r0 = &i
        ld 0x0(r0), r0      # r0 = i
        ld $0xffffffffed, r1 # r1 = -19
        add r0, r1           # r0 = i-19
        bgt r1, 10            # goto 10 if i>19
        br default            # goto default if i<20
10:     ld $0xffffffe9, r1 # r1 = -23
        add r0, r1           # r1 = i-23
        bgt r1, default       # goto default if i>23
        ld $0xfffffec, r1    # r1 = -20
        add r1, r0           # r0 = i-20
        ld $jmptable, r1     # r1 = &jmptable
        j *(r1, r0, 4)       # goto jmptable[i-20]

jmptable: .long case20      # & (case 20)
          .long case21      # & (case 21)
          .long default       # & (case 22)
          .long case23       # & (case 23)

```
- ```

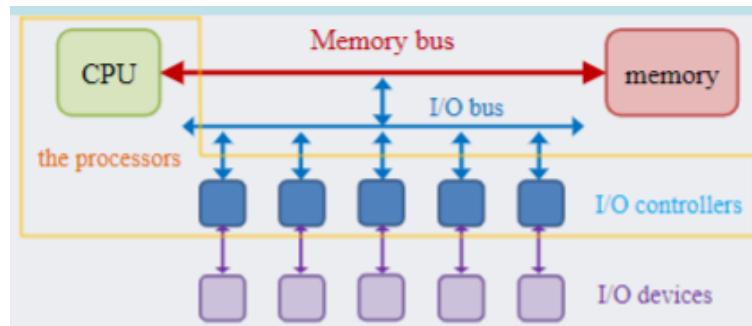
case20:   ld $0xa, r1      # r1 = 10
          br done           # goto done
...
default:   ld $0xe, r1      # r1 = 14
          br done           # goto done
done:     ld $j, r0          # r0 = &j
          st r1, 0x0(r0)    # j = r1
          br cont           # goto cont

```

- ISA

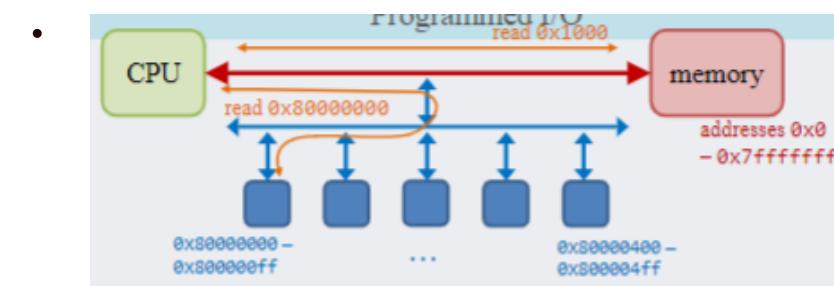
- | Name              | Semantics  | Assembly  | Machine        |
|-------------------|--|-----------|----------------|
| branch            | $pc \leftarrow a$ (or $pc + p^*2$ )                | br a      | 8-pp           |
| branch if equal   | $pc \leftarrow a$ (or $pc + p^*2$ ) if $r[c] == 0$ | beq rc, a | 9cpp           |
| branch if greater | $pc \leftarrow a$ (or $pc + p^*2$ ) if $r[c] > 0$  | bgt rc, a | acpp           |
| jump              | $pc \leftarrow a$                                  | j a       | b--- aaaaaaaaa |
- | Name          | Semantics                                    | Assembly | Machine |
|---------------|--|----------|---------|
| indirect jump | $pc \leftarrow r[t] + o$ (or $r[t] + p^*2$ ) | j o(rt)  | ctpp    |
- | Name                 | Semantics   | Assembly       | Machine   |
|----------------------|---|----------------|-----------|
| dbl-ind jump b+o     | $pc \leftarrow m[r[t] + o]$ (or $m[r[t] + pp^*2]$ ) | j *o(rt)       | dtp       |
| dbl-ind jump indexed | $pc \leftarrow m[r[t] + r[i]*4]$                    | j *(rt, ri, 4) | eti- new! |

## Unit 2a: IO & Asynchrony



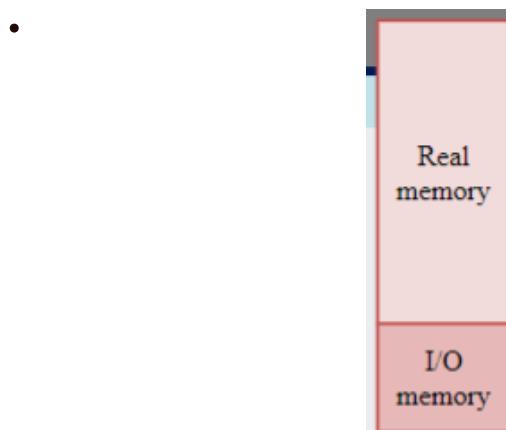
- Memory bus
    - data/control path connecting CPU, main memory, and I/O bus
    - also called the **front-side** bus
  - I/O bus
    - data/control path connecting memory bus and I/O controllers
    - e.g. PCI
  - I/O controller
    - a processor running software (firmware)
    - connects I/O device to I/O bus

- e.g. SCSI, SATA, Ethernet
- I/O device
  - I/O mechanism that generates or consumes data
  - e.g. disk, radio, keyboard, mouse
- Talking to an I/O controller
  - Programmed I/O (PIO)
    - CPU transfers one **word** at a time between CPU and IO controller
    - typically use standard load/store instructions but to IO-mapped memory address
  - IO-mapped memory
    - memory addresses beyond the end of main memory
    - used to name IO controllers (usually configured at boot time)
    - loads and stores are translated into IO bus messages to controller
- Example



- read/write to controller at address 0x80000000

```
ld $0x80000000, r0
st r1, (r0)
ld (r0), r1
```

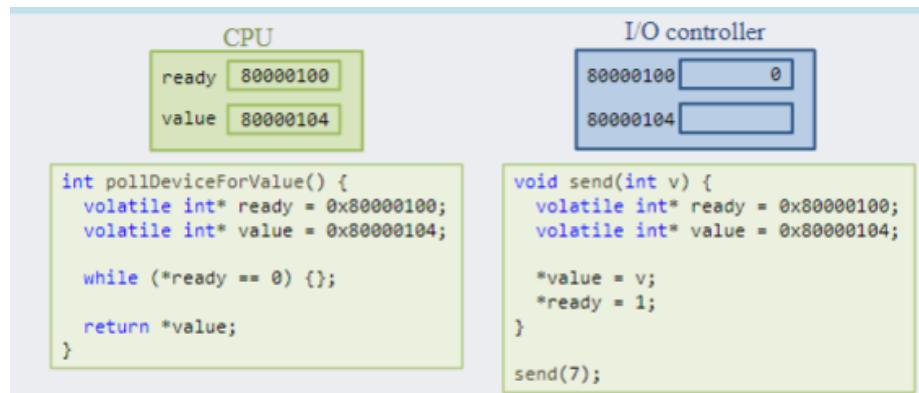


- Limitations of PIO
  - Reading or writing large amount of data slows CPU
    - CPU must transfer one word at a time
    - controller/device is much slower than CPU

- so, CPU runs at controller/device speed, mostly waiting for controller
- IO controller cannot initiate communication
  - sometimes the CPU asks for data
  - but, sometimes the controller devices receives data for the CPU, without CPU asking (mouse click or network packet reception)
- How does controller notify CPU that it has data the CPU should want
- One not-so-good idea
  - CPU periodically asks what controllers offer that CPU needs
  - Drawbacks: asking for data takes time, might still not have any data
  - When is it OK? If there is a high chance of having data ready

## Mar. 13th - Unit 2a: Cont.d

- Polling and IO memory

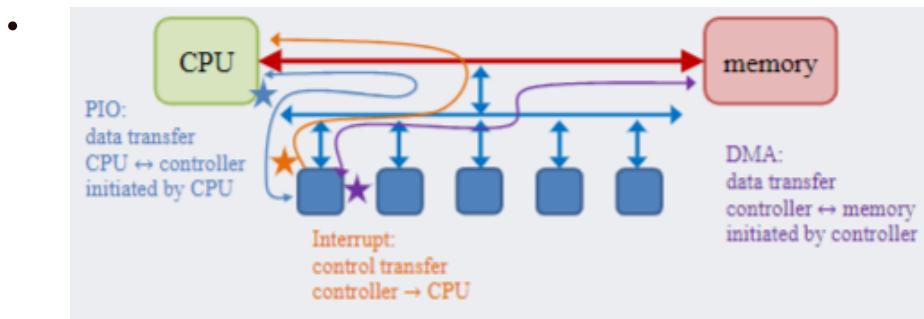


- **volatile:** data is not guaranteed to be steady, can be changed by external factors
- Reading (or writing) IO memory is slow
  - CPU → IO device: give me value
  - IO device → CPU: here is the value
- Polling is OK if poll has low overhead, or has high probability of "hit"

- Key observation

- CPU and IO controllers have independent processors
  - they should be allowed to work in parallel
  - either should be able to initiate data transfer to/from memory
  - either should be able to signal the other to get the other's attention

- Autonomous controller operation
  - **Direct Memory Access (DMA)**
    - controller can send/read data from/to any main memory address
    - the CPU is oblivious to these transfers
    - DMA addresses and sizes are **programmed** by CPU using PIO
  - CPU interrupts
    - controller can signal the CPU
    - CPU checks for interrupts on every cycle (fast, clock-speed poll)
    - CPU jumps to controller's **Interrupt Service Routine** if it is interrupting



- Adding interrupts to SM213
  - New special-purpose CPU registers
    - **isDeviceInterrupting**: set by IO controller to signal interrupt
    - **interruptControllerID**: set by IO controller to identify interrupting device
    - **interruptVectorBase**: **interrupt-handler** jump table, initialized at boot time
  - Modified fetch-execute cycle

```

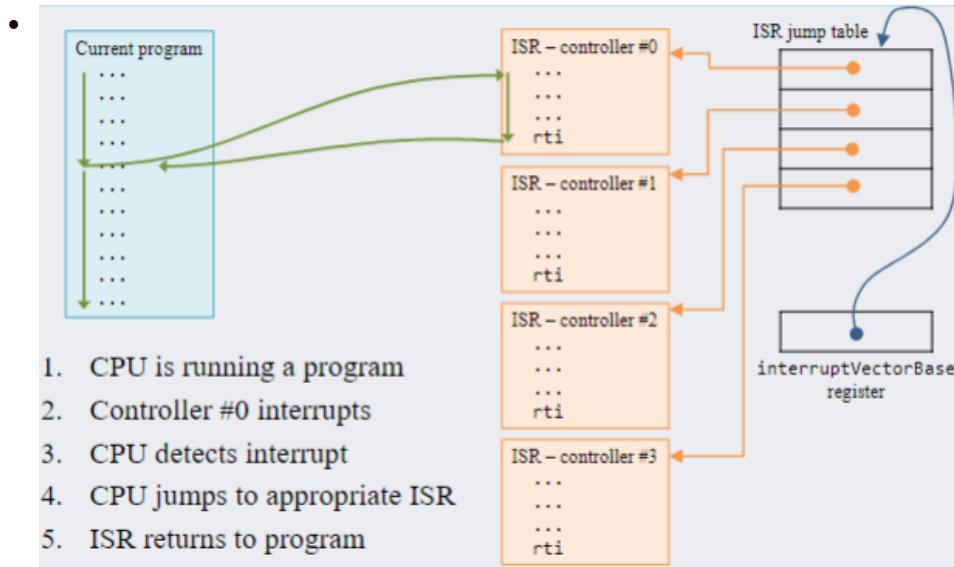
  •
    while (true) {
      if (isDeviceInterrupting) {
        r[5] ← r[5] - 4; m[r[5]] ← r[6];
        r[6] ← pc;
        pc ← interruptVectorBase[interruptControllerID];
      }
      fetch();
      execute();
    }
  
```

RTI: Return from Interrupt Instruction

$t \leftarrow r[6];$   
 $r[6] \leftarrow m[r[5]]; r[5] \leftarrow r[5] + 4;$   
 $isDeviceInterrupting \leftarrow 0;$   
 $pc \leftarrow t;$

"polls" on CPU register instead of I/O memory

- The **if**-block is same as a procedure call
- **interruptVectorBase**: contains the jump table addresses;  
**interruptControllerID**: controller number indexing
- After interruption, the program needs to return from interruption



- Architectural view of interrupts

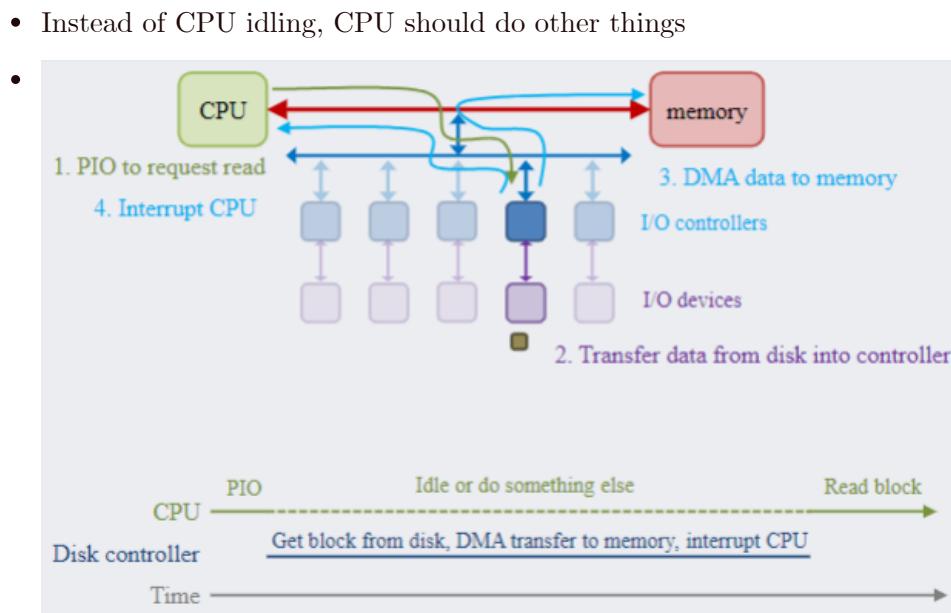
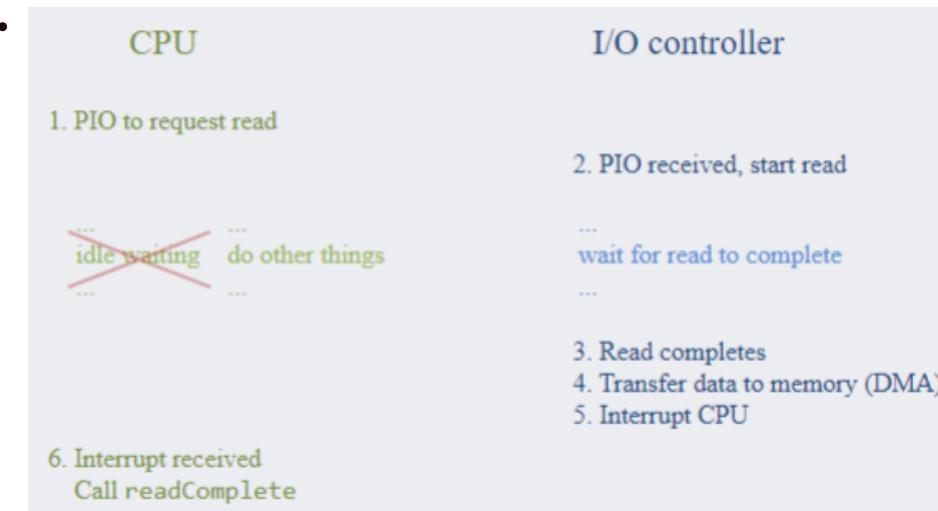


- For which of these devices does [PIO polling](#) present a better use-case than [interrupts](#)?
  - A. Devices that process user input (e.g. keyboard, mouse)
  - B. Programmable timers (e.g., execute a process one second from now)
  - C. Clock (e.g. get current time of day)
  - D. Secondary storage (e.g. hard disks)
  - E. Imaging devices (e.g. cameras, scanners)
- Correct Answer: C
- Explanation: Polling - OK for frequent, regular requests, with high probability of new data (small amounts of data);  
Interrupts - Irregular, infrequent data, large amounts of data, triggered by outside/real-world events.
- A touchscreen monitor needs to notify the CPU that the user touched on a particular position on the screen. Which option is better suited for this operation?
  - A. CPU periodically requests for the last touched position
  - B. Monitor controller uses PIO to notify CPU
  - C. Monitor controller stores location in memory, CPU periodically checks memory for new touch
  - D. Monitor controller interrupts CPU, which runs a pre-determined interrupt handler
  - E. Monitor controller changes a pre-determined register with the position of the last touch
- Correct Answer: D
- Explanation: A - polling, inefficient; B - PIO initiated by CPU, not by IO controller; C - DMA, polling memory, inefficient
- The CPU decides to turn off the NumLock indicator on the keyboard. Which option is better suited for this operation?
  - A. CPU sends a PIO request to the keyboard device
  - B. CPU changes a specific register associated with the indicator
  - C. CPU changes a specific memory location in main memory associated with the indicator
  - D. CPU interrupts the keyboard controller requesting a change in indicator values
  - E. CPU halts until the user hits the Num Lock key

- Correct Answer: A
- Explanation: A - uses PIO, but not repeatedly using PIO, only a single request; B - waste register resources; C - too indirect, inefficient; D - CPU cannot interrupt; E - impractical

## Programming with I/O

- Disk read timeline
- A disk stores blocks of data. Each block has a number. CPU can request read or write to a block.

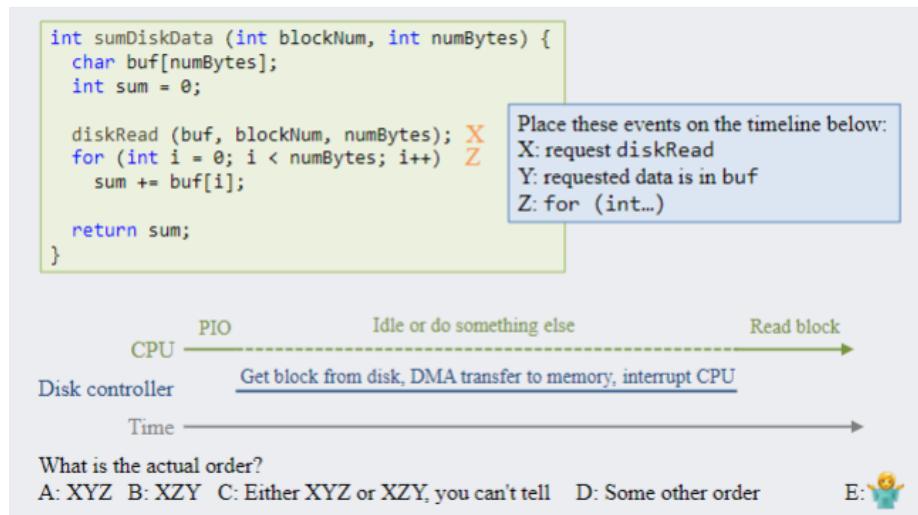


- Sequential execution
- Consider a program that reads data from a disk
  - You can think of the read as having 3 parts
    - figure out what data you want
    - get the data
    - do smth with the data you got
  - these steps must happen in this order

- we think of these steps as executing in sequence
- ```
int sumDiskData(int blockNum, int numBytes) {
    char buf[numBytes];
    int sum = 0;

    diskRead(buf, blockNum, numBytes);
    for (int i = 0; i < numBytes; i += 1) {
        sum += buf[i];
    }

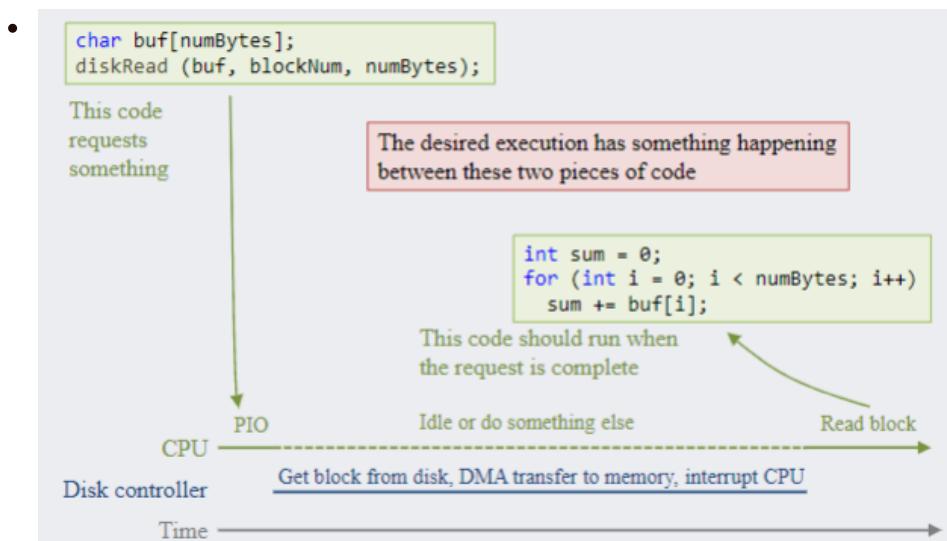
    return sum;
}
```



- Correct Answer: C
- Explanation: because `diskRead` is slow, therefore, the program could execute the `for` loop before completing `diskRead`

## Mar. 15th - Unit 2a: Cont.d

- Making code asynchronous



- Writing asynchronous code in C

- Events and event handlers
  - things that cause asynchronous code to run are called **events**
  - the code that runs when an event occurs is called an **event handler** (or **callback**)
  - handlers are *registered* to a specific event
  - events are *fired* to trigger the execution of the handler

- In code

- request and registration of event handler are listed together in code
- this code continues, does not wait for event
- handler runs asynchronously when event occurs

```
// this function is an event handler
void computeSum() {
    // do smth after disk read
}

int sumDiskBlock(blockNum, numBytes) {
    char buf[numBytes];

    diskRead(buf, blockNum, numBytes,
        computeSum);
}
```

- Disk read PIO

- Expanded disk read with PIO
  - the message sent to disk has several fields
  - each field has different device-memory address
  - access it like a struct
    - writes are to **device memory**
    - they are messages across bus to device controller; they are not writes to main memory

```
• void * diskAddress = (void *) 0x80001000;

#define DISK_OP_READ 1
#define DISK_OP_WRITE 2

struct disk_ctrl {
    int op;
    char * buf;
    int blockNum;
    int numBytes;
}
```

```

void diskRead(char * buf, int blockNum, int
numBytes, ...) {
    struct disk_ctrl * dc = diskAddress;

    // enqueue_handler (whenComplete, buf, blockNum,
    numBytes)
    // the following access use base + offset, each
    populated in IO controller
    // using one st instruction for each
    dc -> op = DISK_OP_READ;
    dc -> buf = buf;
    dc -> blockNum = blockNum;
    dc -> numBytes = numBytes;
}

```

- Implementing disk read (simplified)

- Interrupt vector, device ID, and PIO address
  - initialize before all this starts
  - by operating system when device connects

```

#define MAX_DEVICES
void (*interruptVector [MAX_DEVICES])();

int diskID      = 4;
int* diskAddress = (int*) 0x80001000;

interruptVector [diskID] = diskISR;

```

- Disk Read
  - register event handler
  - request block (PIO)
  - ```
void diskRead (char* buf, int blockNum, int numBytes
                      void (*whenComplete) (char*, int, int)) {
    enqueue_handler (whenComplete, buf, blockNum, numBytes);
    // perform PIO ... more about this later
}
```
- Interrupt handler
  - find event handler and fire event
  - firing means calling the handler procedure

```

void diskISR() {
    struct handler_dsc {
        void (*handler) (char*, int, int);
        char* buf;
        int blockNum;
        int numBytes;
    };
    struct handler_dsc hd;
    dequeue_handler (&hd);
    hd.handler(hd.buf, hd.blockNum, hd.numBytes);
}

```

```

int computeSum (char* buf, int blockNum, int numBytes) {
    int sum = 0;
    for (int i = 0; i < numBytes; i++)
        sum += buf[i];
    return sum;
}

void sumDiskBlock (blockNum, numBytes) {
    char buf [numBytes];

    diskRead (buf, blockNum, numBytes, computeSum);
}

```

What procedure calls `computeSum`?

- A. `sumDiskBlock`
  - B. the procedure that calls `sumDiskBlock`
  - C. another procedure in this program
  - D. the disk interrupt service routine
  - E. something else / I don't know
- Correct Answer: D
  - Explanation: `computeSum` (callback parameter) is called when `diskRead` is complete, which is signaled through a CPU interrupt, launching `computeSum`
  - Connecting asynchrony to program
    - How to use the value computed from the disk block
    - Suppose we want

```

void something() {
    ...
    int s = sumDiskBlock(blk, n);
    printf("%d\n", s);
}

```

- Ordering in asynchronous programs
  - If something has to happen after an event
    - it must be triggered by the event
    - often this means it must be part of (or called by) event's handler
  - To print the sum

```

void computeSumAndPrint (char* buf, int blockNum, int numBytes) {
    int sum = 0;
    for (int i = 0; i < numBytes; i++)
        sum += buf[i];
    printf("%d\n", sum);
    free (buf);
}

void sumDiskBlock (blockNum, numBytes, whenComplete) {
    char* buf = malloc (numBytes);
    diskRead (buf, blockNum, numBytes, whenComplete);
}

void something() {
    sumDiskBlock (1234, 4096, computeSumAndPrint);
}

```

Huge problem:

- often there is a ton of stuff that depend on the returned data
- these must happen after a particular event

- Improving the code

```
void computeSumAndCallback (char* buf, int blockNum, int numBytes,
                           void (*callback) (int)) {
    int sum = 0;
    for (int i = 0; i < numBytes; i++)
        sum += buf[i];
    callback(sum);
    free(buf);
}
```

```
void computeSumAndCallback (... , void (*sumCallback) (int)) {
    int sum = 0;
    for (int i = 0; i < numBytes; i++)
        sum += buf[i];
    sumCallback(sum);
    free(buf);
}

void sumDiskBlock (blockNum, numBytes, whenComplete, sumCallback) {
    char* buf = malloc (numBytes);

    diskRead (buf, blockNum, numBytes, whenComplete, sumCallback);
}

void printInt (int i) {
    printf ("%d\n", i);
}

void something() {
    sumDiskBlock(1234, 4096, computeSumAndCallback, printInt);
}
```

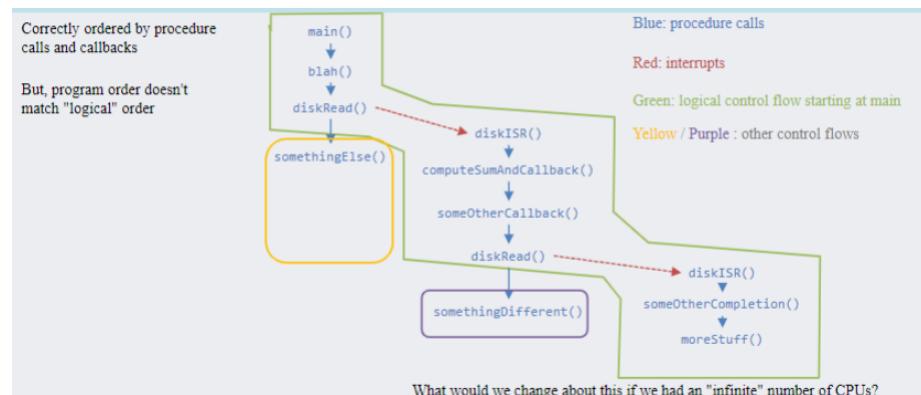
- Top-level function calls need to pass all the proper sequence
- "callback hell"
- Happy system, sad programmer
  - Humans like synchrony
    - we expect each step of a program to complete before the next one starts
    - we use the result of previous steps as input to subsequent steps
    - with disks, for example,
      - we read from a file in one step and then usually use the data we've read in the next step
  - Computer systems are asynchronous
    - the disk controller takes  $10\text{-}20$  milliseconds ( $10^{-3}\text{s}$ ) to read a block
      - CPU can execute 60 million instructions while waiting for the disk to complete one read
      - we must allow the CPU to do other work while waiting for I/O completion
    - many devices send unsolicited data at unpredictable times
      - e.g., incoming network packets, mouse clicks, keyboard-key presses
      - we must allow programs to be interrupted many, many times a second to handle these things
  - Asynchrony makes programmers sad
    - it makes programs more difficult to write and much more difficult to debug
- Possible solutions

- Accept the inevitable
  - use an event-driven programming model
    - event triggering and handling are de-coupled
  - a common idiom in many Java programs
    - GUI programming follows this model
  - CSP is a language boosts this idea to first-class status
    - no procedures or procedure calls
    - program code is decomposed into a set of sequential/synchronous processes
    - processes can fire events, which can cause other processes to run in parallel
    - each process has a guard predicate that lists events that will cause it to run
  - Javascript in web browsers and Node.js embrace asynchrony, albeit awkwardly
- Invent a new abstraction
  - an abstraction that provides programs the illusion of synchrony
  - but, what happens when
    - a program does something asynchronous, like disk read?
    - an unanticipated device event occurs?
- What's the right solution?
  - we still don't know — this is one of the most pressing questions we currently face

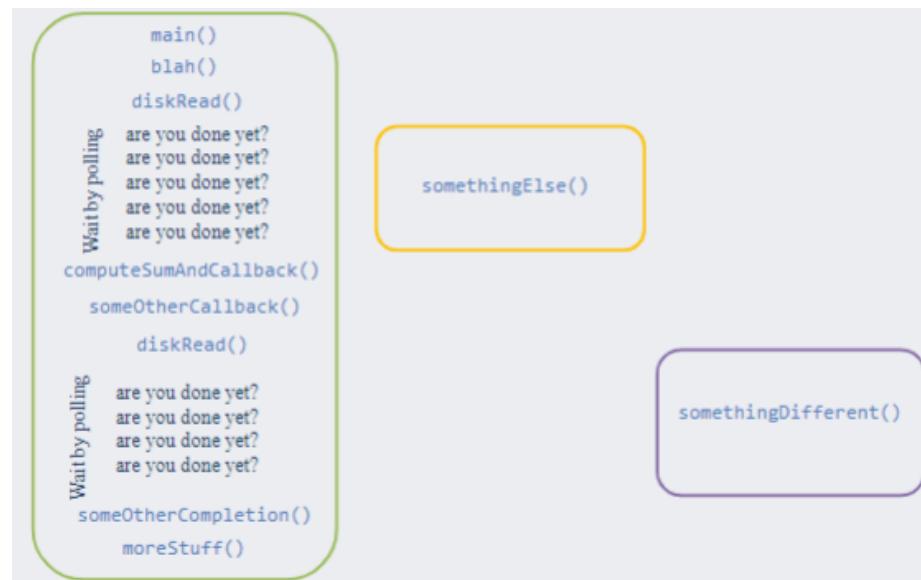
## Mar. 18th - Unit 2b: Virtual Processors

- Issues introduced by IO devices
  - Ordering of program events with IO completion
    - `diskRead` triggers IO controller to start read process
    - program has things that can only run after that process completes
  - Do other things while waiting for IO event
    - need multiple independent streams of execution in program
    - one does read and continues after it completes
    - the other does something else in the meantime
  - Asynchronous Programming
    - Order
      - callback function that is called by completion interrupt
    - Multiple Streams
      - one stream continues with return from `diskRead` without the requested block
      - the other starts with when the interrupt calls the completion callback function
  - If something needs to happen after an event, then it must be triggered by the event
    - It must be part of (or called by) the event's handler
- Problem
  - Usually we request data from IO and need it to proceed
  - Most of executable code depends on returned data

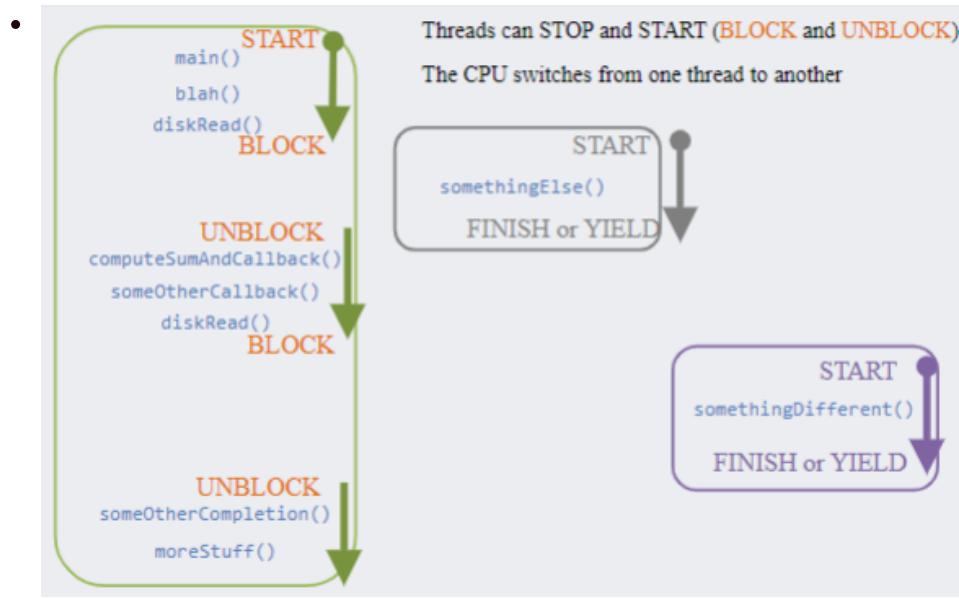
- Sometimes a sequence of requests must be made that depend on each other
- Streams of control in an asynchronous program



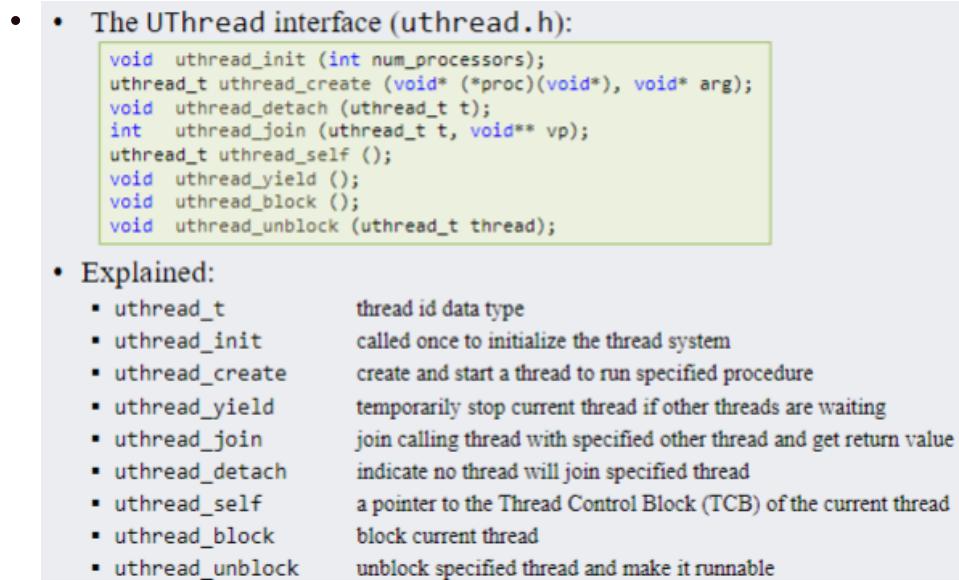
- We can keep logical flow in its own stream, if we have multiple CPUs, OK to wait at the old stream
- Poll the devices: infinite CPUs



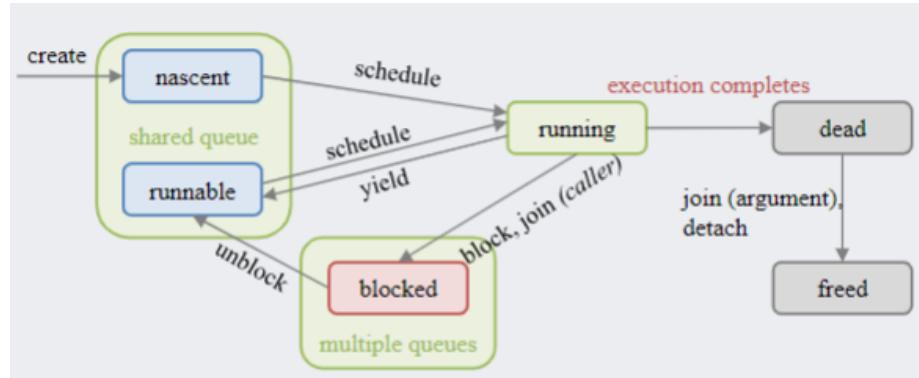
- The Virtual Processor
  - Thread
    - A single stream of synchronous execution of a program
      - the illusion of a single system
    - can be **stopped** and **restarted**
      - stopped when waiting for an event
      - restarted with the event fires
    - can co-exist with other threads sharing a single CPU
      - a scheduler multiplexes processes over processor
      - synchronization primitives are used to ensure mutual exclusion and for waiting and signaling
  - Connecting program, and logical-order with threads
    - Threads can STOP and START (or BLOCK and UNBLOCK)
    - The CPU switches from one thread to another



- UThread: a simple thread system for C



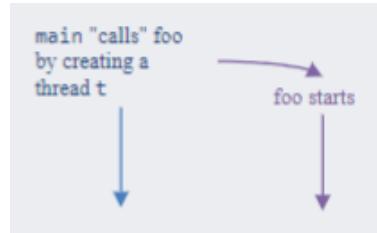
- Thread status / operations DFA



- runnable: waiting and will run when a CPU becomes available
- blocked: multiple queues, because threads can be blocked for multiple reasons, but will not become runnable until explicitly unblocked
- running - join: "I need another thread's produced result, and wait for the other thread to complete"
- dead - join: "Somebody has asked me to join them to give them my produced result"

- Thread operations - Create

- 



- Creating and starting a thread

- "Forks" the control stream

- `uthread_t t = uthread_create(foo, NULL);`

- Starts a new thread to execute a procedure

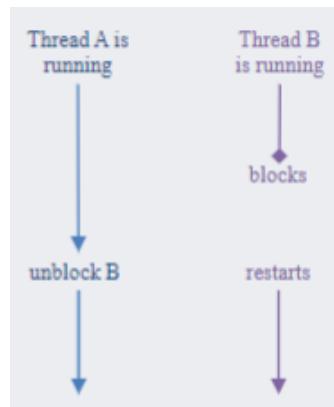
- A created thread may not start right away

- If there aren't enough CPUs available to run all threads, it will start when another thread yields (or when pre-empted)

- created (nascent) threads are considered **Runnable**

- Thread operations - Block and unblock

- 



- Blocking a thread

- Save current thread's state, stop it, switch to a different thread

- `uthread_block();`

- Thread cannot be rescheduled until explicitly unblocked  
(usually some criterion must be satisfied before the thread can be unblocked)

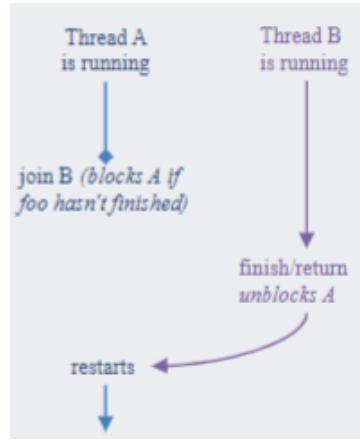
- Unblocking a thread

- Move the thread from its queue of blocked threads into the runnable queue

- `unthread_unblock(t);`

- Thread operations - join

- 

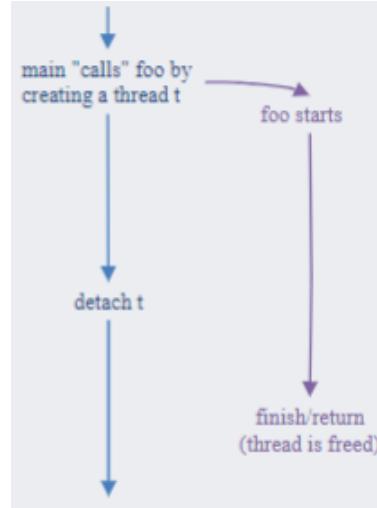


- What happens when a thread finishes
  - Some other thread may want its output
  - We need to keep its information
- Joining with a thread
  - Block current thread until a target thread completes

```
uthread_join(t, void* * rtnvaluePtr);
```

- Can obtain the return value of the target thread
- If target thread has already finished, block is unnecessary
- In effect, join may turn CPU back into a synchronous procedure call
  - But there may be other threads still running besides the joining thread and the target thread
- Thread operations - Detach

- 



- What if no thread is planning to join with it?
  - Detach the thread, notify it that it is not going to be joined
  - If the thread has already completed, free its resources
  - If it is still running, flag it to be freed as soon as it completes
- Thread operation - Yield
  - Yielding a thread
    - Save thread's state and switch to a different runnable thread

- Thread can resume as soon as CPU becomes available (another thread yields or blocks)
- Threads that yield remain runnable
- Rescheduling a thread
  - Restore a thread's state and resume running it
- Threads running on a CPU use resources
  - Registers, memory (activation frame), runtime stack
  - when a thread yields, these resources are backed up
  - when a thread resumes, these resources are restored
- Thread operation - overview
  - Create / Start
    - forks the control stream
    - like an asynchronous procedure call

`t = uthread_create(foo, NULL);`
  - Join
    - joining blocks caller until target has finished
    - join returns result of call that created thread

`uthread_join(t, rtnValuePtr);`
  - Block / Unblock      **\*\*rtnValuePtr is foo()'s return value**
    - stop and restart a thread (so that it can wait)

`uthread_block();`      **stop calling thread until it is unblocked**

`uthread_unblock(t);`      **restart thread t**

## Mar. 20th - Unit 2b: Cont.d

- Example program

```

What does the execution output look like for:
• uthread_init(1)?
• uthread_init(2)?
```

```

void* ping(void* x) {
    for (int i=0; i<NUM_ITERATIONS; i++) {
        printf("|");
        uthread_yield(); // give up CPU if a thread is waiting
    }
    return NULL;
}

void* pong(void* x) {
    for (int i=0; i<NUM_ITERATIONS; i++) {
        printf(".");
        uthread_yield(); // give up CPU if a thread is waiting
    }
    return NULL;
}
```

```

int main(int argc, char* argv[]) {
    uthread_t t0, t1;

    int i;
    uthread_init(1);

    t0 = uthread_create(ping, NULL); // create thread to run ping(NULL) and start if CPU is available
    t1 = uthread_create(pong, NULL); // create thread to run pong(NULL) and start if CPU is available

    uthread_join(t0, 0); // wait until ping thread finishes
    uthread_join(t1, 0); // wait until pong thread finishes

    printf("\n");
}
```

- If we run with 1 virtual processor

- **main**: creates thread **t0**, creates thread **t1**, (**t0, t1** in runnable queue); call **join(t0, 0)**, **main** is blocked, CPU becomes available

- **t0** starts to execute: goes into loop, print |, then yields (moves itself into runnable queue, **t1**, **t0**)
- **t1** starts to execute: goes into loop, print ., then yields (moves itself into runnable queue, **t0**, **t1**)
- **t0** will restart executing ..., then it prints alternating pattern |.|.|.
- **t0** completes executing, joins **main**; **t1** has one last iteration; **t1** is dead first, and then **main** calls join on **t1**
- With 2 (or more) virtual processors
  - Things can get complicated...

The **main** function creates two threads: **foo** and **bar**. Then, while both are running, suppose **foo** calls join on **bar**. What happens?

- A. **bar** is blocked until **foo** completes
- B. **foo** is blocked until **bar** completes
- C. **main** is blocked until **bar** completes
- D. **main** is blocked until both **foo** and **bar** complete
- E. **bar** is forced to stop immediately

- Correct Answer: B
- Revisiting the disk read
  - A program that reads a block from disk
    - Want the disk read to be synchronous

```
read(buf, siz, blkNo);           // read siz bytes
at blkNo into buf
handleData(buf, siz);           // now do
something with the block
```

- but it is actually asynchronous, so we had
 

```
asyncRead(buf, siz, blkNo, handleData);
doSomethingElse();
```
- As a timeline
  - two processors, two separate computations
- "Synchronous" disk read using threads
  - Create two threads that CPU runs, one at a time
    - One for disk read
    - one for doSomethingElse
  - Illusion of synchrony
    - disk read blocks while waiting for disk to complete
    - CPU runs other threads while first thread is blocked
    - disk interrupt restarts the blocked thread

- ```

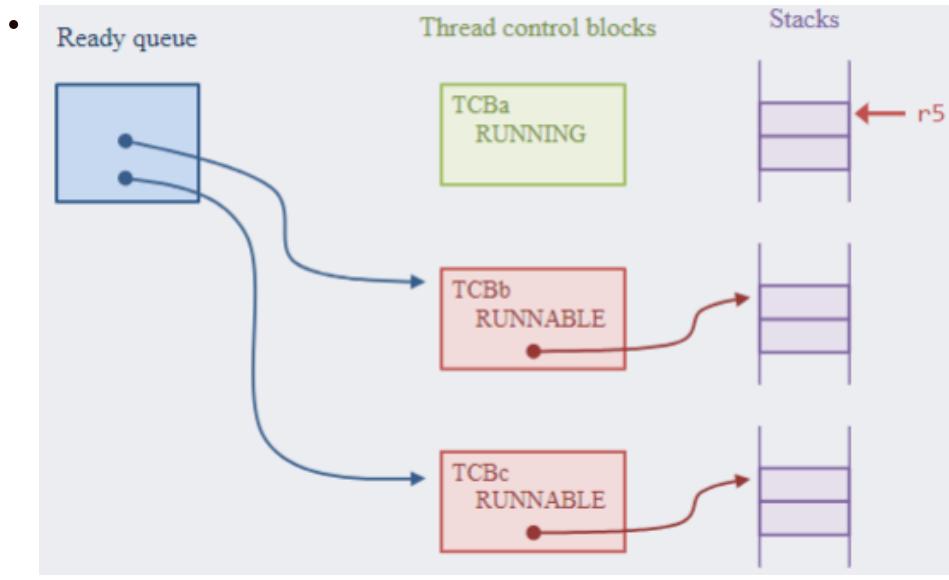
    •
        asyncRead (buf, siz, blkNo);
        blockWaitingForInterrupt(); // let another thread run
        handleData (buf, siz);
    
```
- ```

    interruptHandler() {
        unblockWaitingThread();
    }
    
```
- Now we can do this:

<pre> struct PendingRead {     ...     uthread_t waitingThread; } pendingRead; </pre>	<pre> void foo() {     int sum = sumDiskData(1234, 4096);     doSomethingWithSum(sum); } </pre>
<pre> void diskRead(char* buf, int size, int blkNo) {     pendingRead.waitingThread = uthread_self();     asyncread(buf, size, blkNo, uthread_unblock);     uthread_block(); } </pre>	<pre> void main() {     uthread_create(doSomethingElse, 0);     foo(); } </pre>
<pre> int sumDiskData(int blkNo, int nBytes) {     char buf[nBytes];     int sum = 0;     diskRead(buf, nBytes, blkno);     for (int i = 0; i &lt; nBytes; i++) sum += buf[i];     return sum; } </pre>	<p>In disk interrupt service routine:</p> <pre> uthread_unblock(pendingRead.waitingThread); </pre>

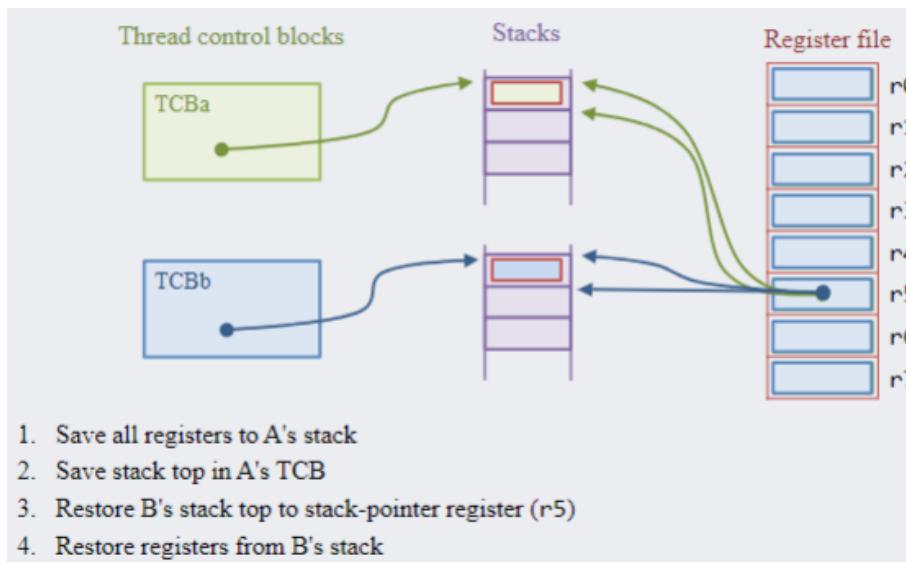
- Implementing threads

- Key concept in implementation: **blocking** and **unblocking**
    - Stop a thread: no longer executing on a CPU
    - What happens to the thread?
    - What happens to the physical processor? CPU has hardware resources for executing processes (registers)
  - Thread management
    - Data structures and operations?
  - Thread state
    - when running: register file (CPU) and runtime stack (in physical memory, activation frames for sequence of called functions)
    - when stopped: thread-control block (TCB) object and runtime stack
  - Thread-control block
    - thread status
    - saved value of thread's stack pointer if it is not running
    - scheduling parameters (priority, quantum, preemptibility)
  - Ready queue
    - list of TCBs of all RUNNABLE threads
  - Blocked queues
    - list of TCBs of BLOCKED threads
    - May be more than one queue, depending on reason for being blocked



## Mar. 22nd - Unit 2b: Cont.d

- Implementing thread yield
  - Thread yield
    - puts current thread on ready queue
    - gets next runnable thread from ready queue
    - switches to next thread
  - `void uthread_yield() {  
 ready_queue_enqueue(uthread_self());  
 uthread_t to_thread = ready_queue_dequeue();  
 assert(to_thread);  
 uthread_switch(to_thread, TS_RUNNABLE);  
}`
- Implementing thread switch
  - Goal
    - Implement a procedure `switch(Ta, Tb)` that stops T<sub>a</sub> and starts T<sub>b</sub>
    - T<sub>a</sub> calls `switch`, but it returns to T<sub>b</sub>
  - Requires
    - saving T<sub>a</sub>'s processor state and setting processor state to T<sub>b</sub>'s saved state
    - state is just registers, which can be saved/restored to/from stack
    - thread-control block has pointer to stack pointer for each thread
  - Implementation
    - save all registers to stack

- save stack pointer to  $T_a$ 's TCB
- set stack pointer to stack pointer  $T_b$ 's TCB
- restore registers from stack
- 
  1. Save all registers to A's stack
  2. Save stack top in A's TCB
  3. Restore B's stack top to stack-pointer register (r5)
  4. Restore registers from B's stack
- Example code
 

• C / x86 assembly

```

asm volatile (
    "pushq %%rbx"
    "pushq %%rcx"
    "pushq %%rdx"
    "pushq %%rsi"
    "pushq %%rdi"
    "pushq %%rbp"
    "pushq %%r8"
    "pushq %%r9"
    "pushq %%r10"
    "pushq %%r11"
    "pushq %%r12"
    "pushq %%r13"
    "pushq %%r14"
    "pushq %%r15"
    "pushfa"
    "movq %%rsp, %0"
    "movq %1, %%rsp"
)
    : "=m" (*from_sp_p)
    : "ra" (to_sp));

```

from\_tcb->saved\_sp ← r[sp]  
r[sp] ← to\_tcb->saved\_sp

The diagram shows the assembly code with annotations:

  - Backing up CPU state of thread A:** Brackets group the first 15 pushq instructions and the pushfa instruction.
  - restoring CPU state of thread B:** Brackets group the popfq instruction followed by the movq and movq instructions.
  - updating stack pointers:** Brackets group the two movq instructions at the bottom of the code.
- The `uthread_switch` procedure saves the `from` thread's registers to the stack, switches to the `to` thread's stack and restores its registers from the stack. But, what does it do with the program counter (PC)?
  - A. It saves the `from` thread's PC to the stack and restores the `to` thread's PC from the stack.
  - B. It saves the `from` thread's PC to its thread control block.
  - C. Nothing. It does not need to change the PC because the `from` and `to` threads' PCs are already saved on the stack before `switch` is called.
  - D. It jumps to the `to` thread's PC value.
  - E. ???
- - Correct Answer: C
  - Explanation:

```

// thread A
void ping() {
    // ...
    uthread_yield();      // needs to call to
    other addresses
                           // return address
                           thus stored on activation
                           // frame
}

// thread B
void pong() {
    // ...
    uthread_yield();
}

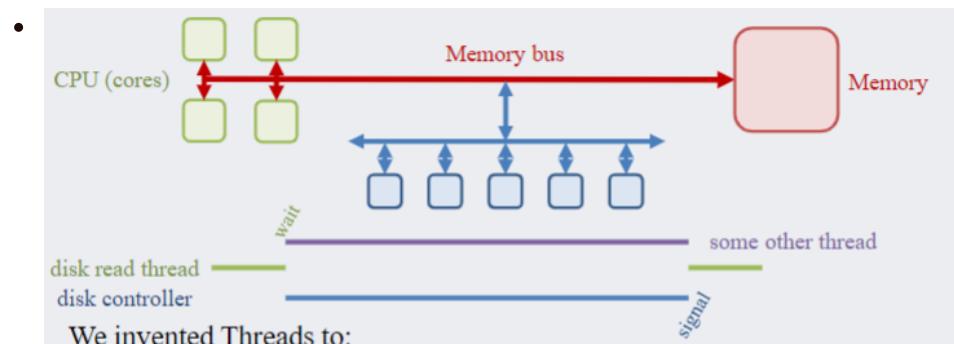
```

- **yield**'s callee prologue save return address to **ping** on its activation frame. Then **yield** calls **switch**, A's stack is switched to B's stack; thus similarly, when B yields, the return address for **pong** is stored on its activation frame
- Thread scheduling
  - Thread scheduling is
    - the process of deciding when threads should run
      - when there are more runnable threads than processors
      - involves a **policy** and a **mechanism**
  - Thread scheduling policy
    - the set of **rules** that determines which threads should be running
  - Things to consider
    - Do some threads have higher **priority** than others?
    - Should threads get **fair access** to the processor?
    - Should threads be guaranteed to **make progress**?
    - Should one thread be able to **pre-empt** another?
- Priority round-robin scheduling policy
  - Priority: number assigned to each thread
    - thread with highest priority goes first
  - When choosing the next thread to run
    - run the highest priority runnable thread
    - when the threads have the same priority, run thread that has waited the longest

- Implementation (mechanism)
  - organize Ready (runnable) Queue as a priority queue
    - highest priority first
    - FIFO among threads of equal **maximum** priority
- Pre-emption
  - This occurs when
    - a "yield" is forced upon the current running thread
    - current thread is stopped to allow another thread to run
  - Priority-based pre-emption
    - when a thread is made runnable (e.g. created or unblocked)
    - if it is higher priority than current-running thread, it pre-empts that thread
  - Quantum-based pre-emption (time-based)
    - each thread is assigned a runtime "quantum" (guaranteed minimum running time on CPU)
    - thread is pre-empted at the end of its quantum
    - How long should quantum be?
      - Disadvantage of too short? Thread may not make much progress
      - Disadvantage of too long? Other threads may result waiting for too long
      - Depends on OS configuration (typically  $\sim 100\text{ms}$ )
- Implementing quantum pre-emption
  - The problem
    - when application thread(s) are running, nothing is watching over them
    - for the system scheduler to control things, it needs a CPU to run on
    - if the application thread is running, the system is not
  - Solution: interrupts (timer device)
    - an IO controller connected to a clock/timer device
    - interrupts processor at regular intervals
  - Timer interrupt handler
    - compares the running time of the current thread to its quantum
    - pre-empts it if quantum has expired

## Unit 2c: Synchronization

- Synchronization



- We invented threads to: **express parallelism** - do things at the same time on different processors

**manage synchrony** - do something else while waiting for IO controller

- Two problems

- Coordinating access to memory (variables) shared among multiple threads
- control flow transfers among threads (wait until notified by another thread)
- **Synchronization** is the mechanism threads use to
  - ensure **mutual exclusion** of critical sections
  - wait for and signal of the occurrence of events

## Mar. 25th - Unit 2c: Cont.d

- Communicating through shared data
  - There will be problems
    - threads **share a data structure**
    - operations involve **multiple memory accesses**
    - these accesses can be **arbitrarily interleaved**
  - Example: a stack implemented as an array

- ```
int n;
int array[SIZE];

void push(i) {
    if (n < SIZE) {
        array[n] = i;
        n += 1;
    }
}

int pop() {
```

```

    if (n > 0) {
        n -= 1;
        return array[n];
    } else {
        return -1;
    }
}

```

- Suppose two threads accessing concurrently
- What happens if both push? both pop? 1 push, 1 pop?
- A linked-list stack

- ```

struct SE {
    struct SE * next;
};

struct SE * top = NULL;

void push_st (struct SE * e) {
    e -> next = top;
    top = e;
}

void SE * pop_st() {
    struct SE * e = top;
    top = (top) ? top -> next : NULL;
    return e;
}

```

- Sequential/synchronous execution of repeated push/pop is OK  
(not running with any threads)

```

void push_driver(long int n) {
    struct SE * e;
    while (n -= 1) {
        push(malloc(...));
    }
}

void pop_driver(long int n) {
    struct SE * e;
    while (n -= 1) {
        do {
            e = pop();
        } while (!e)
    }
    free(e);
}

```

```

push_driver(n);
pop_driver(n);
assert(top == NULL);

```

- but concurrent execution does not always work

```

int main(void) {
    et = uthread_create(push_driver, num);
    dt = uthread_create(pop_driver, num);

    uthread_join(et, 0);
    uthread_join(dt, 0);

    assert(top == NULL);
}

```

- works for `uthread_init(1)` without pre-emption, but not guaranteed to work for `uthread_init(2)`
- The problem
  - Same situation as the array implementation
    - push and pop are critical sections on the shared stack
    - parallel execution leads to potential arbitrary interleaving of operations
    - sometimes, the interleaving corrupts the data structure
- The importance of mutual exclusion
  - Shared data
    - data structure that could be accessed by multiple threads
    - typically, concurrent access to shared data is a bug
  - Critical sections
    - sections of code that access shared data
  - Race condition
    - simultaneous access to critical section by multiple threads (at least one updating)
    - conflicting operations on shared data structure are arbitrarily interleaved
    - unpredictable (non-deterministic) program behaviour - usually a (serious) bug
  - Mutual exclusion
    - a mechanism implemented in software (with some special hardware support)

- to ensure critical sections of a shared data item are executed by only one thread at a time
- reading and writing should be handled differently
- Mutual exclusion using locks
  - lock semantics
    - a lock is either **held** by a thread, or available
    - at most one thread can hold a lock at a time
    - a thread attempting to acquire a lock that is already held is forced to wait
  - lock operations (primitives)
    - **lock**: acquire lock, wait if necessary
    - **unlock**: release lock, allowing another thread to acquire if waiting
  - using locks on the shared stack

```

void push_cs(struct SE * e) {
    lock(&aLock);
    push_et(e);
    unlock(&aLock);
}

struct SE * pop_cs() {
    struct SE * e;
    lock(&aLock);
    e = pop_st();
    unlock(&aLock);
    return e;
}

```

- Questions with locks
  - Thread A calls **push\_cs** while Thread B is executing in **push\_cs**
    - B holds the lock, A must wait until B releases the lock
  - Thread A calls **push\_cs** while Thread B is executing in **pop\_cs**
    - B holds the lock, A must wait until B releases the lock
  - What if **push\_cs** and **pop\_cs** use different locks
    - Suppose A acquires "red" lock, B acquires "blue" lock and may end up with interleaving
  - What if **push\_cs** or **pop\_cs** never call unlock
    - Waiting threads can never resume
- Implementing simple locks

- An initial attempt
  - use a shared global variable for synchronization
  - **lock**: loops until the variable is 0 and then sets it to 1
  - **unlock**: sets the variable to 0

- ```
void lock(int * lock) {
    while (*lock == 1) {      }
    *lock = 1;
}

void unlock(int * lock) {
    *lock = 0;
}
```

- Problem: There is a race in the lock code
  - Suppose two threads both request a lock at (nearly) the same time

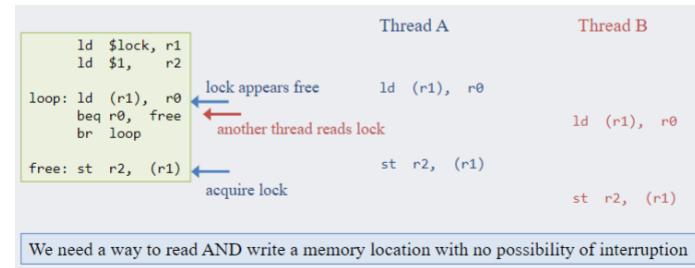
- ```
// Thread A
void lock(int * lock) {
    while (*lock == 1) {      }
    *lock = 1;
}

// Thread B
void lock(int * lock) {
    while (*lock == 1) {      }
    *lock = 1;
}
```

- Suppose the execution occurs as follows:
  1. read `*lock == 0`, exit loop
  2. read `*lock == 0`, exit loop
  3. `*lock = 1`
  4. return with lock held
  5. `*lock = 1`
  6. return with lock held

Both threads think they hold the lock

- The race exists even at the machine code level
  - two instructions are needed to acquire the lock
    - first, read to check that the lock is available
    - second, set the lock to held
  - but a read by another thread can check the lock before the first thread sets the hold



- Atomic memory exchange

- Atomicity
  - is a general property in systems
  - where a group of operations are performed as a single, indivisible unit
- Atomic memory exchange
  - one type of atomic memory instruction
  - group a load and store together atomically
    - no intervening access from other threads
  - exchanging the value of a register and a memory location

Name	Semantics	Assembly
<i>atomic exchange</i>	$r[v] \leftarrow m[r[a]]$ $m[r[a]] \leftarrow r[v]$	xchg (ra), rv

- Implementation
  - Cannot be implemented by CPU alone
    - must synchronize across multiple CPUs
      - accessing the same memory location at the same time
  - Implemented by memory bus
    - memory bus synchronizes every CPU's access to memory
    - the two parts of the exchange (read + write) are coupled on bus
    - bus ensures that no other memory transaction can intervene
    - Higher overhead, slower than normal load and store instructions

## Mar. 27th - Unit 2c: Cont.d

- Spinlock
  - A lock where the waiter **spins**, looping on memory reads until lock is acquired, also called a **busy-waiting** lock
  - Implementation using atomic exchange

- Spin on atomic memory operation
  - That attempts to acquire lock while atomically reading its old value
- ```
ld $lock, r1          # get address of lock
variable
ld $1, r0             # r0 = 1
loop:
    xchg (r1), r0
    beq r0, held      # if (r0 = 0), lock is held
    br loop           # else lock is held by others, spin to wait
held:
    ...

```
- Speeding up spinlocks
  - Spin first on normal read
    - normal reads are fast and efficient compared to exchange
    - use normal read in loop until lock appears free
    - when lock appears free, use exchange to try to grab it
    - if exchange fails, then go back to normal read
  - ```
ld $lock, r1          # r1 = &lock
loop:   ld (r1), r0     # r0 = lock
        beq r0, try      # if (r0 == 0), has a chance
        br loop           # goto try
try:    ld $1, r0         # r0 = 1
        xchg (r1), r0    # atomically swap r0 and lock
        beq r0, held      # goto held, if (r0 == 0), actually obtained
        br loop           # try again if another thread holds lock
held:   ...

```
- Busy-waiting pros and cons
  - Spinlocks are necessary and OK if spinner only waits a short time
  - But, using spinlock to wait for a long time wastes CPU cycles (if running on a single CPU, other threads may not be able to progress)
- Blocking locks

- If a thread might wait for a long time, it should block so that other threads can run
- It will then **unblock** when it becomes runnable, when either the lock is unlocked, or an event is signaled
- Blocking locks for **mutual exclusion**
  - attempting to acquire a held lock blocks calling thread
    - blocked thread's TCB is stored on lock's waiting queue
  - when releasing lock, unblock a waiting thread if there is one
    - remove blocked thread from lock's waiting queue and place it on ready queue
- Blocking for **event notification**
  - wait by blocking, placing TCB on a waiting queue for event
  - signal a specific waiting queue by moving a thread to ready queue
- Spinlocks vs. Blocking locks
  - Spinlocks
    - uncontended lock has low overhead, thread remains on CPU, execute loop a few times
    - waiting for lock has high overhead and wastes resources, thread still remains on CPU
    - used for short expected waits with minimal contention
    - used in implementation of blocking locks
  - Blocking locks
    - lock has fixed overhead, like thread switch, there is backup & restore CPU registers
    - if uncontended, overhead is higher than spinlocks
    - if contended, less waste of CPU resources
    - used when lock may be held for long time, or with high contention
- Monitors (Mutexes) and condition variables
  - MUTual EXclusion plus inter-thread synchronization
    - basic for synchronization primitives in Unix, Java
  - Monitor / Mutex
    - blocking lock to guarantee mutual exclusion
    - Basic operations: **lock** and **unlock** (sometimes called **enter** and **exit**)
  - Condition variable
    - allows threads to synchronize with each other

- **wait**: blocks until another thread signals on the variable
- **signal**: unblocks a waiting thread, if one exists
- **broadcast**: unblocks all threads currently waiting
- can only be accessed from inside of a monitor
- ```

struct uthread_mutex;
typedef struct uthread_mutex* uthread_mutex_t;
struct uthread_cond;
typedef struct uthread_cond* uthread_cond_t;

uthread_mutex_t uthread_mutex_create          ();
void uthread_mutex_lock                      (uthread_mutex_t);
void uthread_mutex_lock_READONLY          (uthread_mutex_t);
void uthread_mutex_unlock                     (uthread_mutex_t);
void uthread_mutex_destroy                   (uthread_mutex_t);

uthread_cond_t uthread_cond_create          (uthread_mutex_t);
void uthread_cond_wait                      (uthread_cond_t);
void uthread_cond_signal                    (uthread_cond_t);
void uthread_cond_broadcast                 (uthread_cond_t);
void uthread_cond_destroy                  (uthread_cond_t);
```

- Using conditions

- Basic formulation
- one thread acquires mutex and may wait for a condition to be established

```

uthread_mutex_lock(aMutex);
while (!aDesiredState) {
    uthread_cond_wait(aCond);
}
doSomethingUseful();
uthread_mutex_unlock(aMutex);
```

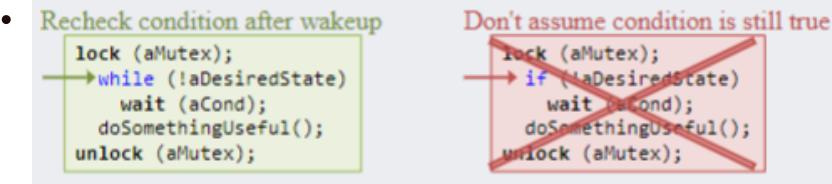
- another thread acquires mutex, establishes condition and signals waiter, if there is one

```

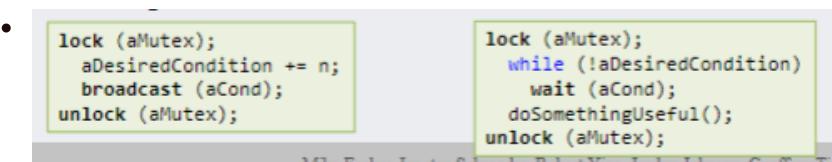
uthread_mutex_lock(aMutex);
aDesiredState = 1;
uthread_cond_signal(aCond);
uthread_mutex_unblock(aMutex);
```

- **wait** releases the mutex and blocks thread
  - before waiter blocks, it automatically releases mutex to allow other threads to acquire it
  - when wait unblocks, it re-acquires mutex, waiting/blocking to enter if necessary
  - note: other threads may have acquired mutex between wait call and return
- **signal** awakens at most one thread

- waiter does not run until signaler releases the mutex explicitly
- a third thread could intervene and acquire mutex before waiter
- waiter must then **re-check** wait condition
- if no threads are waiting, then calling signal has no effect



- **broadcast** awakens all waiting threads
  - may wake up too many
  - that's OK (but wasteful) since threads re-check wait condition and re-wait if necessary



## Apr 3rd - Unit 2c: Cont.d

- Video Streaming with Threads
  - General problem
    - Video playback has 2 parts: fetch/decode, and playback
    - fetch has variable latency, and so we need a **buffer**
      - sometimes you can fetch faster than playback rate
      - but, other times there are long delays
      - buffer hides the delays by fetching ahead of playback position when possible
  - **Bounded buffers** and two independent threads
    - finite buffer of the next few video frames to play
    - maximum size is  $N$  frames
    - goal: keep buffer at least 50% full (or some other threshold of our choosing)
  - Producer thread
    - fetch frame from network and put them in buffer
  - Consumer thread
    - fetch frame from buffer, decode them and send them to video driver
- How are Producer and Consumer connected
  - advantage of this approach is that they are largely decoupled; each has a separate job

- but, it's the consumer that decides when the producer should run
- Step 1 - Request frames

```

struct video_frame;
#define N 100
struct video_frame buf[N];
int buf_length = 0;
int buf_pcur = 0;
int buf_ccur = 0;

void producer() {
    while (1) {
        uthread_lock (mx);
        while (buf_length < N) {
            buf [pcur] = get_next_frame();
            buf_pcur = (pcur + 1) % N;
            buf_length += 1;
        }
        uthread_cond_wait (need_frames);
        uthread_unlock (mx);
    }
}

void consumer() {
    while (1) {
        uthread_lock (mx);
        assert (buf_length > 0);
        show_frame (buf [buf_ccur]);
        buf_ccur = (buf_ccur + 1) % N;
        buf_length -= 1;
        if (buf_length == N/2)
            uthread_cond_signal (need_frames);
        uthread_unlock (mx);
    }
}

```

- Producer update buffer until full
- If buffer drops below threshold, Consumer asks Producer to update more frames
- Step 2 - Delivering frames

```

void producer() {
    uthread_lock (mx);
    while (1) {
        while (buf_length < N) {
            buf [pcur] = get_next_frame();
            buf_pcur = (pcur + 1) % N;
            buf_length += 1;
            uthread_cond_signal (have_frame);
        }
        uthread_cond_wait (need_frames);
    }
    uthread_unlock (mx);
}

void consumer() {
    uthread_lock (mx);
    while (1) {
        while (buf_length == 0)
            uthread_cond_wait (have_frame);
        show_frame (buf [buf_ccur]);
        buf_ccur = (buf_ccur + 1) % N;
        buf_length -= 1;
        if (buf_length < N/2)
            uthread_cond_signal (need_frames);
    }
    uthread_unlock (mx);
}

```

- **uthread\_cond\_signal(have\_frame)**: have at least one frame in the buffer
- **while (buf\_length == 0)**: while buffer is empty, stop and wait for signal

- Full version

```

void producer() {
    uthread_lock (mx);
    while (1) {
        while (buf_length < N) {
            buf [pcur] = get_next_frame();
            buf_pcur = (pcur + 1) % N;
            buf_length += 1;
            uthread_cond_signal (have_frame);
        }
        uthread_cond_wait (need_frames);
    }
    uthread_unlock (mx);
}

void consumer() {
    uthread_lock (mx);
    while (1) {
        uthread_cond_wait (show_next_frame);
        while (buf_length == 0);
            uthread_cond_wait (have_frame);
        show_frame (buf [buf_ccur]);
        buf_ccur = (buf_ccur + 1) % N;
        buf_length -= 1;
        if (buf_length < N/2)
            uthread_cond_signal (need_frames);
    }
    uthread_unlock (mx);
}

```

- **show\_next\_frame**: will be signaled every time a new frame is required for the video driver; every  $\frac{1}{30}$ s
- **uthread\_mutex\_t mx;**  
**uthread\_cond\_t need\_frames;**  
**uthread\_cond\_t have\_frame;**  
**uthread\_cond\_t show\_next\_frame;**
- Producer and Consumer Threads

- General Template

```
int canGoFlag = 1;
uthread_mutex_t mx;
uthread_cond_t canGoCond;
```

Producer (P)

```
uthread_mutex_lock(mx);
uthread_cond_signal(canGoCond);
canGoFlag = 1;
uthread_mutex_unlock(mx);
```

Consumer (C)

```
uthread_mutex_lock(mx);
while (canGoFlag == 0)
    uthread_cond_wait(canGoCond);
canGoFlag = 0;
uthread_mutex_unlock(mx);
```

- Key invariant: C does not complete until it is able to change `canGoFlag` from 0 to 1

- if `canGoFlag` is 0, it waits for P to set it to 1
- This waiting behavior is achieved by C waiting on (and P signaling) the same condition

```
int canGoFlag = 1;
uthread_mutex_t mx;
uthread_cond_t canGoCond;
```

```
void uthread_cond_wait (uthread_cond_t cond) {
    assert (cond->mutex->holder == uthread_self ());
    uthread_enqueue (&cond->waiter_queue, uthread_self ());
    uthread_mutex_unlock (cond->mutex);
    uthread_block ();
    uthread_mutex_lock (cond->mutex);
}
```

Producer (P)

```
uthread_mutex_lock(mx);
uthread_cond_signal(canGoCond);
canGoFlag = 1;
uthread_mutex_unlock(mx);
```

Consumer (C)

```
uthread_mutex_lock(mx);
while (canGoFlag == 0)
    uthread_cond_wait(canGoCond);
canGoFlag = 0;
uthread_mutex_unlock(mx);
```

What would happen if `cond_wait` didn't unlock and lock as it does?

- nothing; it's fine either way
- C might return when it shouldn't
- P might not be able to signal it C correctly in some cases
- P will never be able to successfully signal C under any circumstances
- The instructor sees you playing games during the lecture

- Correct Answer: D
- Explanation: C → acquires mutex, checks condition (succeeds, execute `while` loop) → wait, block (while holding mutex); P → attempts to acquire mutex, fails (C holds) → unable to run, end up waiting and blocked

```
int canGoFlag = 1;
uthread_mutex_t mx;
uthread_cond_t canGoCond;
```

Producer (P)

```
uthread_mutex_lock(mx);
uthread_cond_signal(canGoCond);
canGoFlag = 1;
uthread_mutex_unlock(mx);
```

Consumer (C)

```
uthread_mutex_lock(mx);
while (canGoFlag == 0)
    uthread_cond_wait(canGoCond);
canGoFlag = 0;
uthread_mutex_unlock(mx);
```

Why not this? (CX)

```
uthread_mutex_lock(mx);
if (canGoFlag == 0)
    uthread_cond_wait(canGoCond);
canGoFlag = 0;
uthread_mutex_unlock(mx);
```

Suppose P and CX are the only threads accessing the mutex.  
Which statement is correct?

- CX always works
- CX never works
- CX works if and only if there is a single P thread
- CX works if and only if there is a single CX thread
- Churros are tasty

- Correct Answer: D

- Explanation:

CX → acquires mutex, check condition, executes `while` loop → wait, release mutex

P1 → acquire mutex, set condition, signal → release mutex

CX → attempt to acquire mutex, fails, wait, go back to sleep

P2 → acquires mutex, set condition, signal → release mutex

- CX → wake up, proceed as if condition passes → set condition to 0 → release mutex
  - In D, shared data may become corrupted if there are multiple (CX) consumer threads
- CX1 → acquire mutex, check condition, wait on condition → release mutex
- P → acquire mutex, set condition, signal condition → release mutex
- CX1 and CX2 may have a race to acquire the mutex; assume CX2 obtain mutex first
- CX1 → attempt to acquire mutex (failed), block, on mutex queue
- CX2 → acquires mutex, check condition, set `canGoFlag` to 0 → release mutex
- CX1 → acquires mutex, complete `wait`, proceeds **as if** the `if` condition is true, set `canGoFlag` = 0 → release mutex

## Apr. 5th - Unit 2c: Cont.d

- Beer for everyone
  - Beer pitcher is a shared data structure with these operations
    - `pour()`: pours one glass from pitcher (reduces beer left by 1)
    - `refill()`: adds more beer to pitcher (increases beer left by  $N$ )
  - Implementation goal
    - synchronize access to pitcher
    - pouring from empty pitcher requires waiting for it to be filled
    - filling pitcher releases waiting threads
  - We will use
    - `glasses`: amount of beer left in the pitcher, in glasses (`int`)
    - `mx`: mutex
    - `hasBeer`: condition indicating there is at least one beer left
  - Setup

```

struct BeerPitcher {
    int glasses;
    uthread_mutex_t mx;
    uthread_cond_t hasBeer;
};

struct BeerPitcher* createAndInitializePitcher() {
    struct BeerPitcher* p = malloc(sizeof(struct
Beerpitcher));

```

```

    p -> glasses = 0;
    p -> mx = uthread_mutex_create();
    p -> hasBeer = uthread_cond_create(p -> mx);
    return p;
}

// Pouring a glass
void pour(struct BeerPitcher* p) {
    uthread_mutex_lock(p -> mx);
    while (p -> glasses == 0) {
        // similar to streaming: block playback
when buffer is empty
        uthread_cond_wait(p -> hasBeer);
    }
    p -> glasses = 1;
    uthread_mutex_unlock(p -> mx);
}

// Refilling the pitcher (pitcher has unlimited
capacity)
void refill(struct BeerPitcher* p, int n) {
    uthread_mutex_lock(p -> mx);
    p -> glasses += n;
    for(int i = 0; i < n; i += 1) {
        uthread_cond_signal(p -> hasBeer);
        // or uthread_cond_broadcast, if most
waiters will be woken up
    }
    uthread_mutex_unlock(p -> mx);
}

```

- Signal and Mutex Race

Thread A: pour; Thread B: refill; Thread C: pour

Start with `glasses == 0`

A → acquires mutex, checks condition → waits, release mutex

B → acquires mutex, sets `glasses += 1`, signals condition → releases mutex

A and C race to get mutex, A fails

C → acquires mutex, sets `glasses -= 1` → releases mutex

A → acquires mutex, checks condition → waits, releases mutex

- What if we want to refill automatically?
  - a pitcher has capacity `maxGlasses` and current volume `glasses`
  - pouring removes one glass if there is enough beer, and waits otherwise

- refilling loops forever, waiting for pitcher to be empty, and when it is, it refills the pitcher to its full capacity and awakens any pourer

- ```
// add uthread_cond_t isEmpty =
uthread_cond_create(p -> mx);
```
- ```
void pour(...) {
    ...
    if (p -> glasses == 0) {
        // similar to streaming: buffer is
        // empty, make Producer get more frames
        uthread_cond_signal(p -> isEmpty);
    }
}

void refill(...) {
    ...
    while(1) {
        while(p -> glasses > 0) {
            uthread_cond_wait(p -> isEmpty);
        }

        // similar to streaming: fill buffer,
        // signal Consumers that frames available
        p -> glasses += p -> maxGlasses;
        for (int i = 0; i < maxGlasses; i +=
1) {
            uthread_cond_signal(p ->
hasBeer);
        }
    }
    ...
}
```

- Back to disk reads

- Suppose we write an async disk read, and we want to block/wait with conditions

```
void read(char * buf, int bytes, int blockno) {
    // some thread can intervene here, ISR may occur
    // before read thread resumes
    scheduleRead(buf, nbytes, blockno);

    uthread_mutex_lock(mx);
    uthread_cond_wait(disk_op_complete);
    uthread_mutex_unlock(mx);
}
```

```

void readCompleteCalledByISR() {
    uthread_mutex_lock(mx);
    uthread_signal(disk_op_complete);
    uthread_mutex_unlock(mx);
}

```

- Wait-signal problem
  - wait condition check / trigger and wait are not atomic
  - signal could occur before wait, thus waiter could miss signal and never wake up
- Solution
  - Ensure that condition check / trigger and wait are atomic
  - So that wait is ordered before signal
- Naked notify
  - Naked signal / notify
    - A signal called outside of a monitor
    - Should be avoided (can cause wait-signal race)
      - signal usually needed due to modifying shared data, if outside of mutex, shared data can also be modified outside
  - Sometimes, it is necessary
    - when blocking is not allowed
    - e.g. in some asynchronous event handlers (IO operations are slow)
      - if an IO controller holds a mutex, then CPU must wait a very long time to acquire mutex
- Reader-writer monitors
  - Critical sections could be classified as
    - **reader**: if only reads the shared data
    - **writer**: if it updates the shared data
    - we can weaken the mutual exclusion constraint
      - writer requires an **exclusive access** to the monitor
      - but a group of readers can **access monitor concurrently**
  - Read-write monitors
    - monitor states: **free, held for reading, held for writing**
    - If held for reading, multiple readers can access simultaneously
      - we will need to know how many readers are accessing, and when they are done
  - Operations

- `mutex_lock()`: lock for writing
  - only acquires lock if it is free
  - sets state to **held for writing**
- `mutex_lock_read_only()`:
  - if lock is free, set its state to **held for reading**
  - if lock already held for reading, allow a new reader
  - increments a reader count
- `mutex_unlock()`:
  - if **held for writing**, set state to **free**
  - if **held for reading**, decrement reader count until zero, then set to **free**
- Fair access to reader-writer block
  - Policy question
    - if monitor state is held for reading and
      - A calls `mutex_lock()`, and blocks while waiting for monitor to be free
      - B calls `mutex_lock_read_only()`
  - Option 1 - Disallow new readers while a writer is waiting
    - affects a thread that could be running but now must wait
    - provides fair access to monitor (writer may have been waiting longer)
  - Option 2 - Allow new readers while a writer is waiting
    - increases concurrency, allows more threads to run
    - writer may need to wait for a long time to get access (starvation)
  - Solution
    - tradeoffs, depend on situation and application

## Apr. 8th - Unit 2c: Cont.d

- Semaphores - Edsger Dijkstra (THE system)
  - A non-negative atomic counter
    - any attempt to make counter negative will block the calling thread
    - no operation to read value; only to change it
  - **P(s)**: wait, dec, sub
    - From Dutch: *prober te verlagen* - "try lowering"
    - atomically blocks until  $s > 0$ , then decrements  $s$

- will be automatically woken up for resumption when  $s > 0$
- **V(s)**: signal, inc, add
  - From Dutch: *verhogen* - "to increase"
  - atomically increase  $s$ , and unblock threads waiting in  $P$  appropriate
- UThread semaphores

```

struct uthread_sem;
typedef struct uthread_sem * uthread_sem_t;

uthread_sem_t uthread_sem_create(int
initial_value);
void uthread_sem_destroy(uthread_sem_t);
void uthread_sem_wait(uthread_sem_t);
void uthread_sem_signal(uthread_sem_t);

```

- Drinking beer with semaphores
  - Use semaphore to store number of glasses held by pitcher
    - set initial value to empty (zero) when creating it

```

uthread_sem_t glasses =
uthread_sem_create(0);

```

- Pour and Refill no longer require a monitor
  - since the semaphore atomically changes the counter already

```

void pour(uthread_sem_t glasses) {
    uthread_sem_wait(glasses);
}

void refill(uthread_sem_t glasses, int n) {
    for (int i = 0; i < n; i += 1) {
        uthread_sem_signal(glasses);
    }
}

```

- Implementing monitors using semaphores
  - Implementing a mutex using semaphores
    - create semaphore with initial value 1 (free, representing an available mutex)
      - range of semaphore should be [0, 1]
    - lock is **P()**, **wait()**
    - unlock is **V()**, **signal()**
  - Implementing condition variables

- Difficult
  - In condition variables, signal without wait will be ignored / no effect
  - In semaphores, signal can unlock a future wait
- Example: disk read

```
// Asynchronous read request
void read(char* buf, int nbytes, int blockno) {
    scheduleRead(buf, nbytes, blockno);
    // what if thread gets pre-empted here?
    // even if signal is signaled before, the
    // following wait() can still execute as intended

    // semaphore is readComplete, initialized to 0
    // thread for read becomes blocked
    uthread_sem_wait(readComplete);
    // handle data
}

// Read completion (called by disk ISR)
void onReadComplete() {
    // increase readComplete and wake up a waiter
    uthread_sem_signal(readComplete);
}
```

- No critical section, no wait-signal race problem
- Blocking queue using semaphores

- Semaphores:
  - **mx** - starts at 1 (like a mutex, range [0, 1])
  - **length** - starts at 0
- **lock**  $\iff$  **wait**
- **unlock**  $\iff$  **signal**

- ```
void enqueue (queue_t* queue, item_t item) {
    uthread_sem_wait (queue->mx);
    item->next = NULL;
    if (queue->tail)
        queue->tail->next = item;
    queue->tail = item;
    if (queue->head == NULL)
        queue->head = queue->tail;
    uthread_sem_signal (queue->mx);
    uthread_sem_signal (queue->length);
}
```

- ```

item_t dequeue (queue_t* queue) {
    // locks if queue is empty
    item_t item = NULL;
    uthread_sem_wait (queue->length);
    uthread_sem_wait (queue->mx);
    item = queue->head;
    queue->head = queue->head->next;
    if (queue->head) == NULL)
        queue->tail = NULL;
    item->next = 0;
    uthread_sem_signal (queue->mx);
    return item;
}

```

- Why are there no loops?
  - "State variable" is a semaphore, read/write occurs atomically, no chance of state being changed since waking up no need to re-check state
- Why wait on length outside critical section?
  - avoid deadlock
- Why signal on length outside critical section?
  - It's fine, dequeue go back to sleep briefly on `mx`

## Apr. 10th - Unit 2c: Cont.d

- Ordering threads using semaphores
  - If thread B must wait for thread A to finish
    - Initialize semaphore `b` to 0, i.e.
    - `uthread_sem_t b = uthread_sem_create(0)`
    - Thread A: `uthread_sem_signal(b)`
    - Thread B: `uthread_sem_wait(b)`
  - If both threads need to wait for each other (rendezvous)
    - Initialize two semaphores `a`, `b` with 0
      - ```

// Thread A
uthread_sem_signal(a);
uthread_sem_wait(b);

// Thread B
uthread_sem_signal(b);
uthread_sem_wait(a);

```
- If the order of wait and signal are reversed for either (or both) threads
  - If reversed for A, both threads run fine
  - If reverse both of them, then deadlock
- Barriers using semaphores

- Local, non-reusable barrier
  - Like thread join, but for any thread rather than a specific one
  - Create semaphore with initial value 0
  - Create  $N$  threads, each thread calls signal on the semaphore
  - Main thread can wait for any created thread to finish
  - To wait for all of them, just wait  $N$  times
  - Can also be used to wait for some milestone
- Analogy: job application with reference contact information, potential employer receives application and contacts references directly, but will not process application until hearing back from 2 references
- Problems with concurrency
  - Race condition
    - Competing, unsynchronized access to shared variable
      - from multiple threads
      - at least one of the threads is attempting to update the variable
    - Solved with synchronization
      - guaranteeing mutual exclusion for competing access
      - but the language does not help you see what data might be shared - can be very hard
  - Deadlock
    - multiple competing actions wait for each other
    - all actions are prevented from completing
  - Livelock
    - Threads respond to actions by other threads, but other threads also respond
    - Threads are unable to make progress
    - ```
while (true) {
    lock1.lock();
    if (!lock2.tryLock()) {
        lock1.unlock();
        lock2.lock();
        if (!lock1.tryLock()) {
            lock2.unlock();
            continue;
        }
    }
    break;
}
```

- Starvation
  - A thread is unable to gain regular access to a shared resource
  - "Greedy" threads may lock resource for a long time
  - Threads with higher priority may skip ahead
  - Reader threads may acquire reader-writer lock in front of writer
- Systems with multiple mutexes
  - Consider a system with 2 mutexes: **a**, **b**
  - ```
void foo() {
    uthread_mutex_lock(a);
    uthread_mutex_unlock(a);
}

void bar() {
    uthread_mutex_lock(b);
    uthread_mutex_unlock(b);
}

void x() {
    uthread_mutex_lock(a);
    bar();
    uthread_mutex_unlock(a);
}

void y() {
    uthread_mutex_lock(b);
    foo();
    uthread_mutex_lock(b);
}
```
  - If **x()** calls **foo()**: thread becomes blocked by itself
  - With **y()**: threads for **x**, **y**, if each acquire a/b before calling bar/foo → deadlock
- Concurrency problems with recursion
  - If a recursive call attempts to acquire a lock
    - Lock is already being held, so thread blocks
    - Thread is waiting for lock to be released
    - But that same thread is already holding the lock
  - If we release lock before calling recursively
    - Could break critical section's problem
  - Solution: Re-entrant mutex

- Allows a lock to be acquired more than once (only if it is acquired again by the same thread)
- Unlock only releases the lock if called as many times as lock was called
- Each lock operation (from the same thread) increments counter
- Unlock decrements it, and releases lock when counter is 0
- Identifying deadlocks with waiter graphs
  - Waiter graph
    - Vertices for threads and locks
    - edge from lock to thread, if lock is **HELD** by thread
    - edge from thread to lock, if lock is **WAITING** for lock
    - a cycle indicates deadlock
- The dining philosophers problem
  - 5 computers competed for access to 5 shared tape drivers
  - Description
    - 5 philosophers sit at a round table with one fork placed in between each
    - Each philosopher is either eating, or thinking
      - not doing two things simultaneously
      - they never speak to each other
    - A large bowl of spaghetti in the middle of the table requires 2 forks to serve
  - Deadlock
    - Assume that philosophers always start with the left fork
    - Also assume that all philosophers decide to start eating at the same time
      - Everyone is able to get left fork
      - But everyone waits for right fork
  - Livelock
    - Assume that, if philosophers can't get second fork, they release the first fork, then wait on second
    - If the process is repeated, and all philosophers are synchronized
      - Philosophers will repeatedly get one fork at a time
      - All are busy, but unable to eat
- Avoiding deadlock
  - If not necessary, don't use multiple threads
    - will have many idle CPU cores and write asynchronous code
  - If not necessary, don't use shared variables

- Use local variables and parameters whenever possible
  - if threads don't access shared data, no need for synchronization
- Avoid unnecessary locks
  - When possible, use atomic data structures and lock-free synchronization
  - If possible, use only one lock at a time
- Evaluate the scope of the lock
  - can possibly avoid deadlock by reducing portion of code that needs lock?
- Organize locks into precedence hierarchy
  - each lock is assigned a unique precedence number
  - before thread  $X$  acquires a lock  $i$ , it must hold all higher precedence locks
  - ensures that any thread holding lock  $i$  can not be waiting for  $X$ 
    - so that locks are always acquired in the same order
- Limit wait time on locks
  - Provide an alternate action if lock cannot be acquired
- Detect and destroy
  - If cannot avoid deadlock, detect when it has occurred
  - break deadlock by interrupting/terminating threads